

# Computer Vision

Bogdan Alexe

[bogdan.alexe@fmi.unibuc.ro](mailto:bogdan.alexe@fmi.unibuc.ro)

University of Bucharest, 2<sup>nd</sup> semester, 2020-2021

# Course structure

## 1. Features and filters: low-level vision

Linear filters, color, texture, edge detection

## 2. Grouping and fitting: mid-level vision

Fitting curves and lines, robust fitting, RANSAC, Hough transform, segmentation

## 3. Multiple views

Local invariant feature and description, epipolar geometry and stereo, object instance recognition

## 4. Object Recognition: high – level vision

Object classification, object detection, part based models, bovw models

## 5. Video understanding

Object tracking, background subtraction, motion descriptors, optical flow

# Lab class 2

## Automatic grading of multiple choice tests

- a simpler version of the last  
year's first project



### TEST GRILĂ

INFORMATICĂ

FIZICĂ

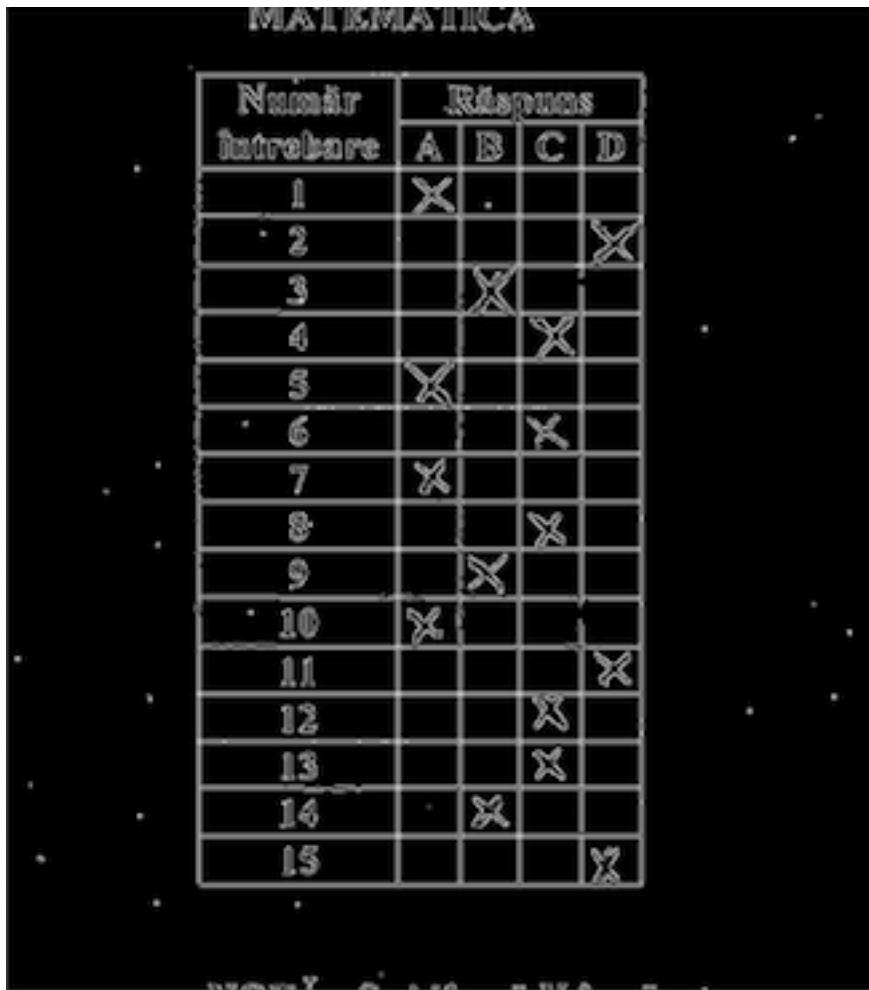
#### MATEMATICĂ

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2		X		
3		X		
4	X			
5		X		
6		X		
7	X			
8			X	
9			X	
10	X			
11		X		
12		X		
13			X	
14			X	
15			X	

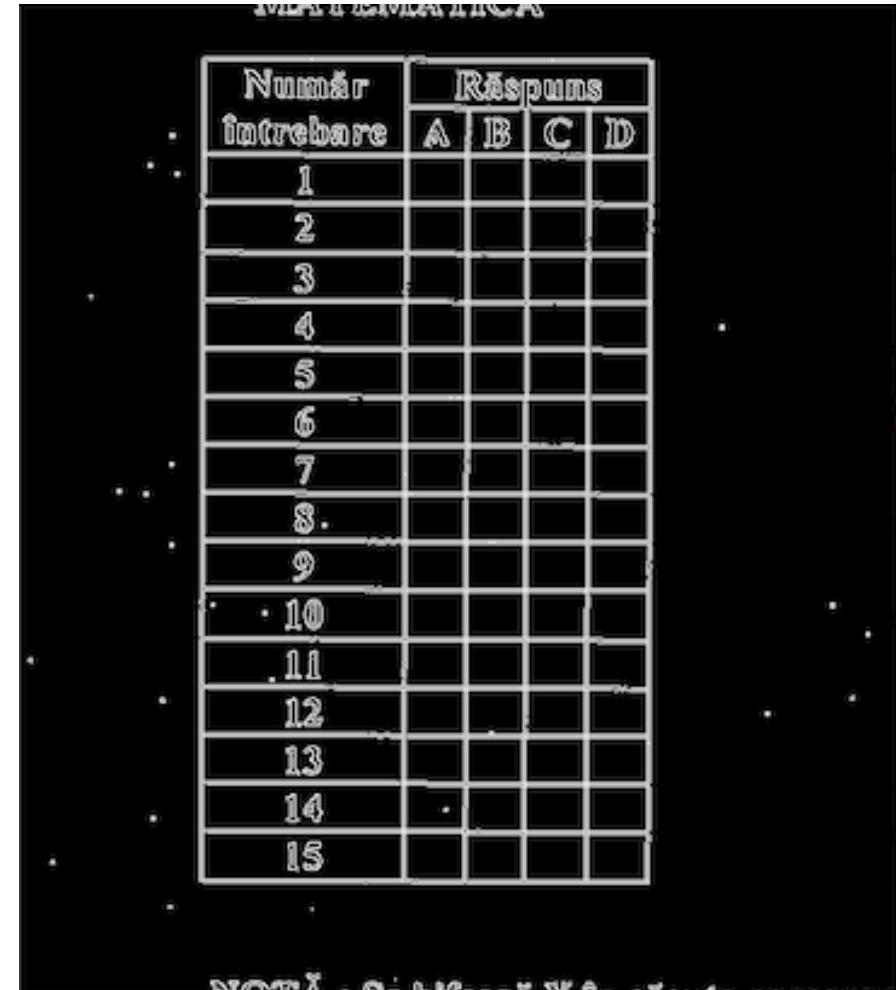
Număr întrebare	Răspuns			
	A	B	C	D
1	X			
2			X	
3			X	
4	X			
5			X	
6	X			
7	X			
8		X		
9			X	
10			X	
11			X	
12			X	
13			X	
14	X			
15			X	

NOTĂ : Se bifează X în căsuța corespunzătoare răspunsului corect.

# Lab class 2 – Automatic grading of multiple choice tests



Edge image for query



Edge image for template

# Lab class 2 – Automatic grading of multiple choice tests

MATEMATICA

Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2	X			
3		X		
4	X			
5		X		
6		X		
7	X			
8			X	
9		X		
10	X			
11		X		
12	X			
13		X		
14			X	
15			X	

NOTĂ : Se bifă o săgeată în cadrul

Horizontal lines found with Hough transform

MATEMATICA

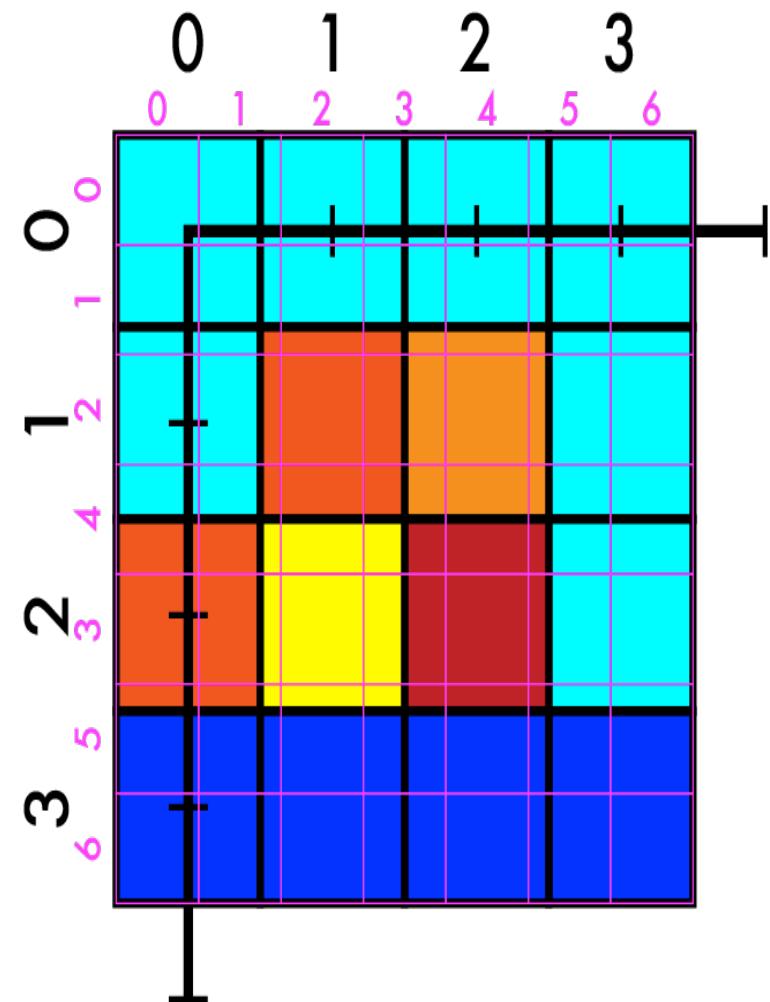
Număr întrebare	Răspuns			
	A	B	C	D
1			X	
2	X			
3			X	
4	X			
5			X	
6			X	
7	X			
8				X
9			X	
10	X			
11		X		
12	X			
13			X	
14			X	
15			X	

NOTĂ : Se bifă o săgeată în cadrul

Vertical lines found with Hough transform

# Resize 4 x 4 to 7 x 7

- Create our new image
- Match up coordinates

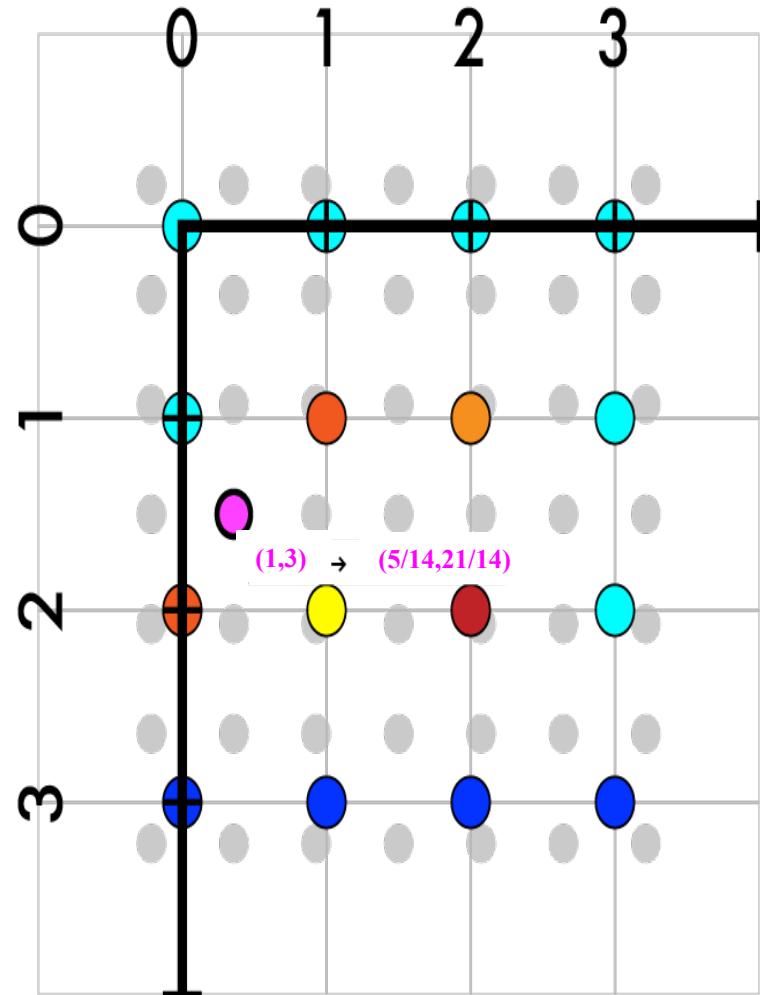


# Resize 4 x 4 to 7 x 7

- Create our new image
- Match up coordinates

$$\begin{pmatrix} x_1 \\ y_1 \end{pmatrix} = \begin{pmatrix} \frac{4}{7} & 0 \\ 0 & \frac{4}{7} \end{pmatrix} \begin{pmatrix} x_2 \\ y_2 \end{pmatrix} + \begin{pmatrix} -\frac{3}{14} \\ -\frac{3}{14} \end{pmatrix}$$

- Iterate over new points
  - Map to old coordinates
  - Take point (1,3)
  - Map it to (5/14, 21/14)
  - Interpolate old values



# Bilinear interpolation

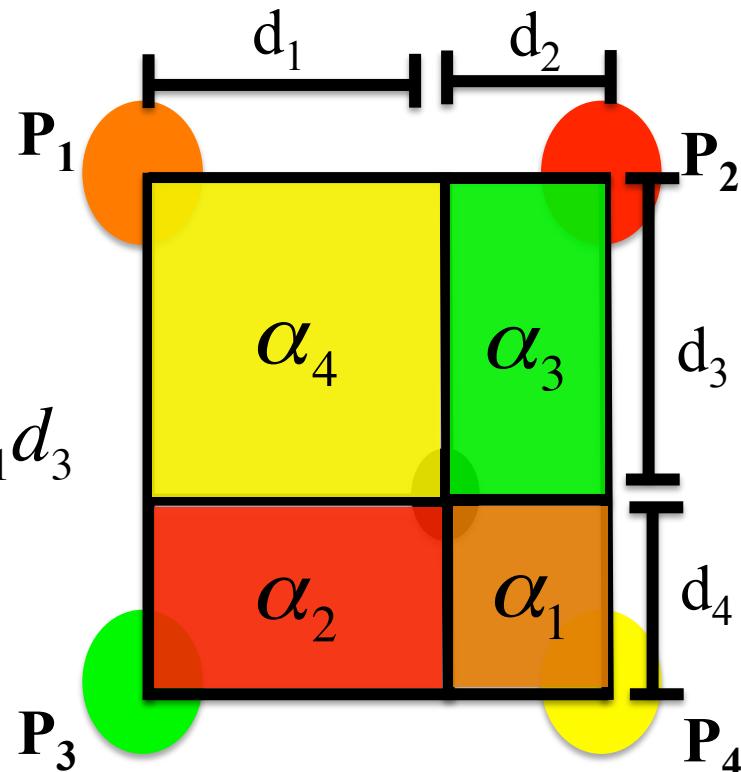
Find the closest pixels in a box

Determine coefficients  $\alpha_1, \alpha_2, \alpha_3, \alpha_4$   
such that:

$$q = \alpha_1 P_1 + \alpha_2 P_2 + \alpha_3 P_3 + \alpha_4 P_4$$

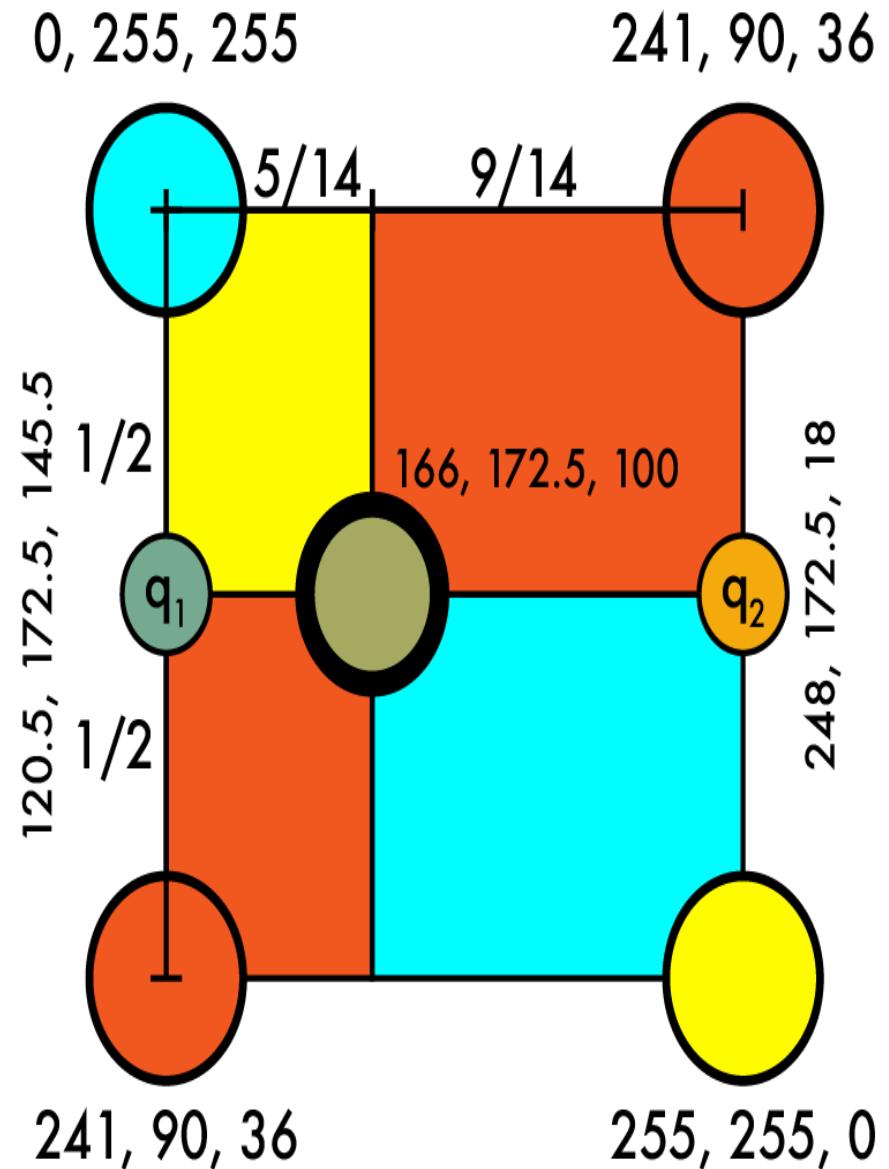
$$\alpha_1 = d_2 d_4, \alpha_2 = d_1 d_4, \alpha_3 = d_2 d_3, \alpha_4 = d_1 d_3$$

$\alpha_i$  - area of the opposite rectangle to  
pixel  $P_i$



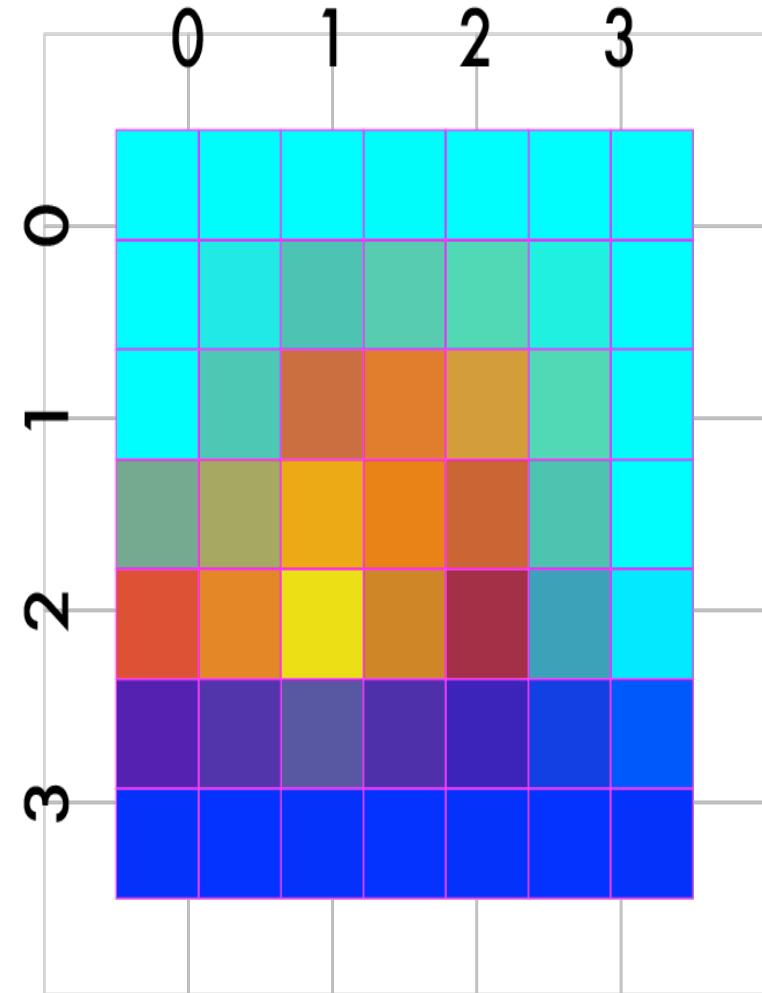
# Resize 4 x 4 to 7 x 7

- Create our new image
- Match up coordinates
- Iterate over new points
  - Map to old coordinates
  - Take point  $(3,1)$
  - Map it to  $(21/14, 5/14)$
  - Interpolate old values
    - Size of opposite rectangles
    - OR find  $q_1$  and  $q_2$ , then interpolate between them
  - $q_1 = (120.5, 172.5, 145.5)$
  - $q_2 = (248, 172.5, 18)$
  - $q = (166, 172.5, 100)$

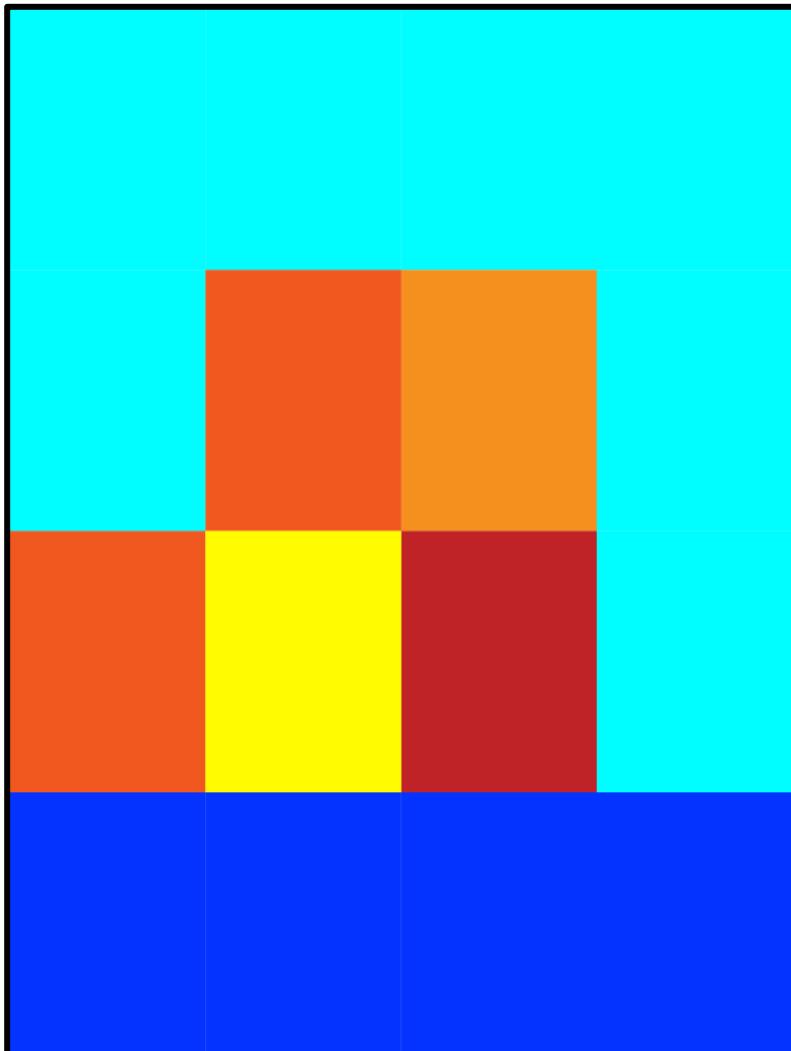


# Resize 4 x 4 to 7 x 7

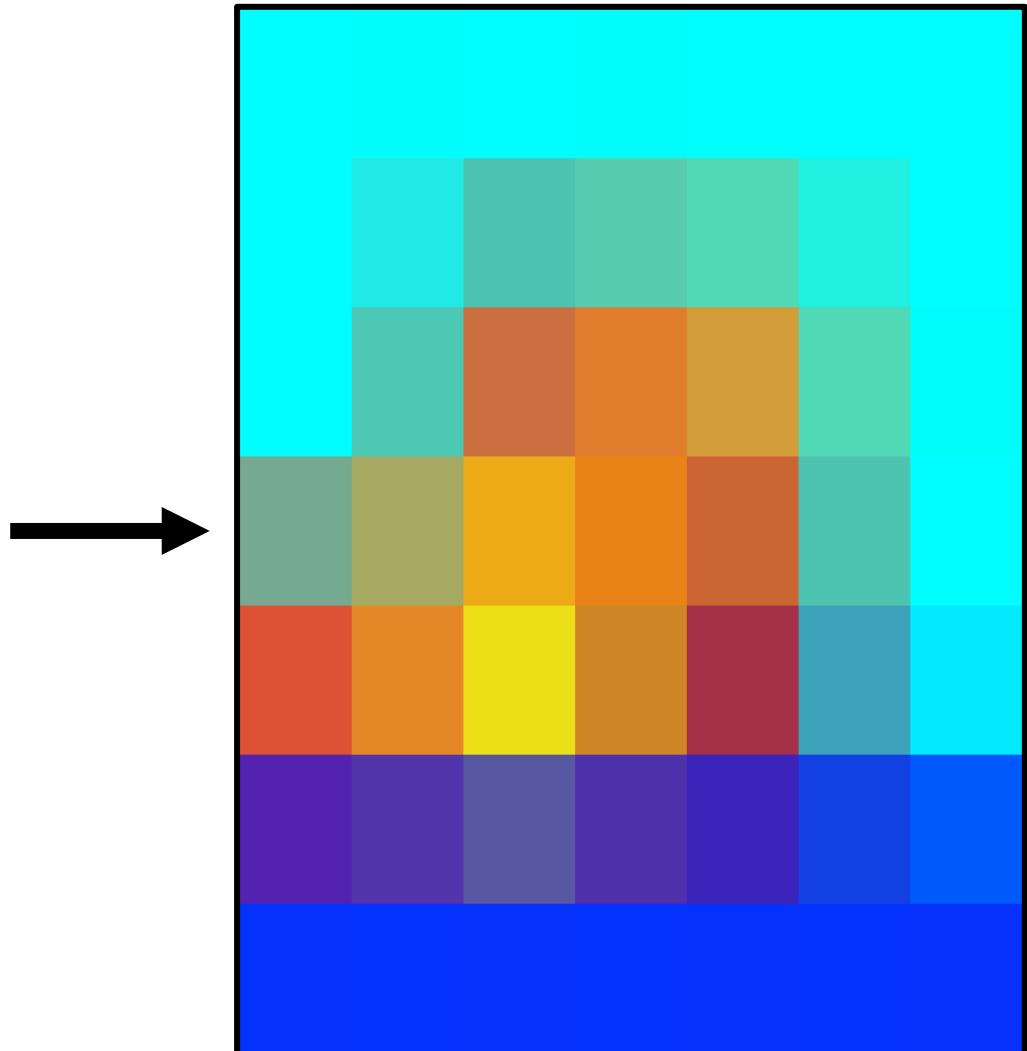
- Create our new image
- Match up coordinates
- Iterate over new points
  - Map to old coordinates
  - Take point  $(3,1)$
  - Map it to  $(21/14, 5/14)$
  - Interpolate old values
    - $q = (166, 172.5, 100)$
- Fill in the rest



# Final result!

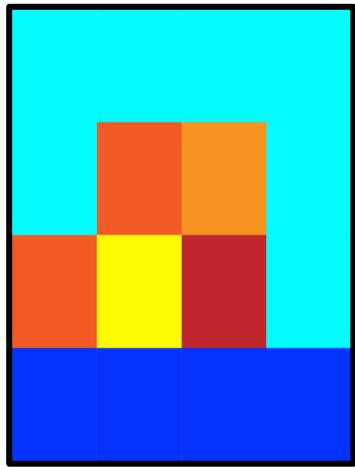


Original image  $4 \times 4$

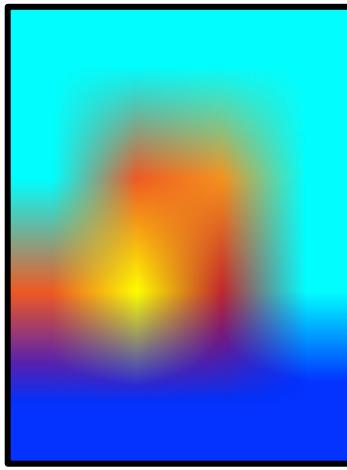


Resized image  $7 \times 7$  using  
bilinear interpolation

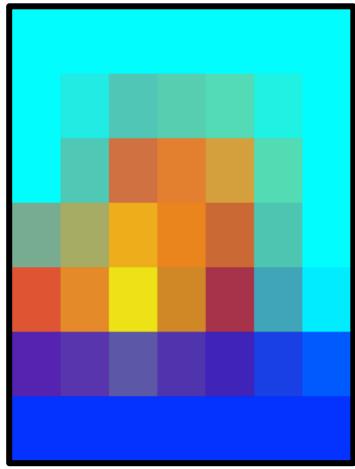
# Different scales



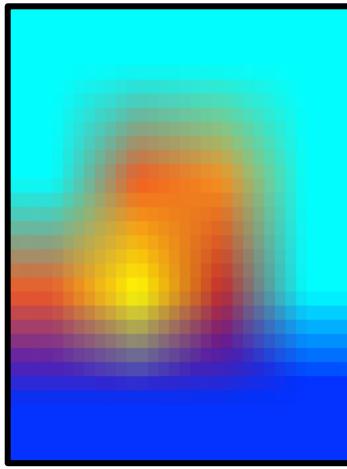
Original image  $4 \times 4$



Resized image  $256 \times 256$   
using bilinear interpolation



Resized image  $7 \times 7$  using  
bilinear interpolation

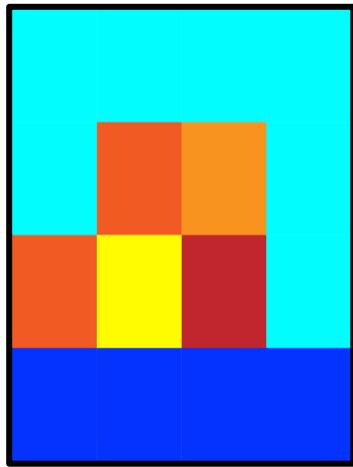


Resized image  $32 \times 32$   
using bilinear interpolation

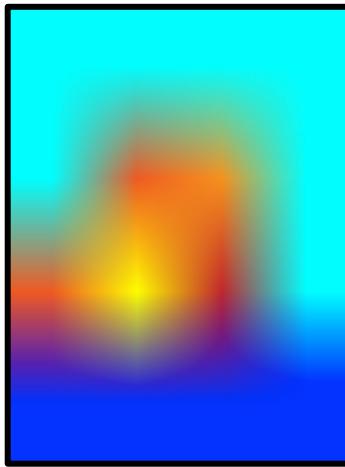


Star like pattern

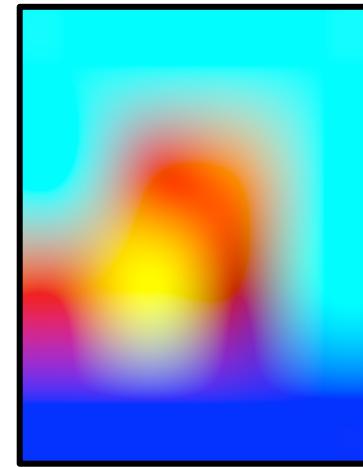
# Different methods



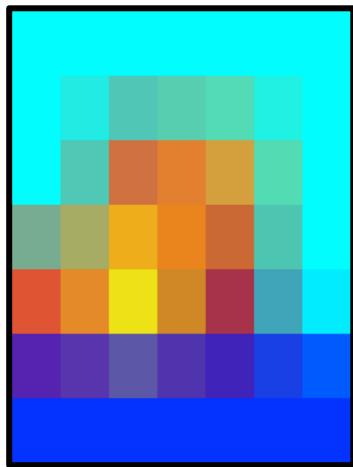
Original image  $4 \times 4$



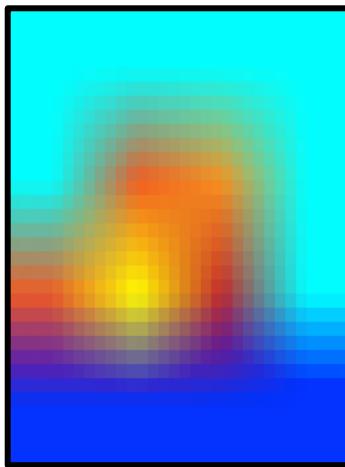
Resized image  $256 \times 256$   
using bilinear interpolation



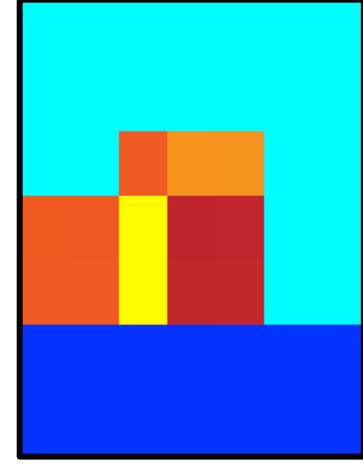
Resized image  $256 \times 256$   
using bicubic interpolation



Resized image  $7 \times 7$  using  
bilinear interpolation

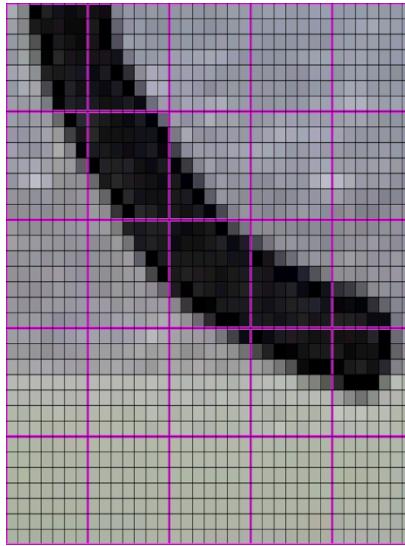


Resized image  $32 \times 32$   
using bilinear interpolation

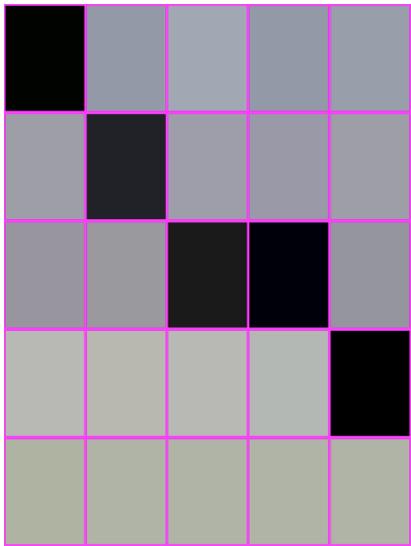


Resized image  $7 \times 7$  using  
nearest neighbor interpolation

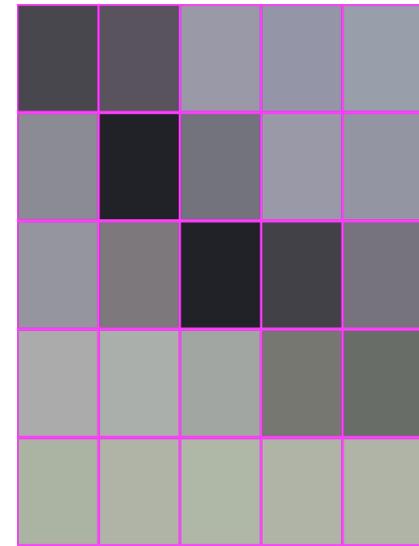
# Shrinking an image: from 448 x 448 to 64 x 64



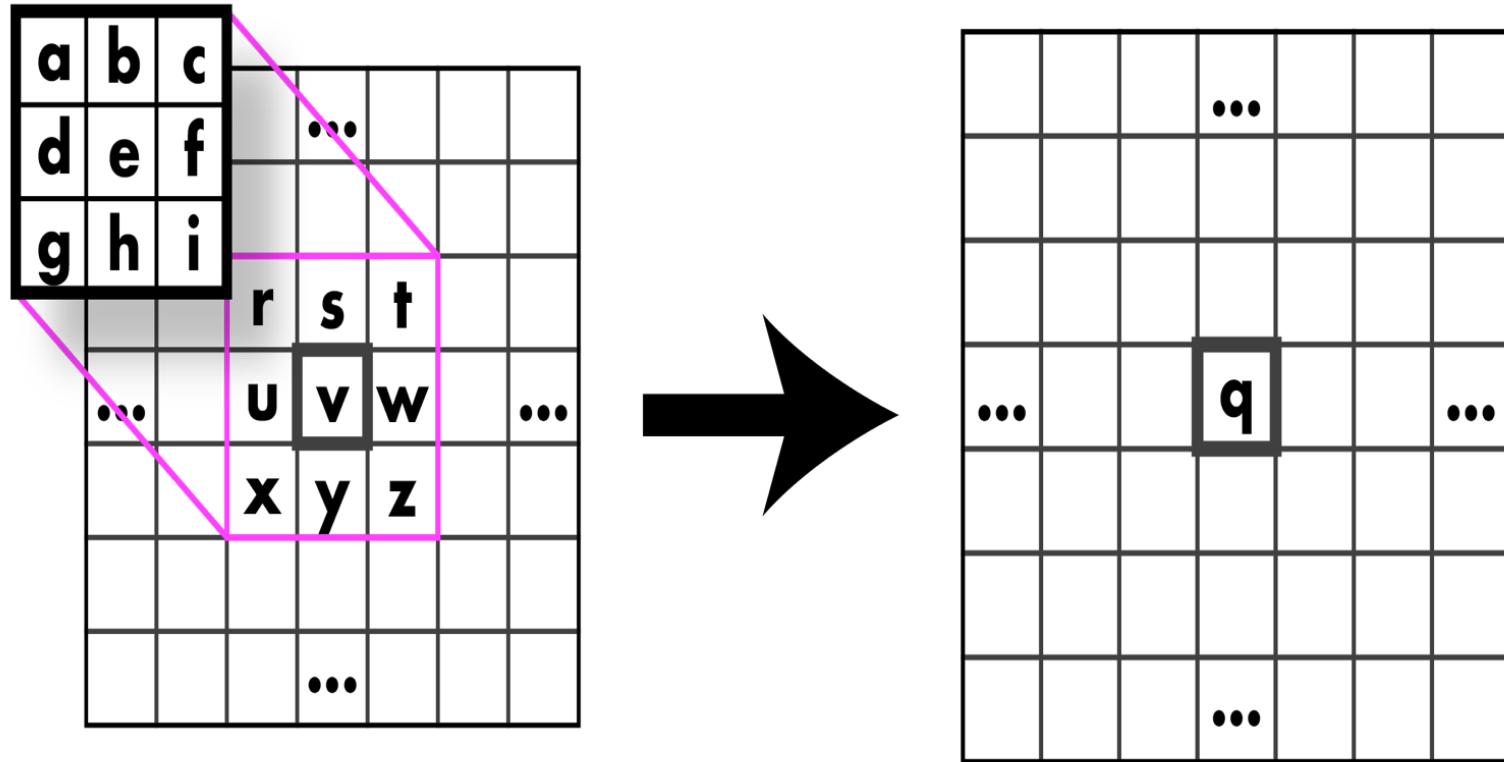
Nearest neighbor or  
bilinear “interpolation”



averaging

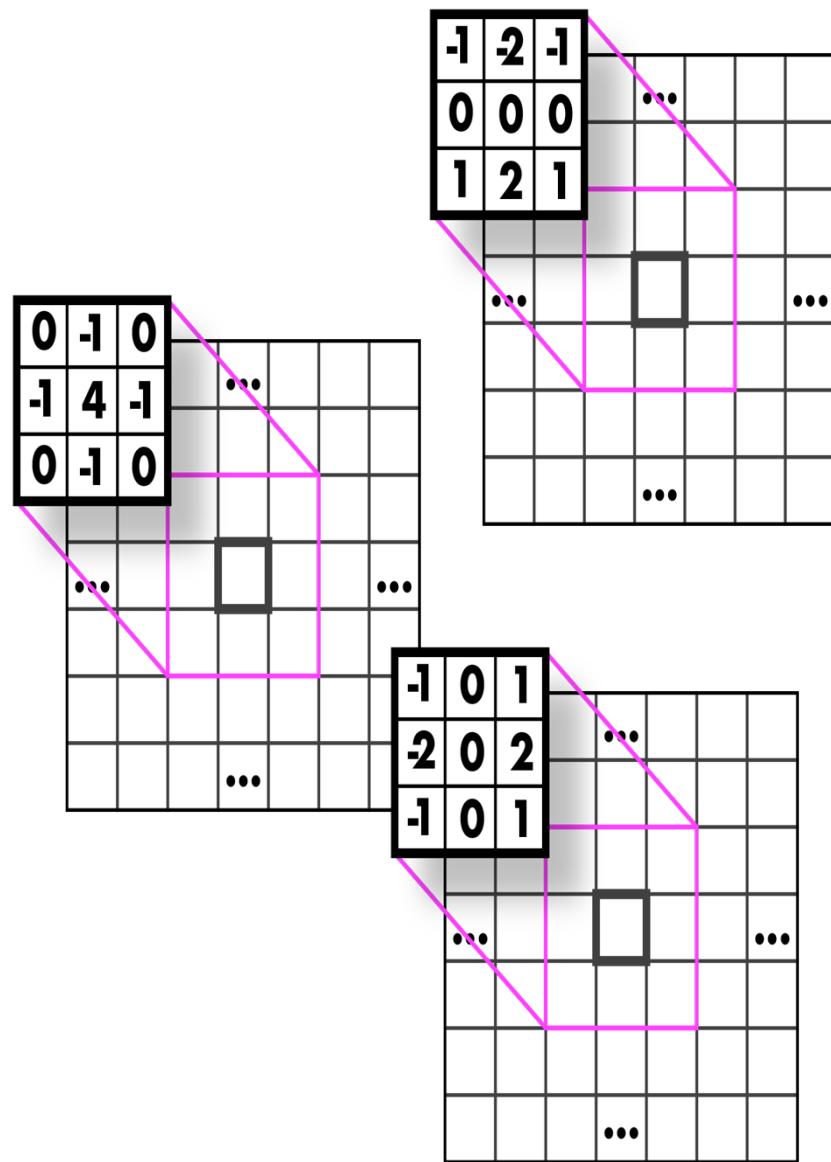
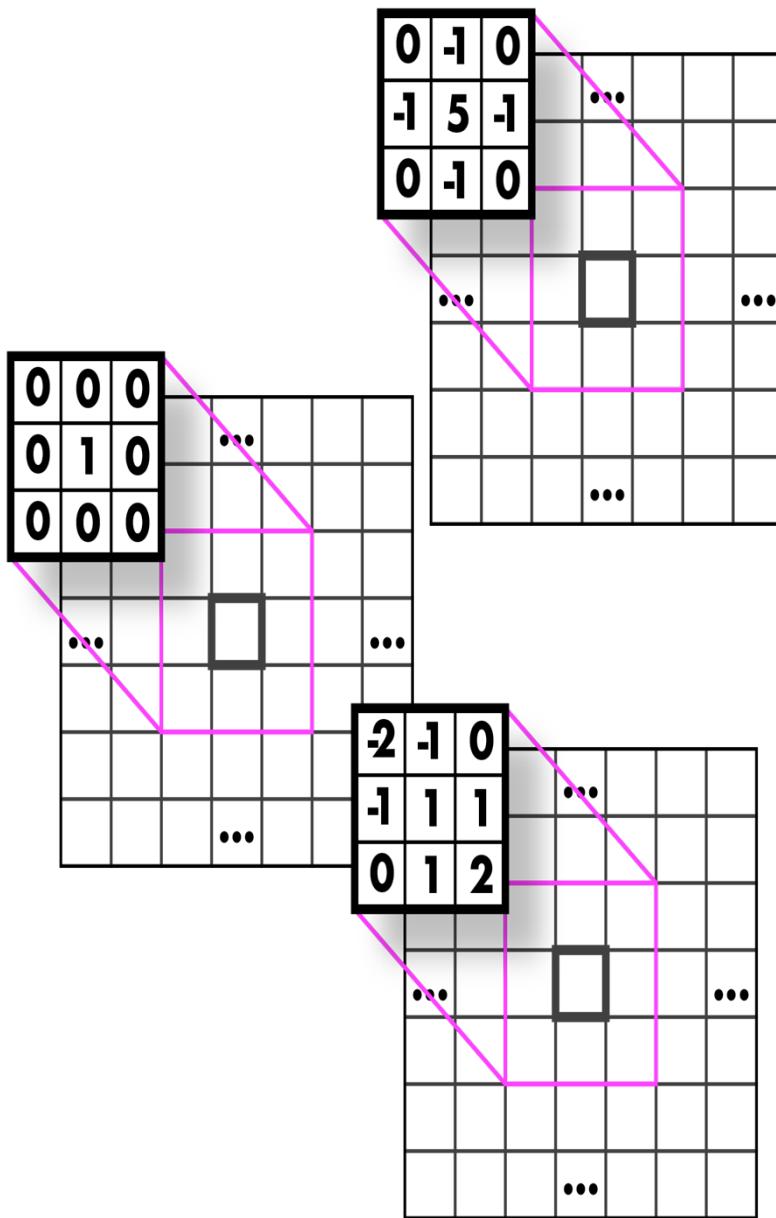


# Convolution: Weighted sum over pixels



$$q = a \times r + b \times s + c \times t + d \times u + e \times v + f \times w + g \times x + h \times y + i \times z$$

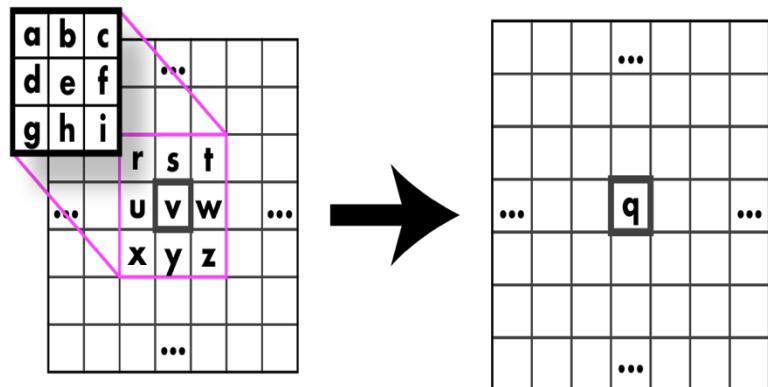
# Filters



# Cross-Correlation vs Convolution

## Cross-Correlation

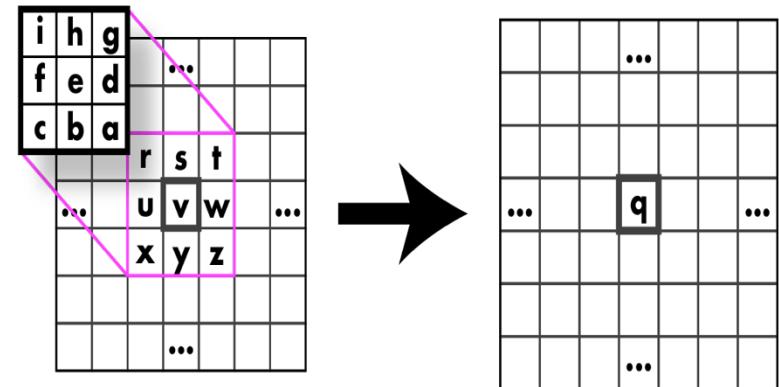
$$\left( \begin{array}{ccc} a & b & c \\ d & e & f \\ g & h & i \end{array} \right) \star \left( \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline \end{array} \right)$$



$$q = a \times r + b \times s + c \times t + d \times u + e \times v + f \times w + g \times x + h \times y + i \times z$$

## Convolution

$$\left( \begin{array}{ccc} a & b & c \\ d & e & f \\ g & h & i \end{array} \right) * \left( \begin{array}{|c|c|c|c|c|c|c|c|c|c|} \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline & & & & & & & & & \\ \hline \end{array} \right)$$



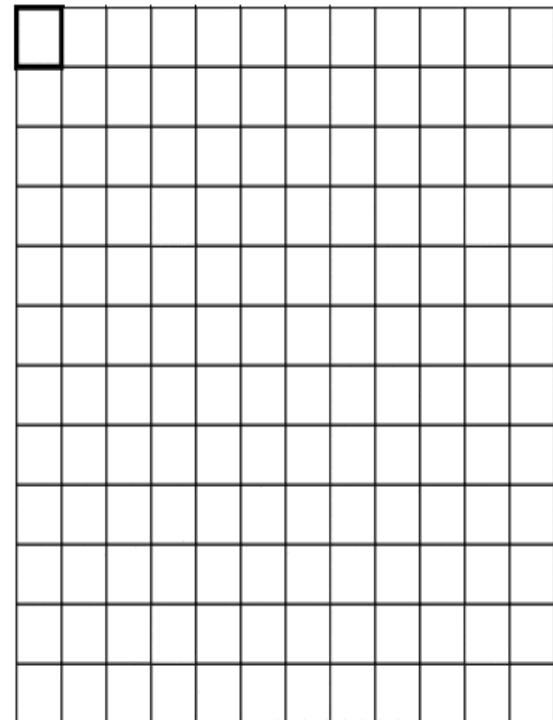
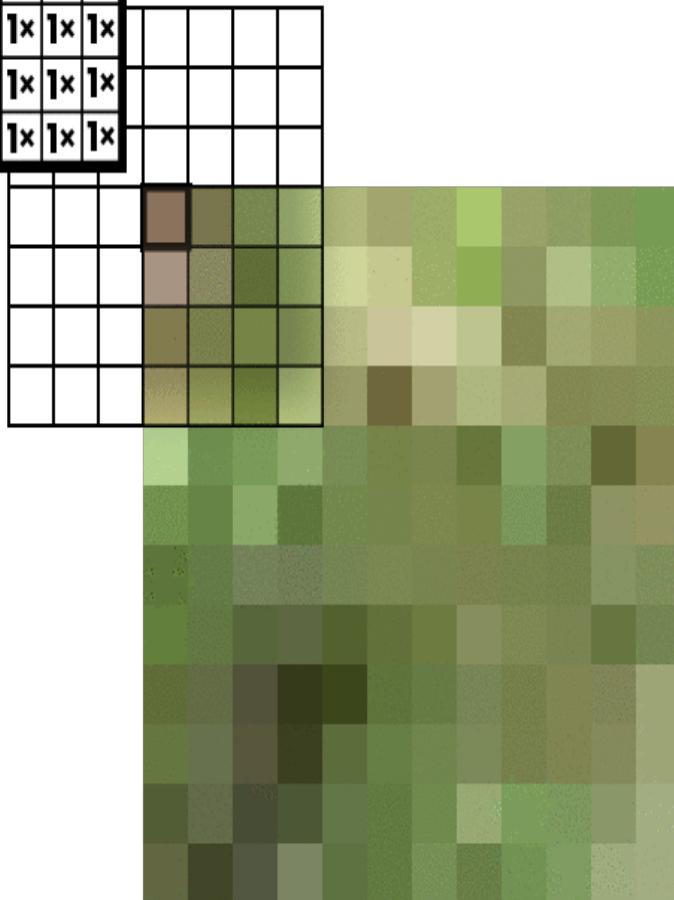
$$q = i \times r + h \times s + g \times t + f \times u + e \times v + d \times w + c \times x + b \times y + a \times z$$

Rotate the filter:  
left-right and up-down

# Kernel slides across image

$\frac{1}{49}$

1x							
1x							
1x							
1x							
1x							
1x							
1x							



# Image filtering

- For each pixel, compute function of local neighborhood and output a new value
  - Same function applied at each position
  - Output and input image are typically the same size

# Image filtering

- Linear filtering: function is a weighted sum/difference of pixel values
- Really important!
  - Enhance images
    - Denoise, smooth, increase contrast (sharpening), etc.
  - Extract information from images
    - Texture, edges, distinctive points, etc.
  - Detect patterns
    - Template matching

# Example: box filter

$$\frac{1}{9} \begin{matrix} g[\cdot, \cdot] \\ \begin{array}{|c|c|c|} \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline 1 & 1 & 1 \\ \hline \end{array} \end{matrix}$$

# Image filtering

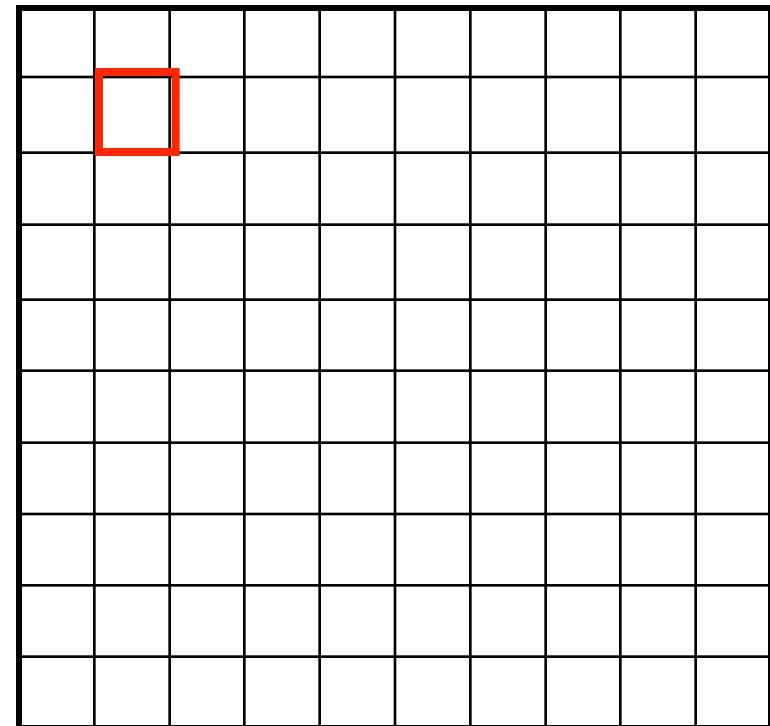
$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

$h[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0



$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

0	10								

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] = \frac{1}{9} \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix}$$

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$


$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

$$f[.,.]$$

$$h[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0


$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$


0    10    20    30    30

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

1	1	1
1	1	1
1	1	1

$f[.,.]$

0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	90	0	90	90	90	0	0	0
0	0	0	90	90	90	90	90	0	0	0
0	0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0

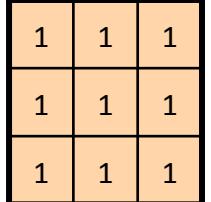
$h[.,.]$

			0	10	20	30	30			

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$


$f[.,.]$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$h[.,.]$

			0	10	20	30	30		

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

# Image filtering

$$g[\cdot, \cdot] \frac{1}{9}$$

$$f[.,.]$$

0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	90	0	90	90	90	0	0
0	0	0	90	90	90	90	90	0	0
0	0	0	0	0	0	0	0	0	0
0	0	90	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0

$$h[.,.]$$

	0	10	20	30	30	30	20	10	
	0	20	40	60	60	60	40	20	
	0	30	60	90	90	90	60	30	
	0	30	50	80	80	90	60	30	
	0	30	50	80	80	90	60	30	
	0	20	30	50	50	60	40	20	
	10	20	30	30	30	30	20	10	
	10	10	10	0	0	0	0	0	

$$h[m, n] = \sum_{k,l} g[k, l] f[m + k, n + l]$$

Slide credit: Steve Seitz

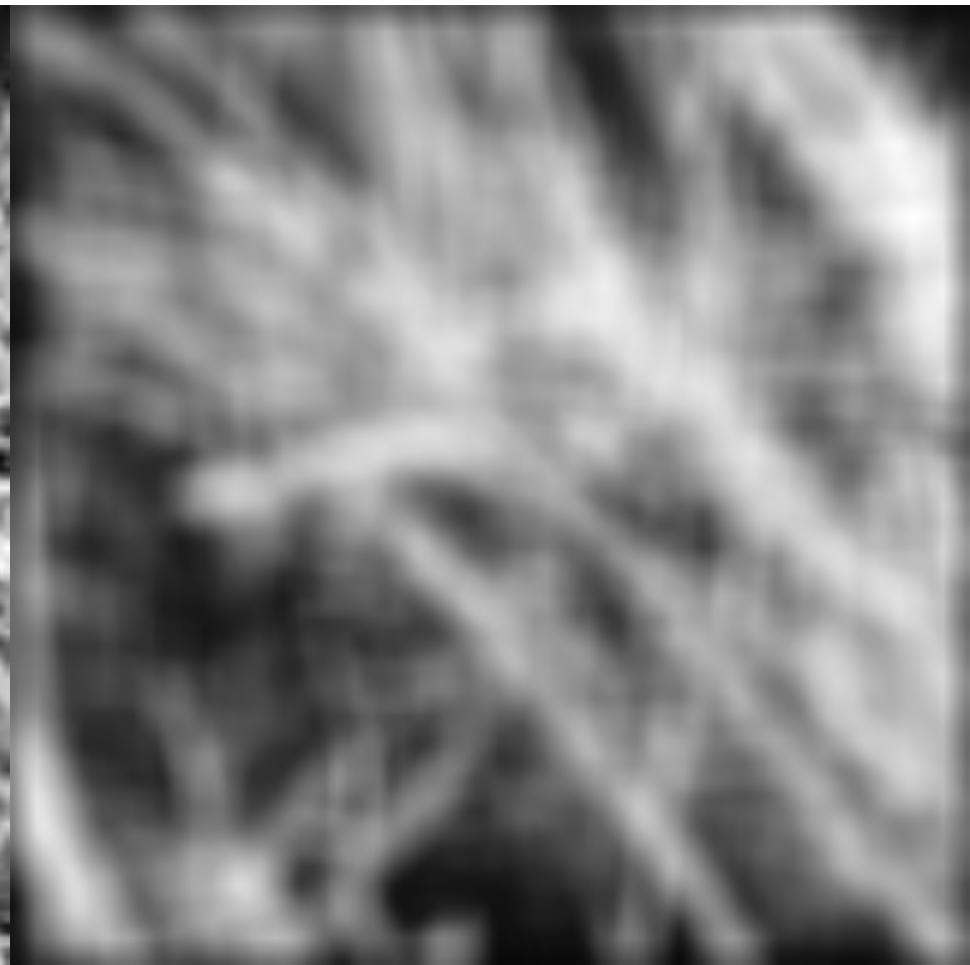
# Box filter

What does it do?

- Replaces each pixel with an average of its neighborhood
- Achieve smoothing effect  
(remove sharp features)

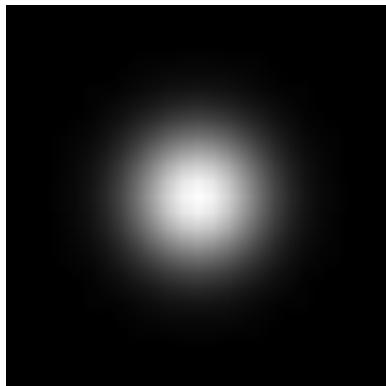
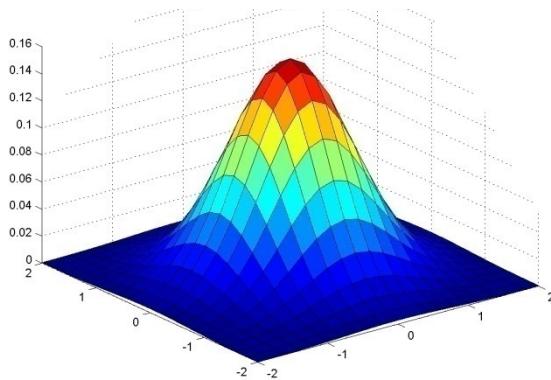
$$g[\cdot, \cdot] = \frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$

# Smoothing with box filter



# Important filter: Gaussian filter

- Spatially-weighted average



0.003	0.013	0.022	0.013	0.003
0.013	0.059	0.097	0.059	0.013
0.022	0.097	0.159	0.097	0.022
0.013	0.059	0.097	0.059	0.013
0.003	0.013	0.022	0.013	0.003

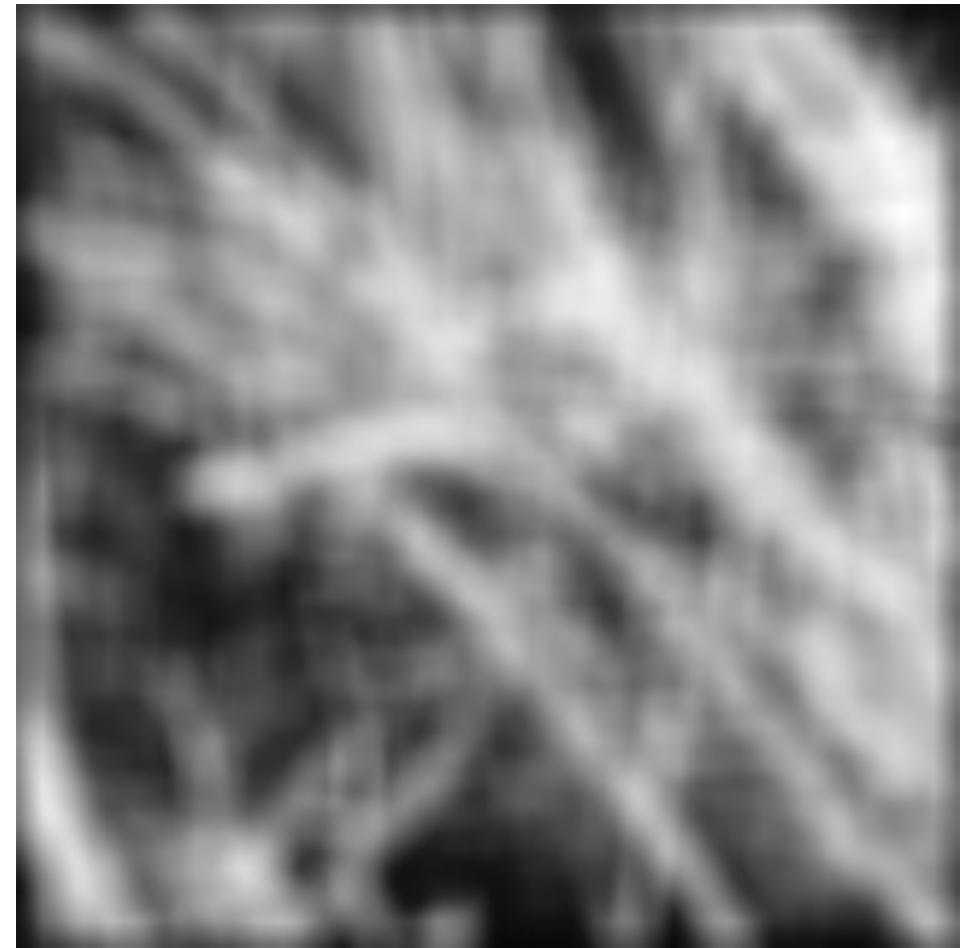
$5 \times 5, \sigma = 1$

$$G_\sigma = \frac{1}{2\pi\sigma^2} e^{-\frac{(x^2+y^2)}{2\sigma^2}}$$

# Smoothing with Gaussian filter



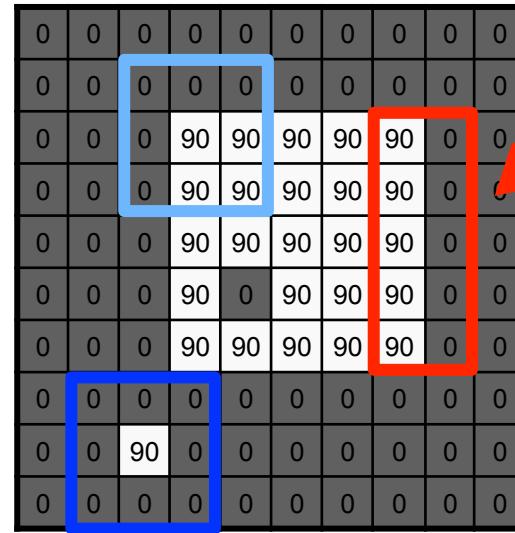
# Smoothing with box filter



# Smoothing an image

Box filter

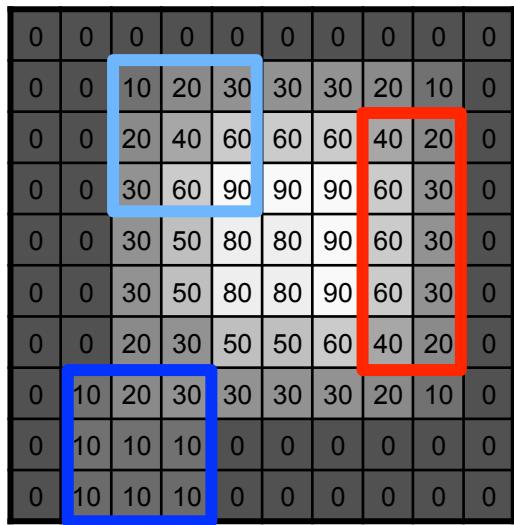
$$\frac{1}{9} \begin{matrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{matrix}$$



Vertical edge

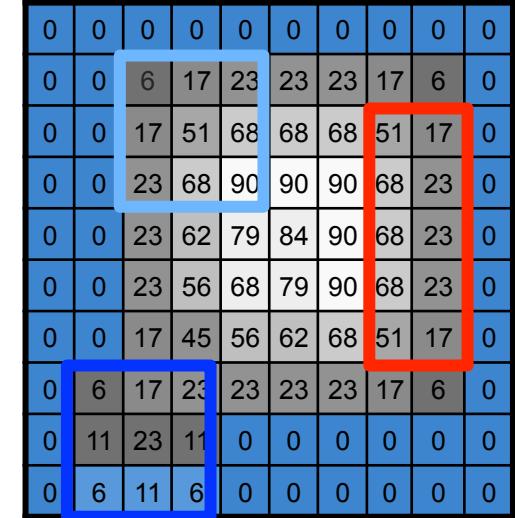
Gaussian filter

$$\frac{1}{16} \begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix}$$



OBSERVATIONS:

1. Remove regions with rapid changes in image intensity function (high frequencies regions) - edges
2. Distances between pixel values are getting smaller



# Gaussian filters

- Remove “high-frequency” components from the image (low-pass filter)
  - Images become more smooth
- Convolution with self is another Gaussian
  - So can smooth with small-width kernel, repeat, and get same result as larger-width kernel would have
  - Convolving two times with Gaussian kernel of width  $\sigma$  is same as convolving once with kernel of width  $\sigma\sqrt{2}$
- *Separable* kernel
  - Factors into product of two 1D Gaussians

# Separability of the Gaussian filter

$$\begin{aligned} G_\sigma(x, y) &= \frac{1}{2\pi\sigma^2} \exp^{-\frac{x^2 + y^2}{2\sigma^2}} \\ &= \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{x^2}{2\sigma^2}} \right) \left( \frac{1}{\sqrt{2\pi}\sigma} \exp^{-\frac{y^2}{2\sigma^2}} \right) \end{aligned}$$

The 2D Gaussian can be expressed as the product of two functions, one a function of  $x$  and the other a function of  $y$

In this case, the two functions are the (identical) 1D Gaussian

# Separability example

2D filtering  
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform filtering  
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 \\ 18 \\ 18 \end{matrix}$$

Followed by filtering  
along the remaining column:

# Separability example

2D filtering  
(center location only)

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix}$$

The filter factors  
into a product of 1D  
filters:

$$\begin{matrix} 1 & 2 & 1 \\ 2 & 4 & 2 \\ 1 & 2 & 1 \end{matrix} = \begin{matrix} 1 \\ 2 \\ 1 \end{matrix} \times \begin{matrix} 1 & 2 & 1 \end{matrix}$$

Perform filtering  
along rows:

$$\begin{matrix} 1 & 2 & 1 \end{matrix} * \begin{matrix} 2 & 3 & 3 \\ 3 & 5 & 5 \\ 4 & 4 & 6 \end{matrix} = \begin{matrix} 11 \\ 18 \\ 18 \end{matrix}$$

Followed by filtering  
along the remaining column:

$$\begin{matrix} 1 \\ 2 \\ 1 \end{matrix} * \begin{matrix} 11 & & \\ 18 & & \\ 18 & & \end{matrix} = \begin{matrix} 65 \end{matrix}$$

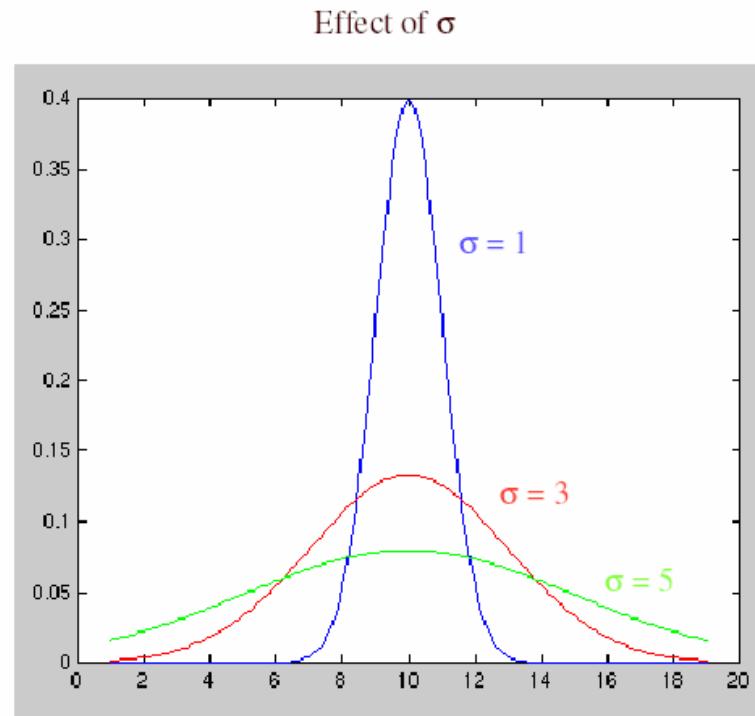
# Separability

- Why is separability useful in practice?
  - Separable filters are more efficient – they employ only  $O(n)$  multiplications instead of  $O(n^2)$  multiplications by non-separable filters ( $n$  is the size of the filter)

# Practical matters

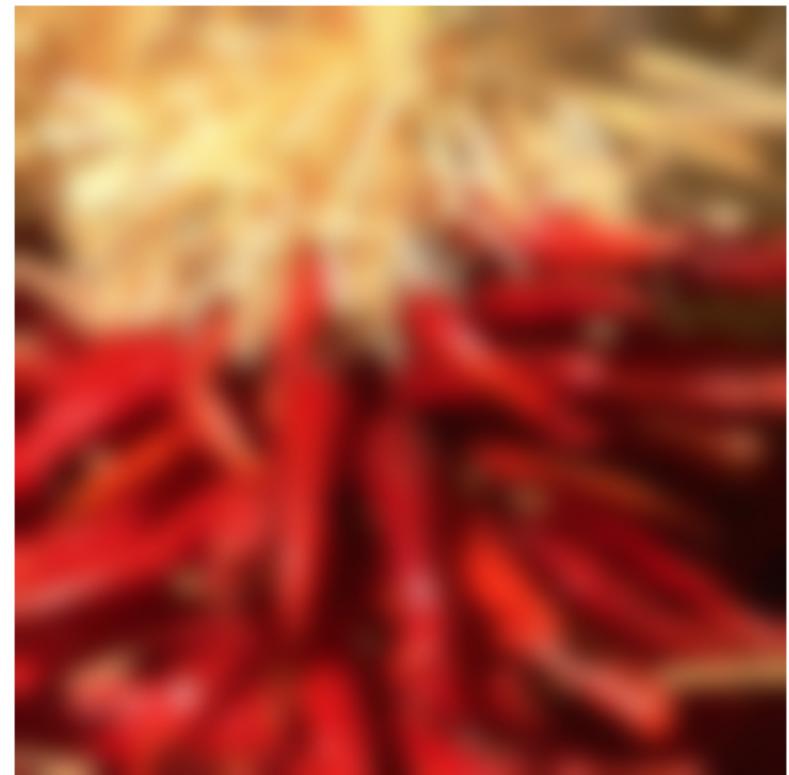
How big should the filter be?

- Values at edges should be near zero  $\leftarrow$  important!
- Rule of thumb for Gaussian: set filter half-width to about  $3 \sigma$



# Practical matters

- What about near the edge?
  - the filter window falls off the edge of the image
  - need to extrapolate
  - methods:
    - clip filter (black)
    - wrap around
    - copy edge
    - reflect across edge

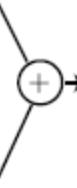


# Properties of smoothing filters

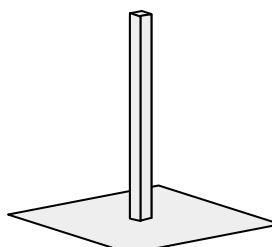
- Smoothing
  - Values positive
  - Sum to 1 → constant regions same as input (e.g. white wall)
  - Amount of smoothing proportional to mask size
  - Remove “high-frequency” components; “low-pass” filter

# Application: Hybrid Images

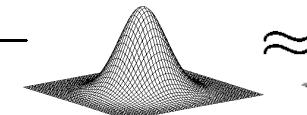
Gaussian Filter



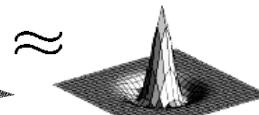
Laplacian Filter



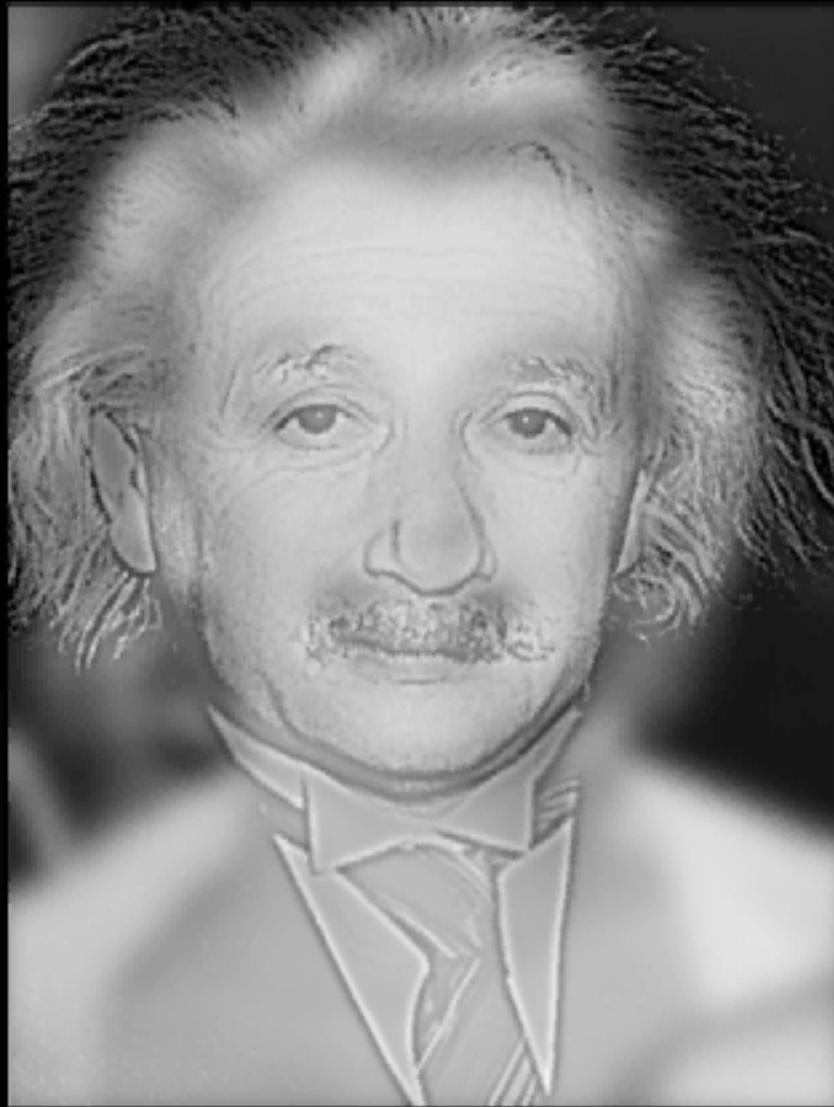
unit impulse



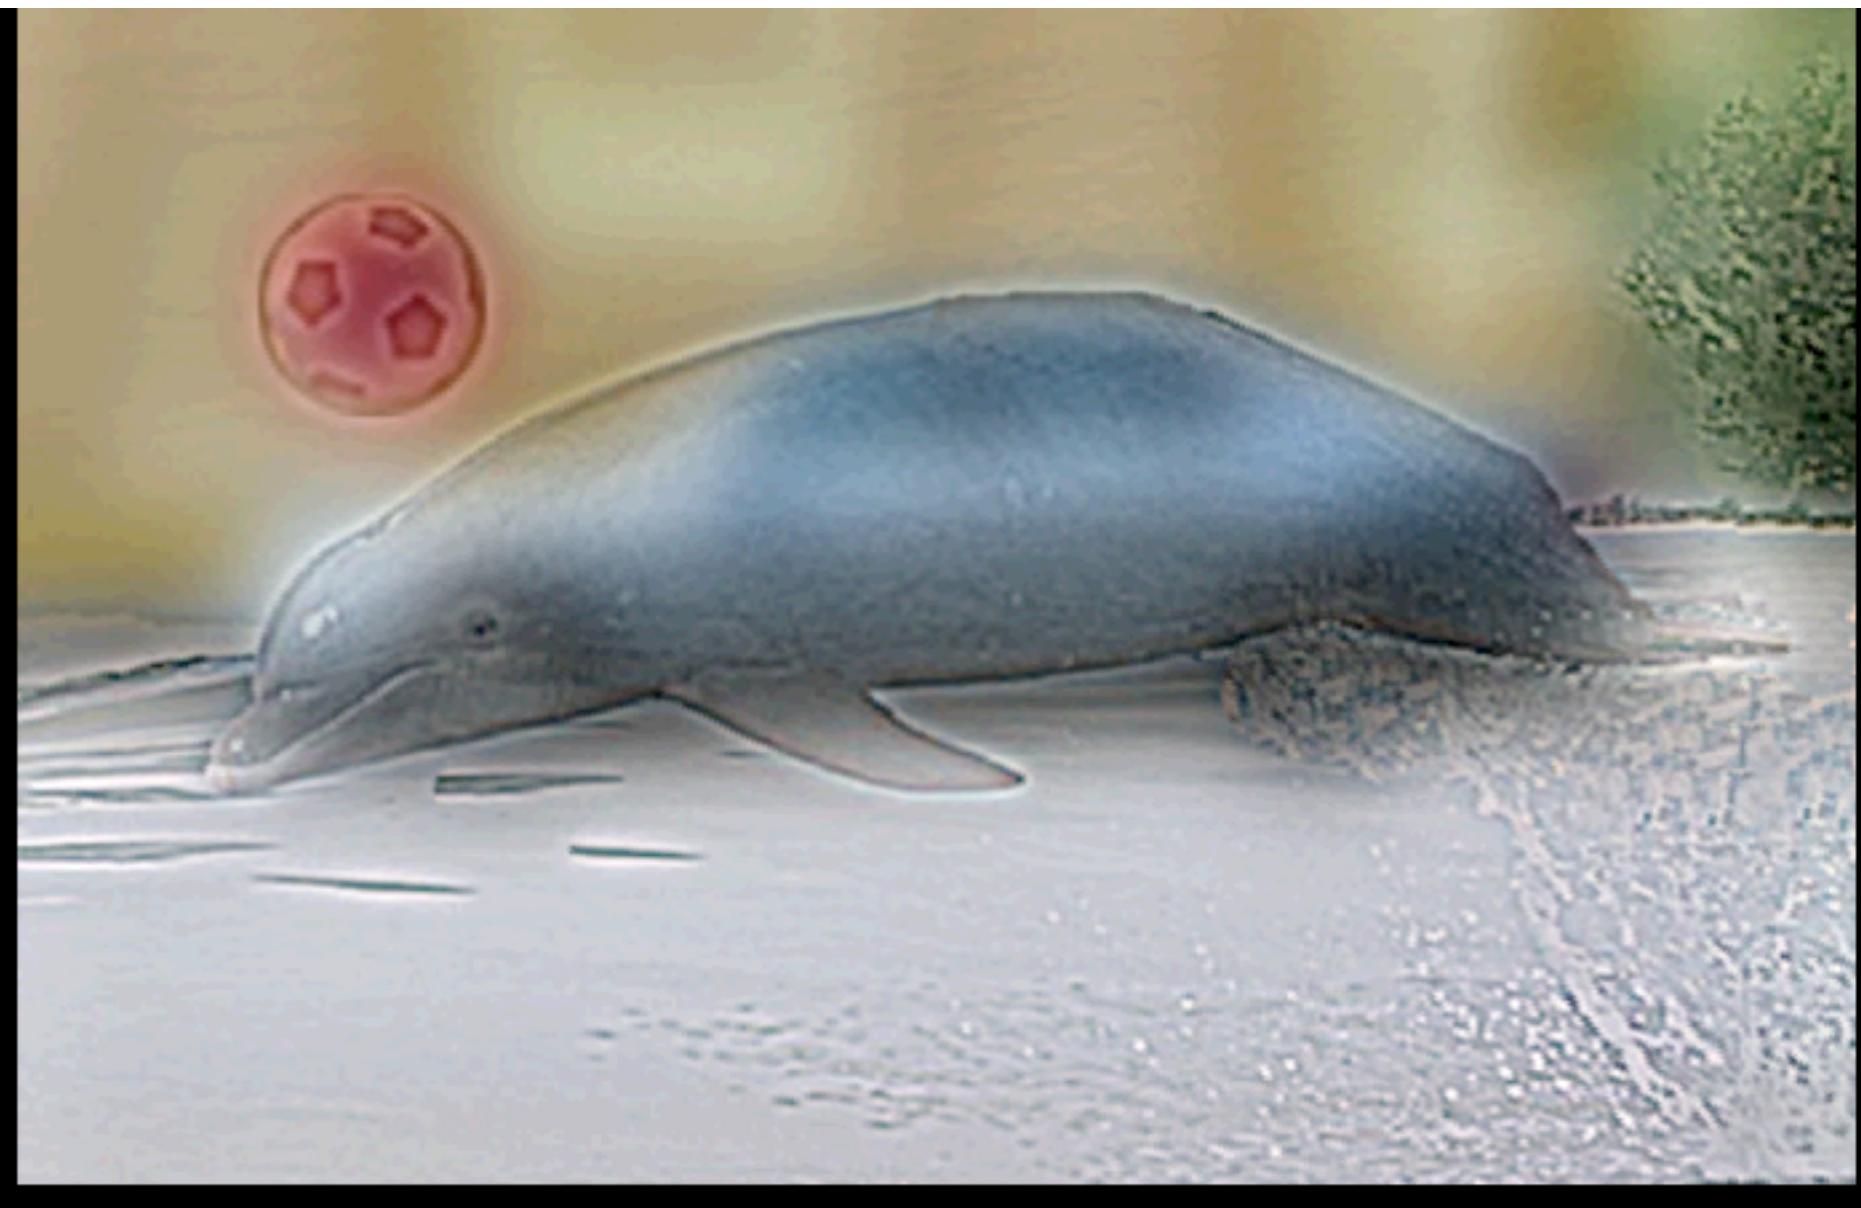
Gaussian



Laplacian of Gaussian







# Filtering examples: sharpening

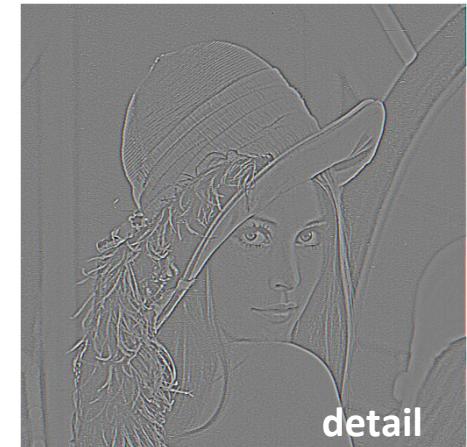
- What do we lose after smoothing an image?



-



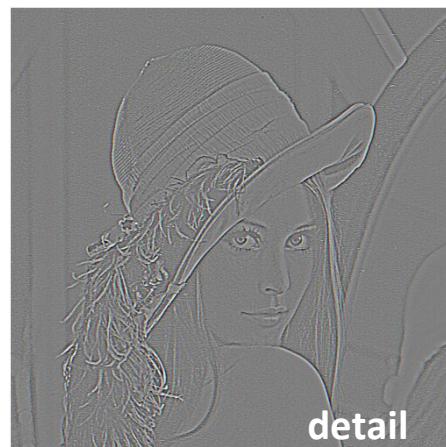
=



- Adding the detail:



+



=



# Filtering examples: sharpening

$$\left[ \begin{array}{ccc} 0 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 0 \end{array} \right] - \frac{1}{9} \left[ \begin{array}{ccc} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{array} \right]$$

(Sum = 1)



input

?

output

# Filtering examples: sharpening

$$\frac{1}{9} \begin{array}{|c|c|c|} \hline -1 & -1 & -1 \\ \hline -1 & 17 & -1 \\ \hline -1 & -1 & -1 \\ \hline \end{array}$$

(Sum = 1)

\*



input

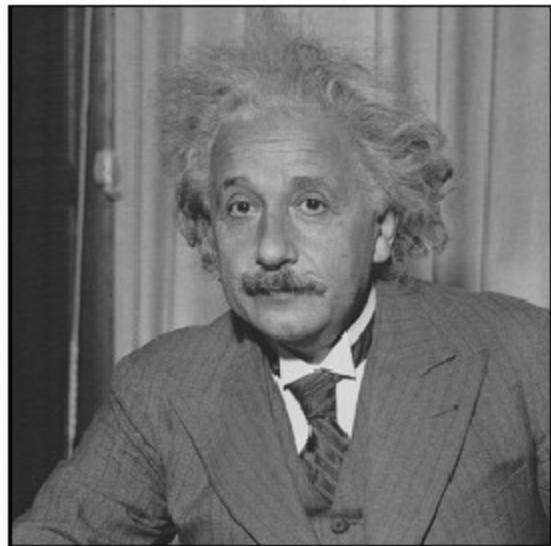
=



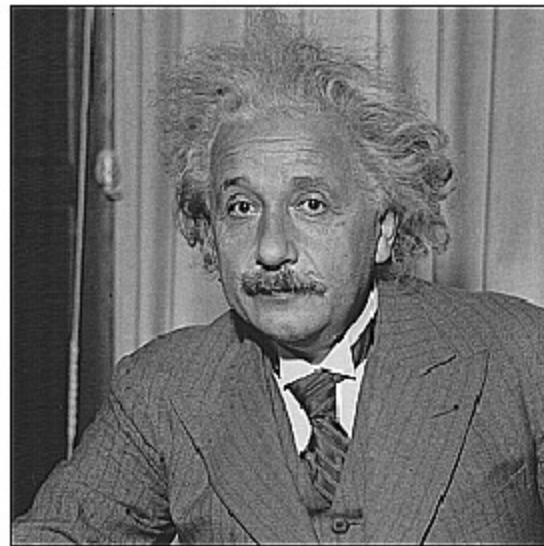
output

Sharpening filter - accentuate differences with local averages

# Filtering examples: sharpening

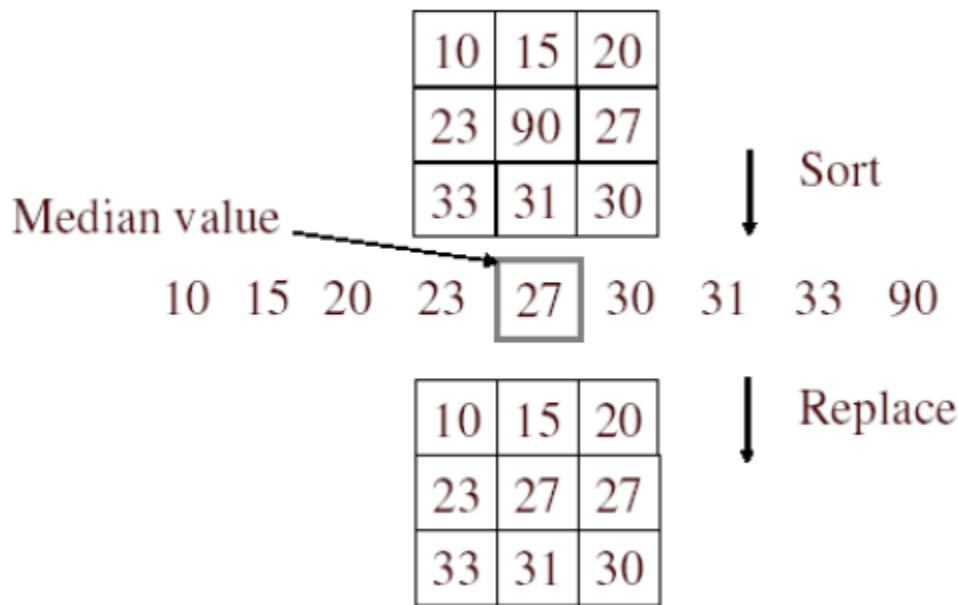


**before**



**after**

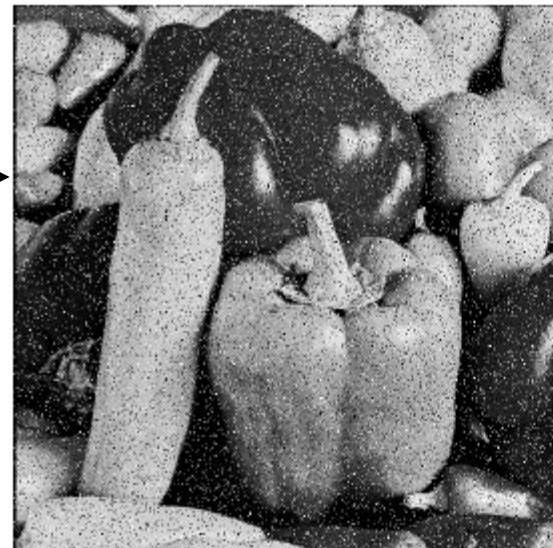
# Median filter



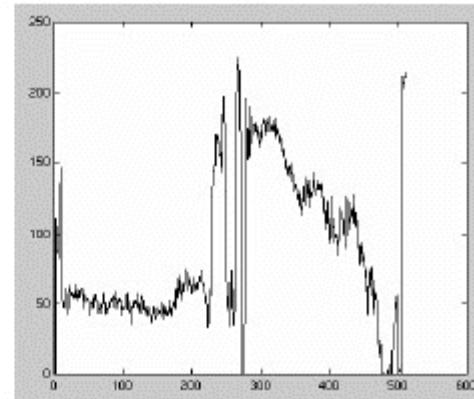
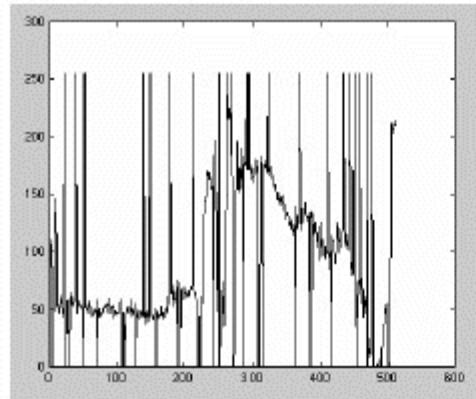
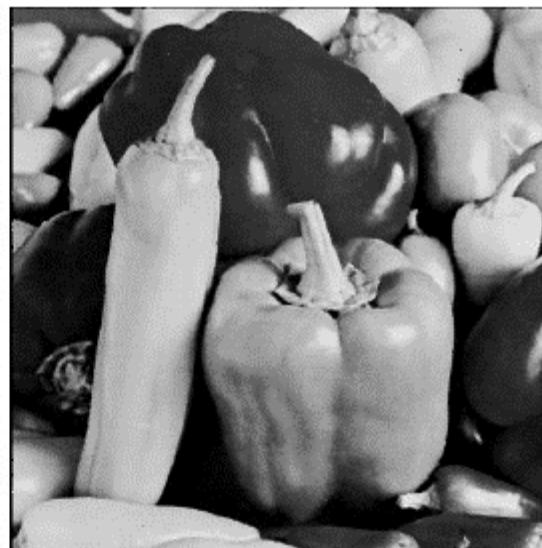
- No new pixel values introduced
- Removes spikes: good for removing image noise
- Non-linear filter

# Median filter

Salt and  
pepper  
noise



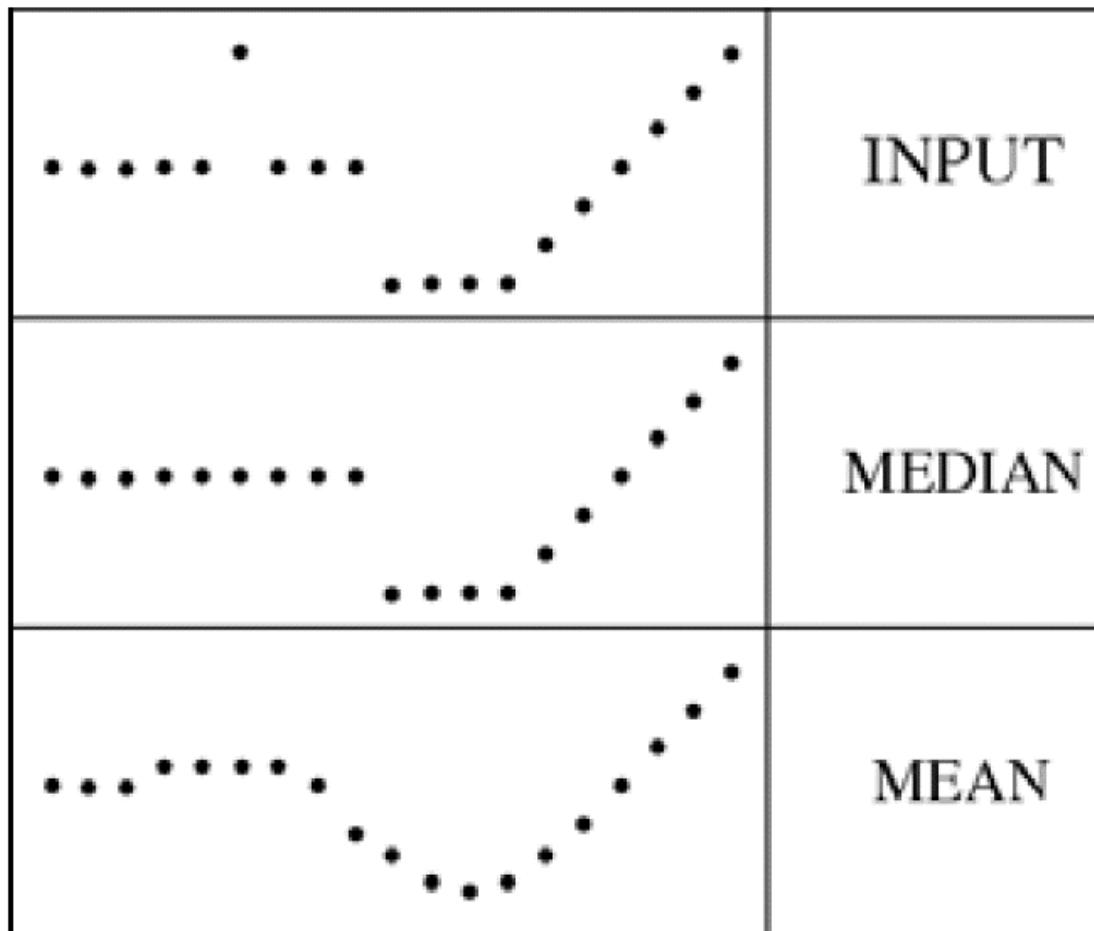
Median  
filtered



Plots of a row of the image

# Median filter

- Median filter is edge preserving



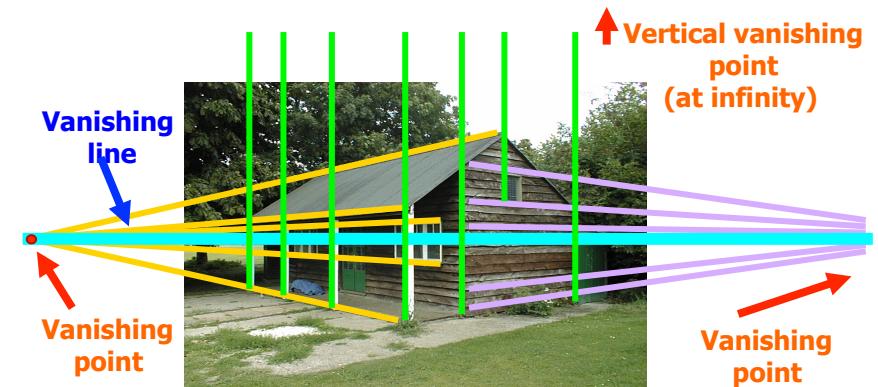
# Gradients & edges

# Why do we care about edges?

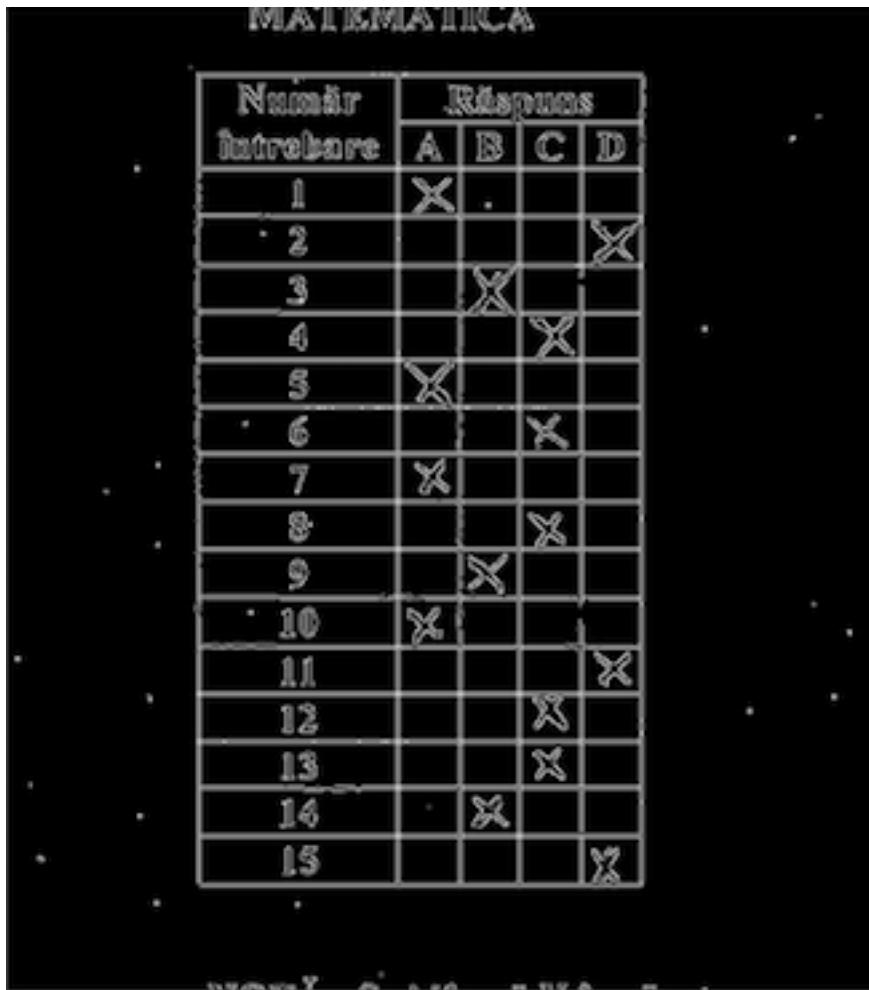
- Extract information,  
recognize objects



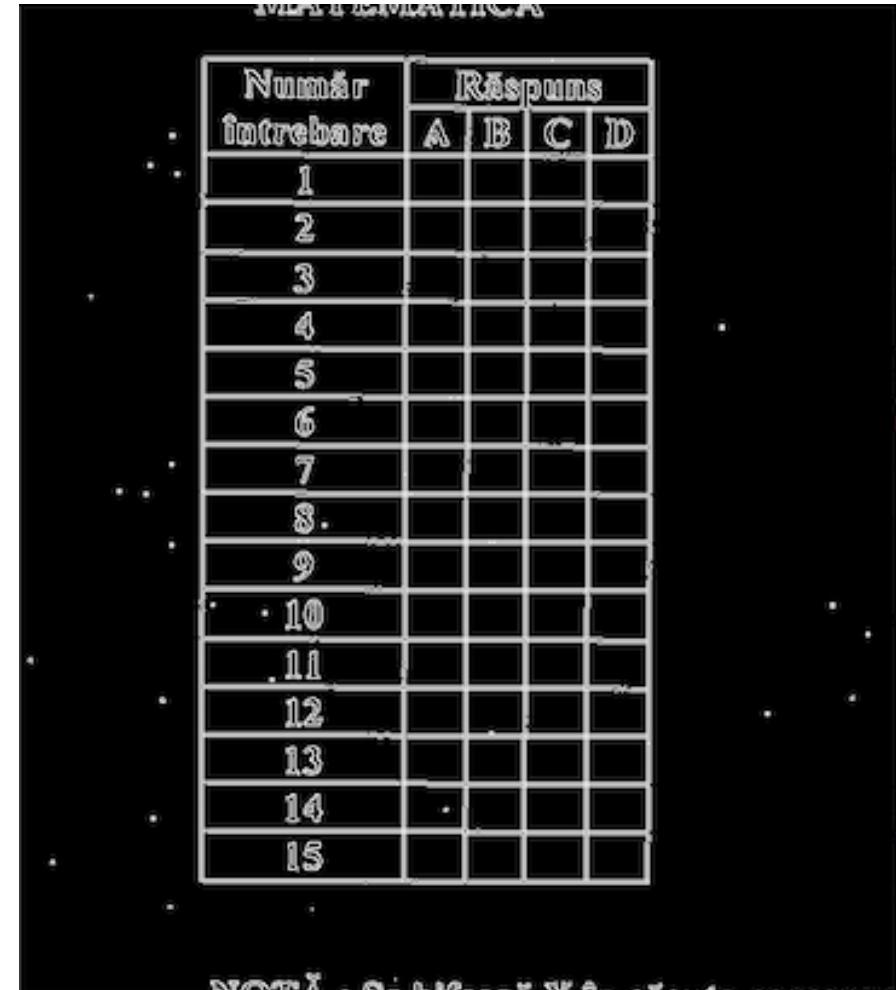
- Recover geometry and  
viewpoint



# Lab class 2 – Automatic grading of multiple choice tests



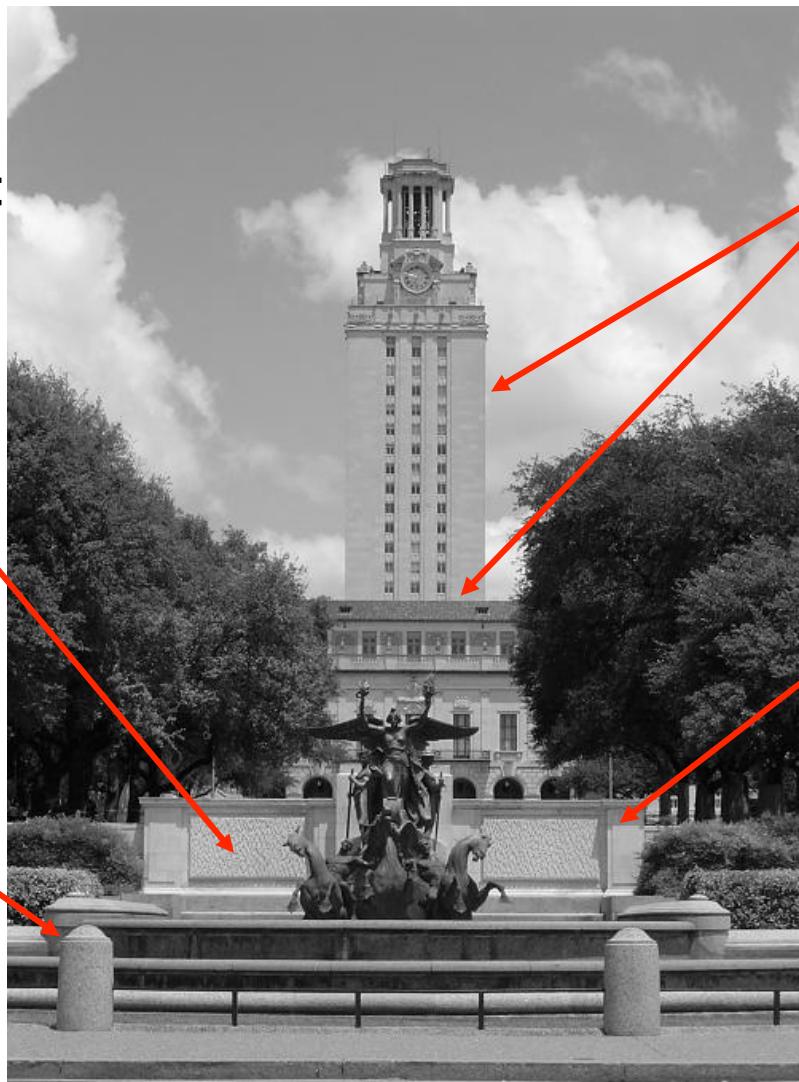
Edge image for query



Edge image for template

# What causes an edge?

Reflectance change:  
appearance  
information, texture

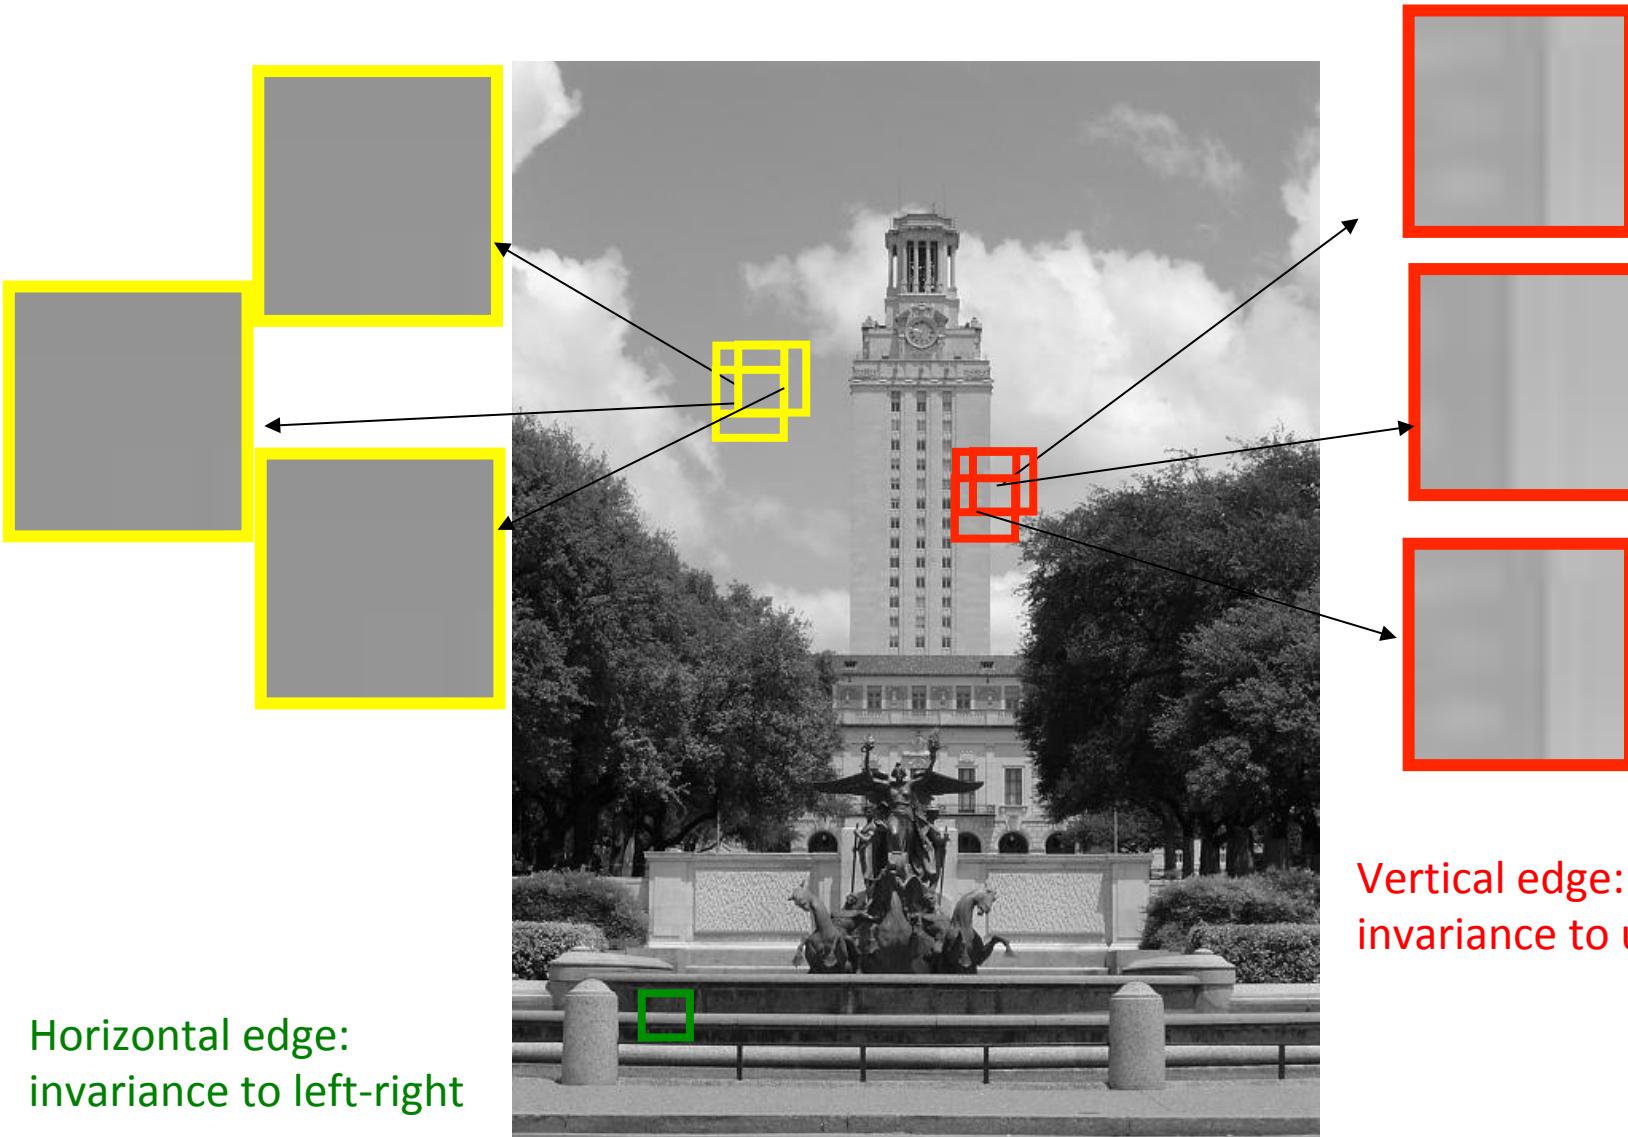


Change in surface  
orientation: shape

Depth discontinuity:  
object boundary

Cast shadows

# Edges/gradients and invariance



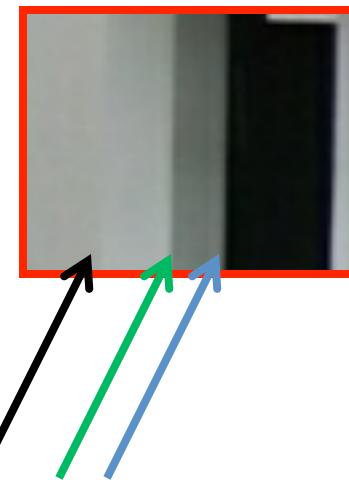
Horizontal edge:  
invariance to left-right

Vertical edge:  
invariance to up-down

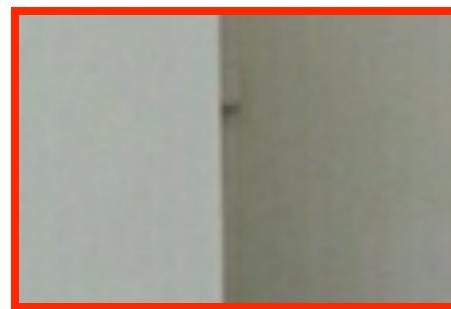
# Closeup of edges



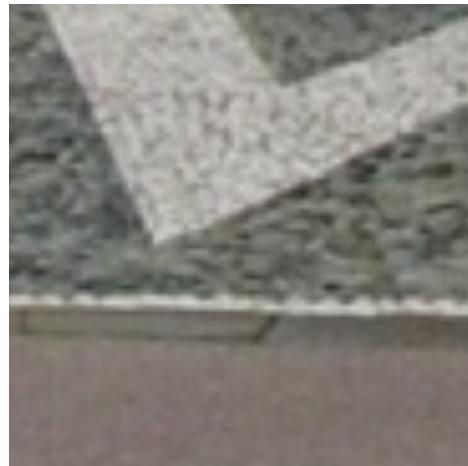
# Closeup of edges



# Closeup of edges

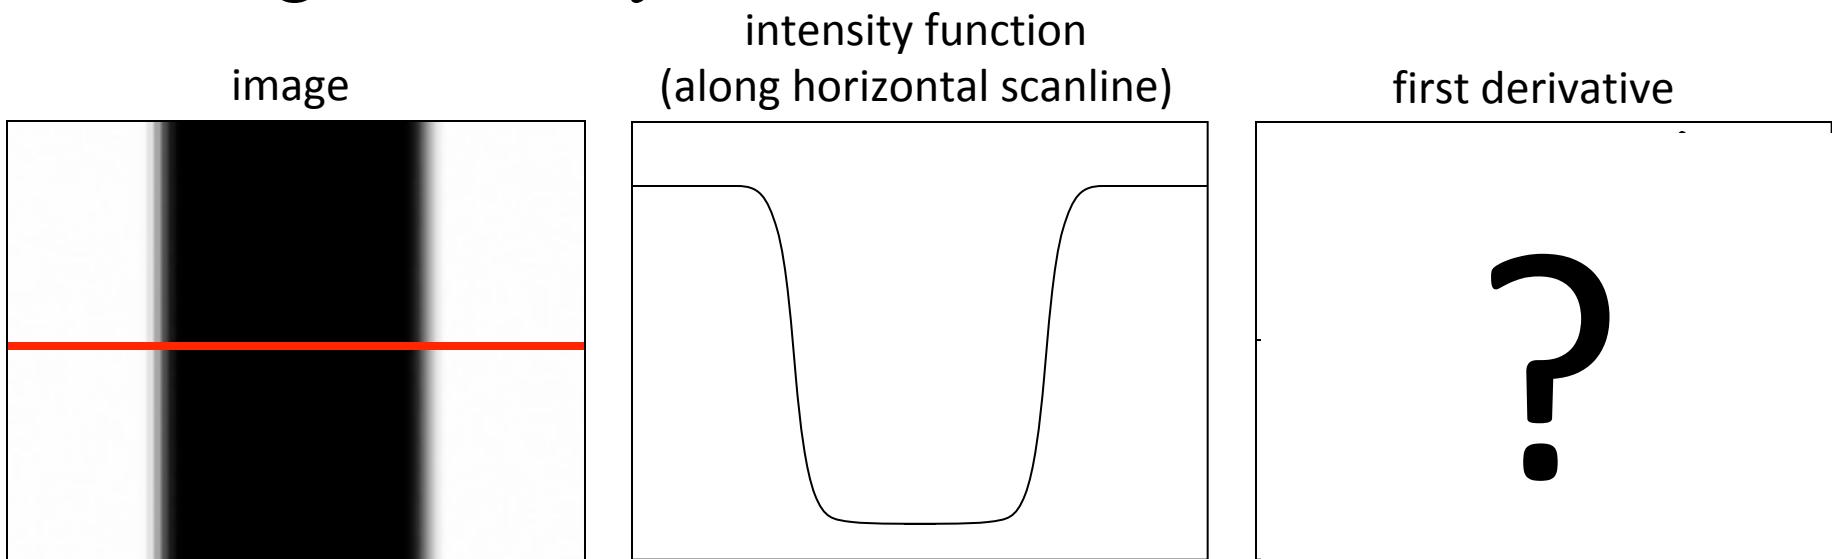


# Closeup of edges



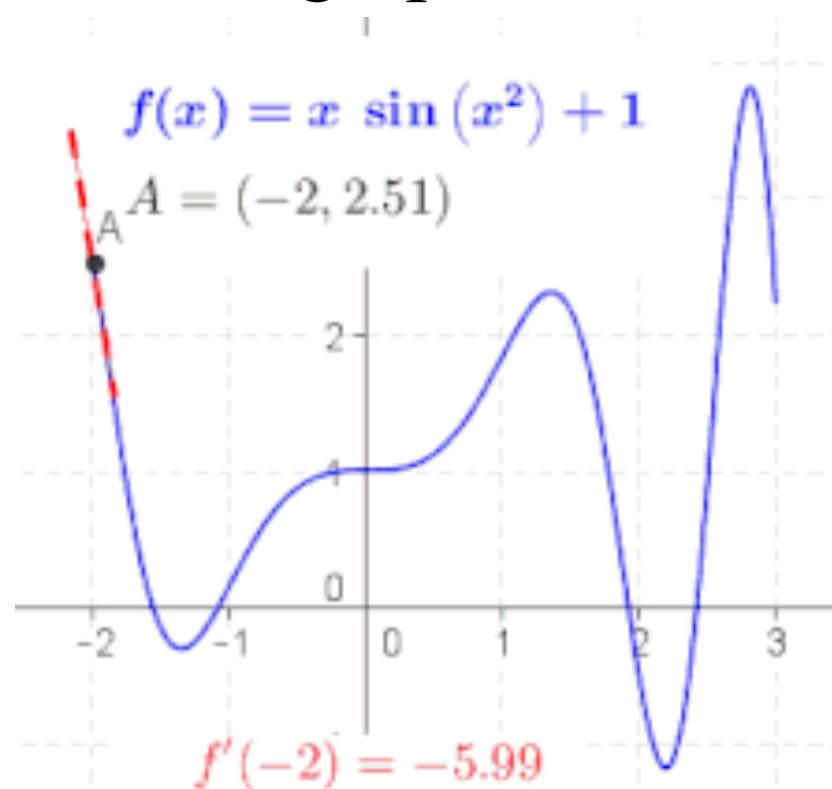
# Characterizing edges

- An edge is a place of rapid change in the image intensity function



# Derivative of a function

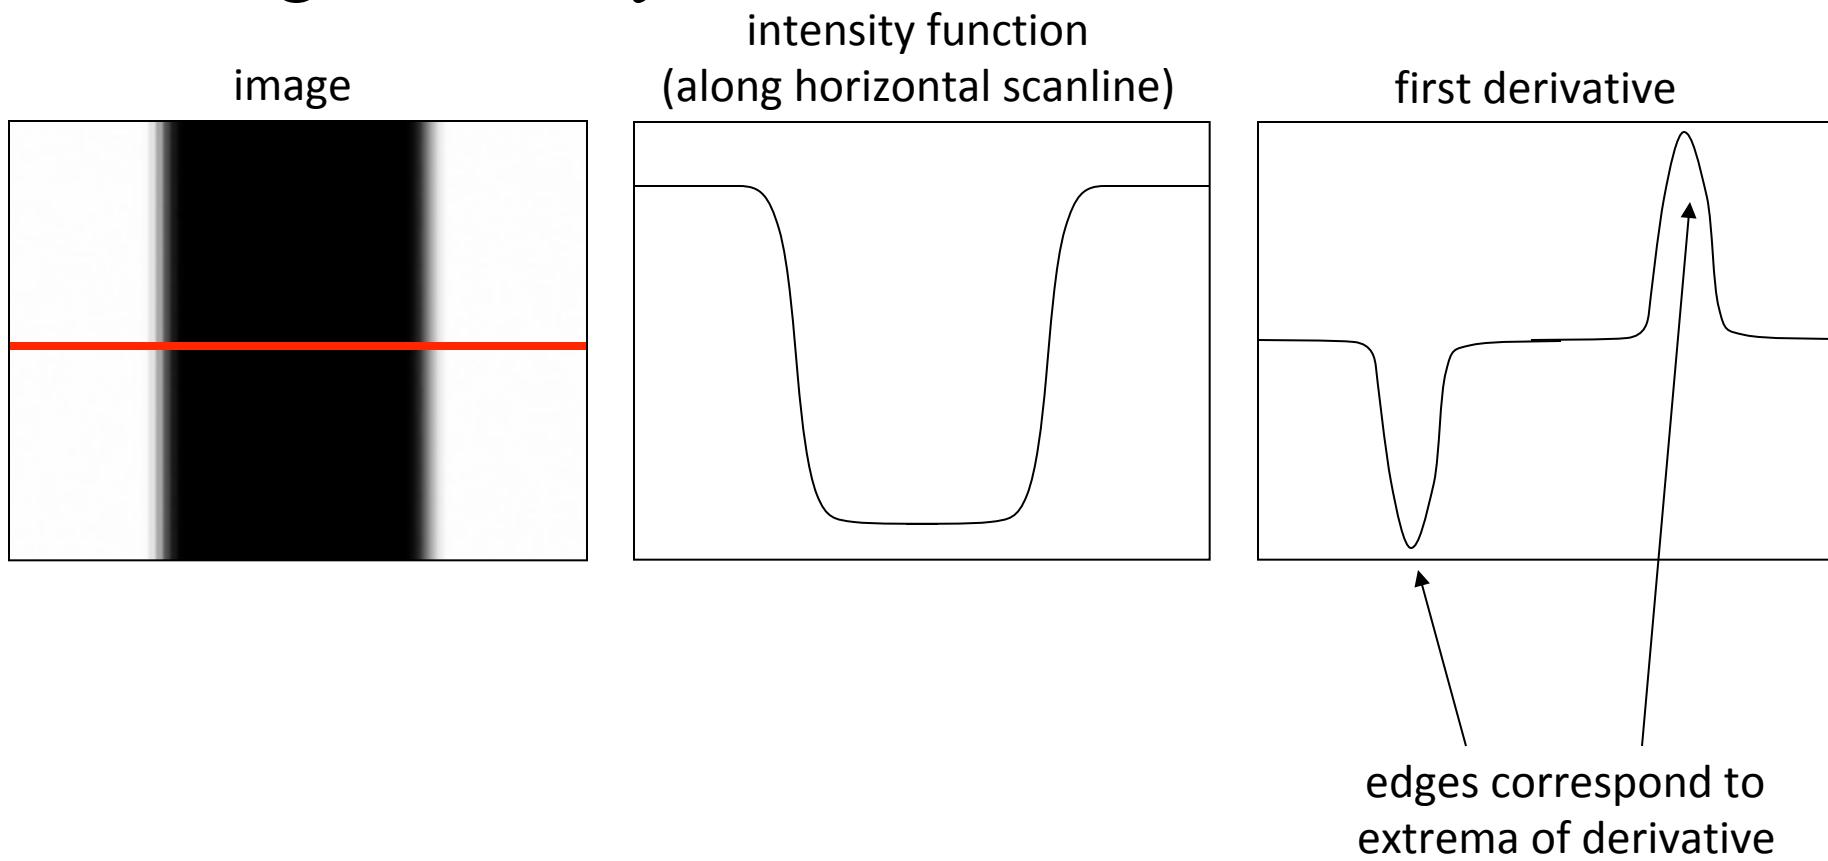
- slope of the tangent line to the graph of the function at that point
- **green**  $> 0$
- **black**  $= 0$
- **red**  $< 0$



<https://en.wikipedia.org/wiki/Derivative>

# Characterizing edges

- An edge is a place of rapid change in the image intensity function



# Derivatives with convolution

For 2D function,  $f(x,y)$ , the partial derivative is:

$$\frac{\partial f(x, y)}{\partial x} = \lim_{\varepsilon \rightarrow 0} \frac{f(x + \varepsilon, y) - f(x, y)}{\varepsilon}$$

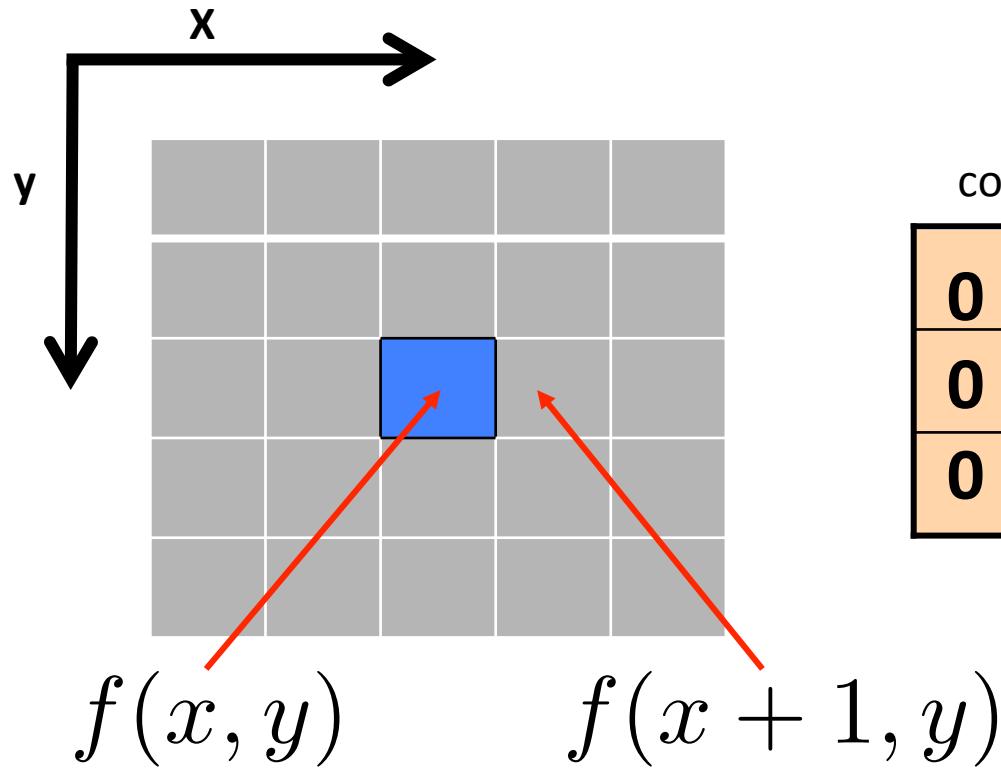
For discrete data, we can approximate using finite differences:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$

To implement above as convolution, what would be the associated filter?

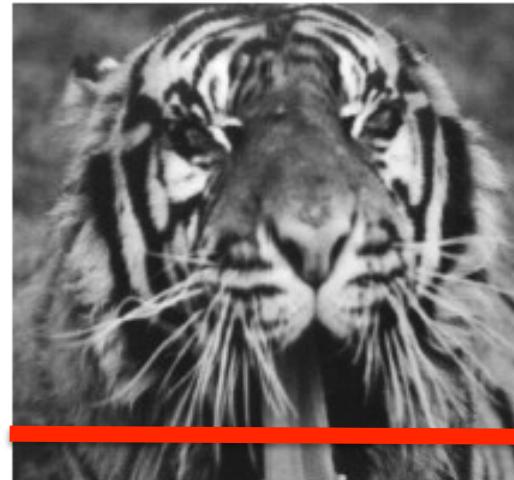
# Derivatives with convolution

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x + 1, y) - f(x, y)}{1}$$



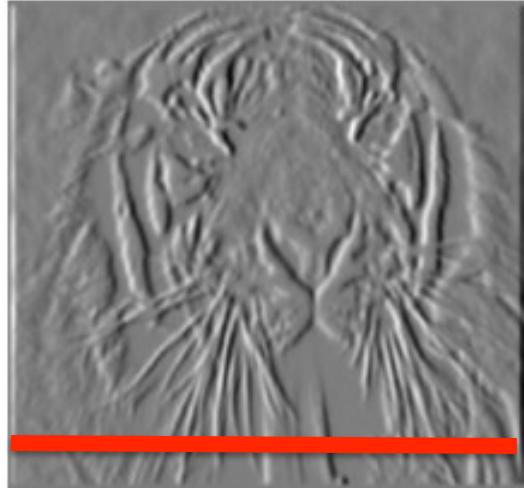
correlation			convolution		
0	0	0	1	-1	
0	0	1	1	-1	
0	0	0	0	0	

# Partial derivatives of an image



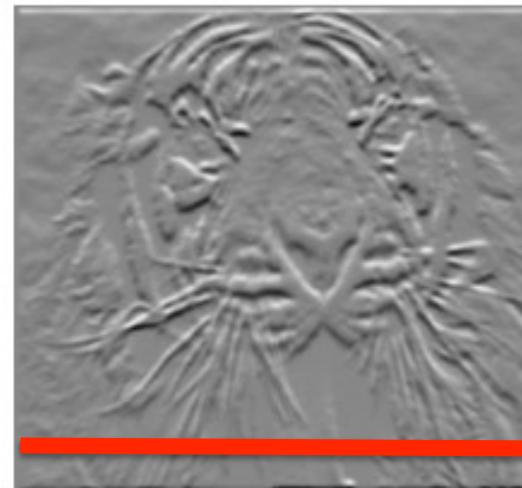
$$\frac{\partial f(x, y)}{\partial x}$$

-1	1
----	---



$$\frac{\partial f(x, y)}{\partial y}$$

-1
1



Which shows changes with respect to x?

(showing filters for correlation)

Slide credit: Kristen Grauman

# Sobel filters

There exists other approximations for computing partial derivatives:

- Sobel filters: compute derivatives considering larger neighborhoods (3 x 3)

-1	0	1
-2	0	2
-1	0	1

Vertical Sobel filter  
for computing the partial derivative

$$\frac{\partial f(x, y)}{\partial x}$$

1	2	1
0	0	0
-1	-2	-1

Horizontal Sobel filter  
for computing the partial derivative

$$\frac{\partial f(x, y)}{\partial y}$$

# Other filters for finite differences

Prewitt:  $M_x = \begin{bmatrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$

Sobel:  $M_x = \boxed{\begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}}$ ;  $M_y = \begin{bmatrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{bmatrix}$

Roberts:  $M_x = \begin{bmatrix} 0 & 1 \\ -1 & 0 \end{bmatrix}$  ;  $M_y = \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix}$

```
img = cv2.imread('simona.jpg',0)
sobelx = cv2.Sobel(img,cv2.CV_64F,1,0,ksize=3)
plt.imshow(sobelx,cmap = 'gray')
plt.show()
```

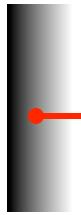


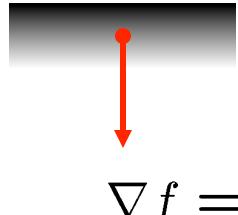
# Image gradient

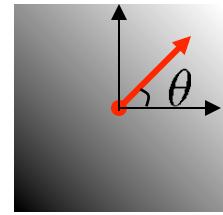
The gradient of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity


$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The **gradient direction** (orientation of edge normal) is given by:

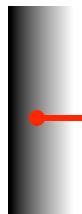
$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

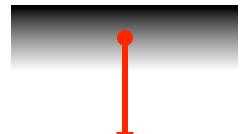
# Image gradient

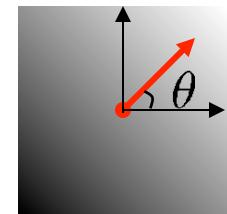
The gradient of an image:

$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity


$$\nabla f = \left[ \frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[ 0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[ \frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

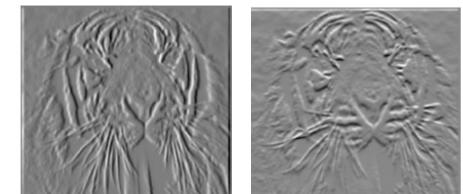
The **gradient direction** (orientation of edge normal) is given by:

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The **edge strength** is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

or  $\|\nabla f\| = \left| \frac{\partial f}{\partial x} \right| + \left| \frac{\partial f}{\partial y} \right|$



Slide credit Steve Seitz

# Computing the image gradient

1. compute the partial derivatives:

$$\frac{\partial f(x, y)}{\partial x} \quad \frac{\partial f(x, y)}{\partial y}$$

2. compute the gradient magnitude:

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2} \quad \text{or} \quad \|\nabla f\| = \left|\frac{\partial f}{\partial x}\right| + \left|\frac{\partial f}{\partial y}\right|$$

# Computing the image gradient

1. compute the partial derivative  $\frac{\partial f(x, y)}{\partial x}$

Use the vertical  
Sobel filter

$$M_x =$$

-1	0	1
-2	0	2
-1	0	1

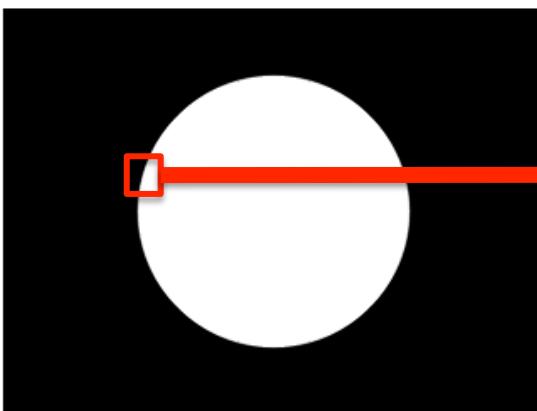
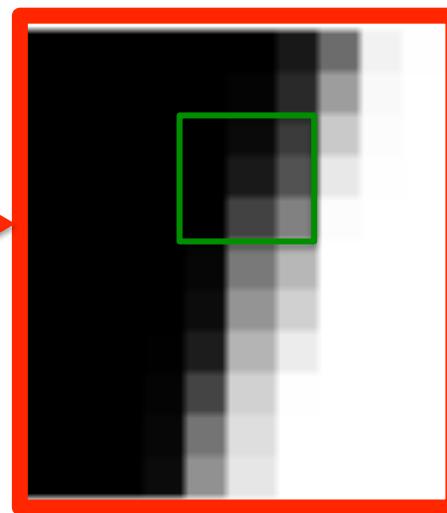


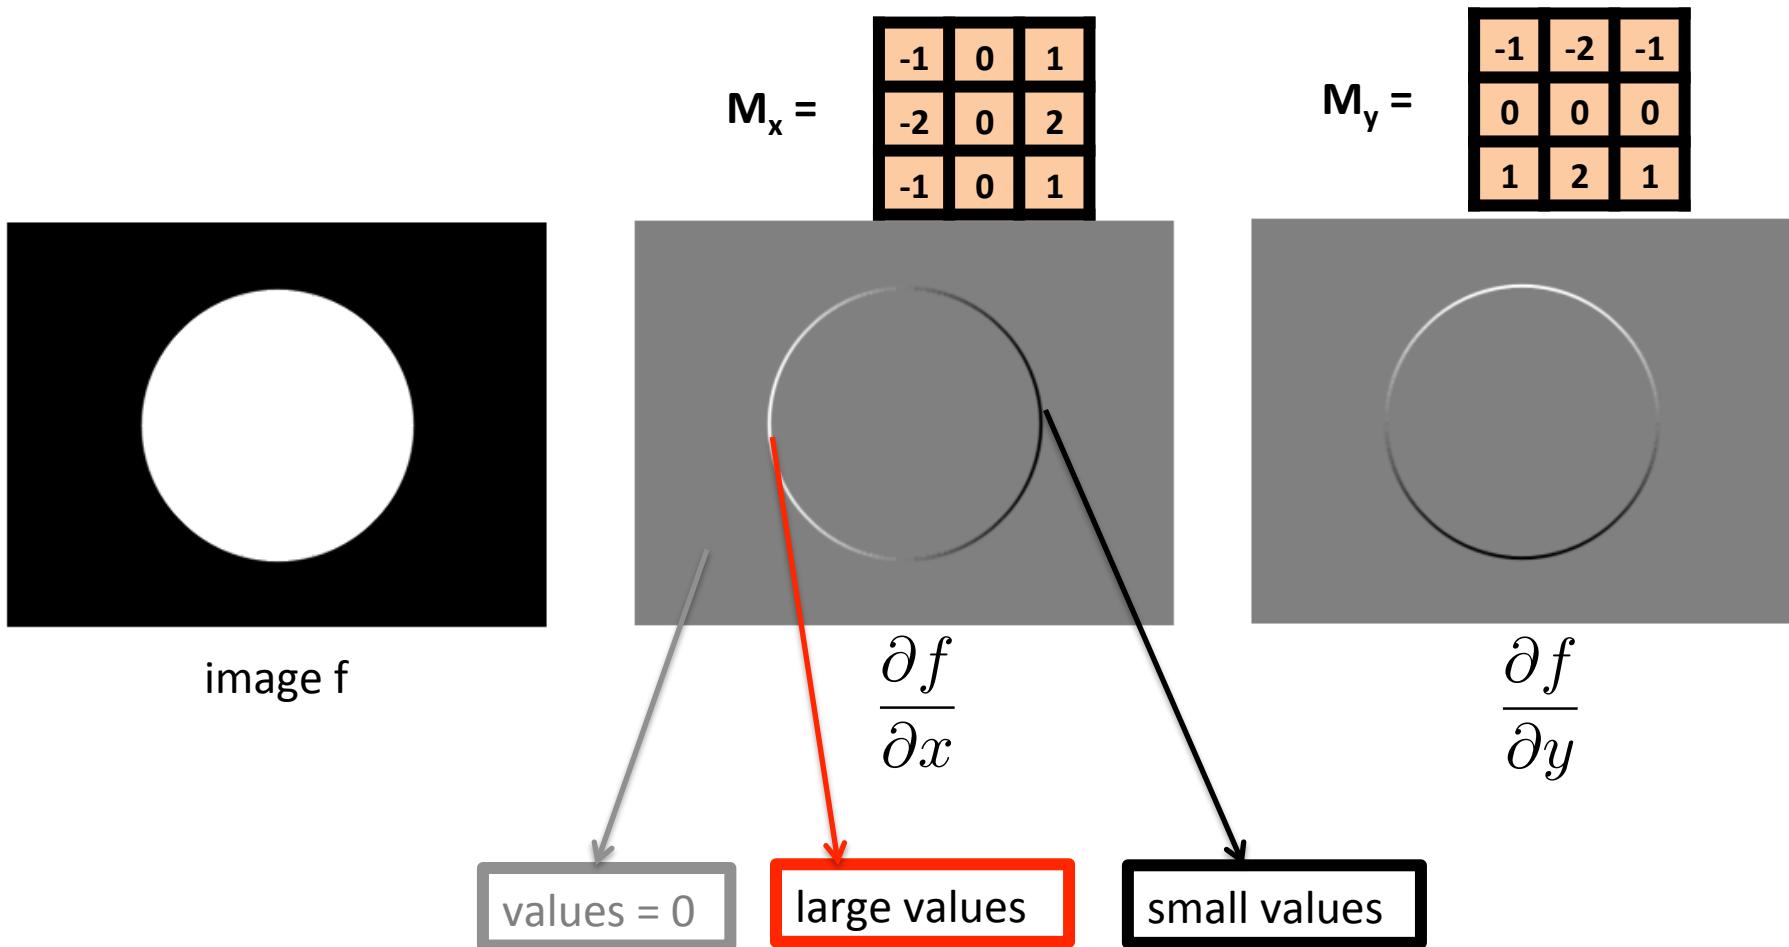
image f



0	0	0	0	4	41	157	249	255
0	0	0	0	10	59	201	252	255
0	0	0	0	25	82	232	254	255
0	0	0	0	66	129	252	255	255
0	0	0	6	121	183	255	255	255
0	0	0	1	148	208	255	255	255
0	0	1	2	180	236	255	255	255
0	0	4	6	209	254	255	255	255
0	0	7	11	222	255	255	255	255

$$\begin{aligned} &= (-1) * 0 + 0 * 0 + 1 * 25 + \\ &+ (-2) * 0 + 0 * 0 + 2 * 66 + \\ &+ (-1) * 0 + 0 * 6 + 1 * 121 \\ &= 278 \end{aligned}$$

# Computing the image gradient



# Computing the image gradient

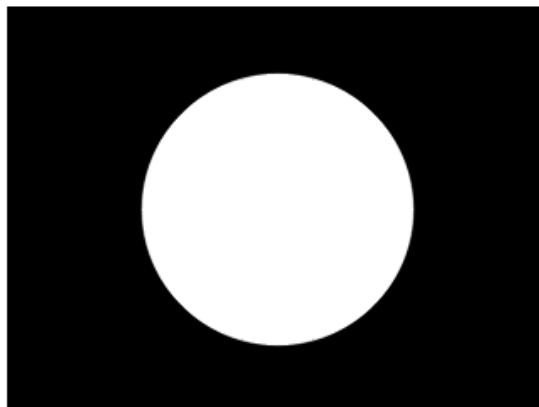
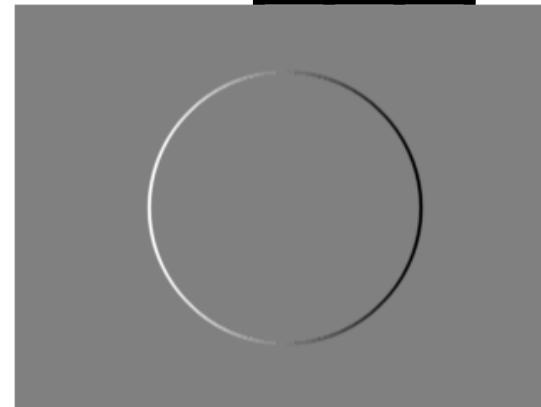


image  $f$

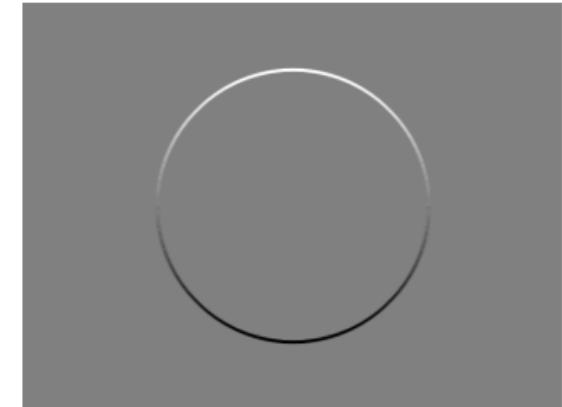


$$M_x =$$

-1	0	1
-2	0	2
-1	0	1

$$M_y =$$

-1	-2	-1
0	0	0
1	2	1

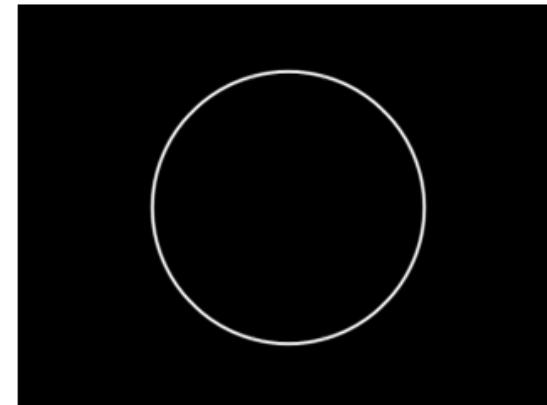


$$\frac{\partial f}{\partial x}$$

$$\frac{\partial f}{\partial y}$$

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Gradient magnitude



# Computing the image gradient

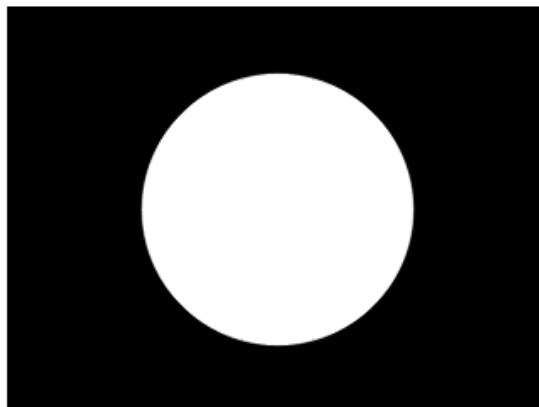
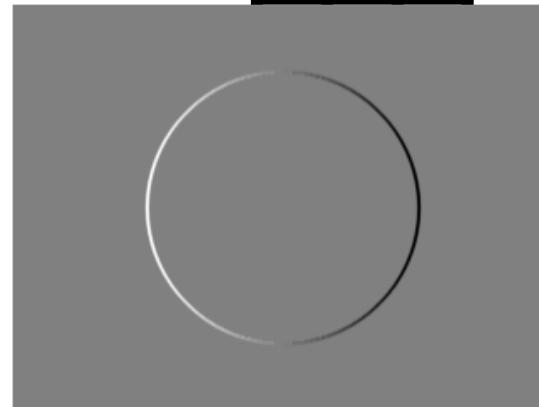
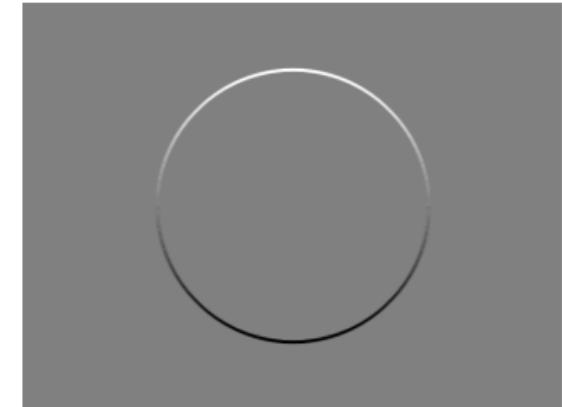


image  $f$



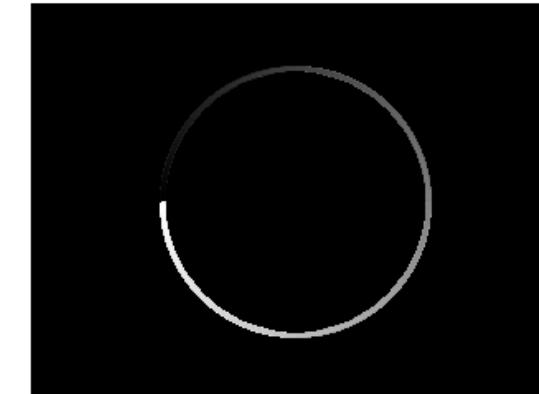
$$M_x = \begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$$

$$M_y = \begin{matrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{matrix}$$

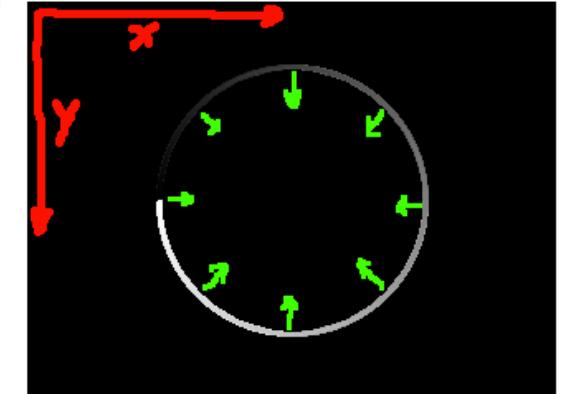


$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

Gradient orientation



$$\frac{\partial f}{\partial x} \quad \frac{\partial f}{\partial y}$$



# Computing the image gradient

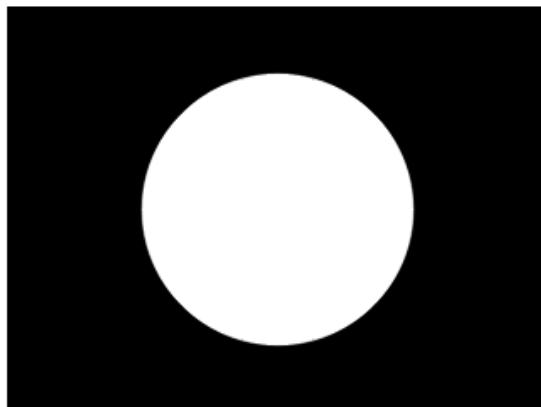
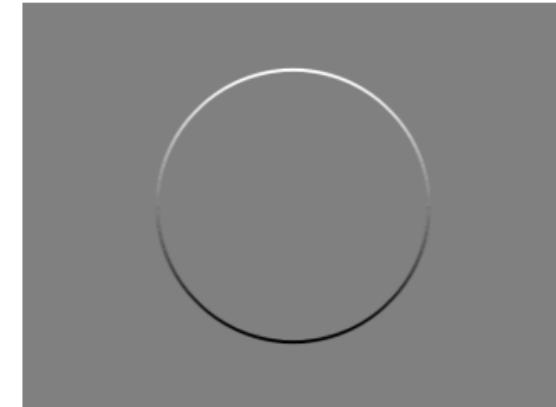


image  $f$



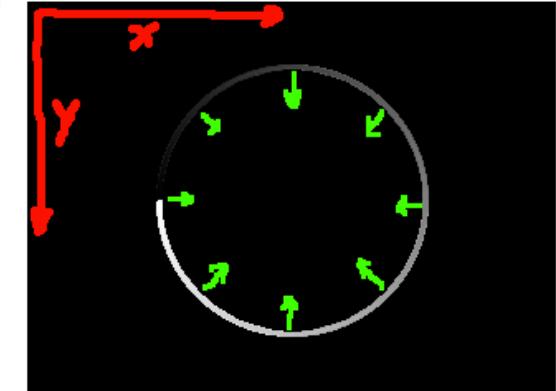
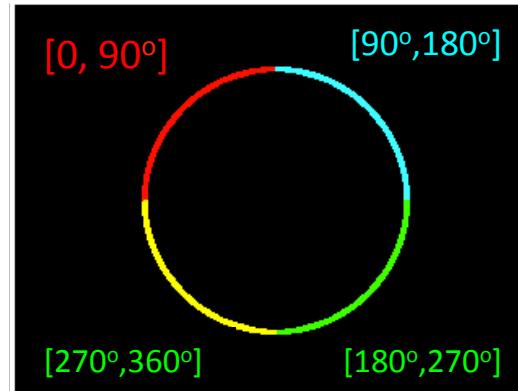
$$\frac{\partial f}{\partial x}$$



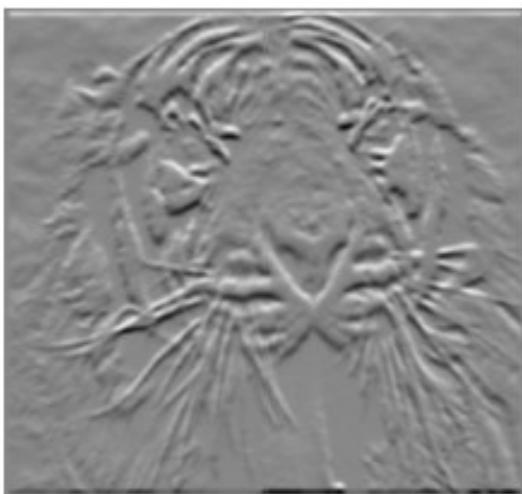
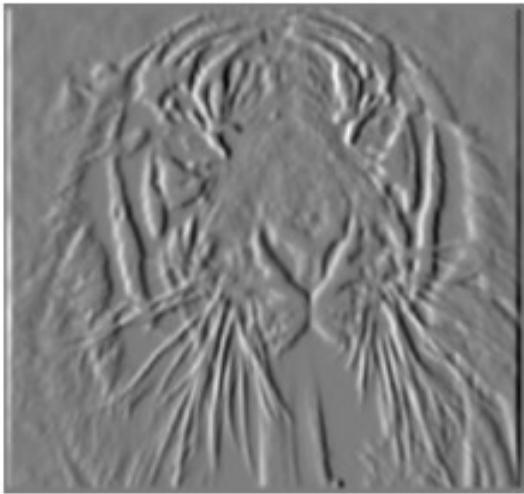
$$\frac{\partial f}{\partial y}$$

$$\theta = \tan^{-1} \left( \frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

[0, 90°]  
[90°, 180°]  
[180°, 270°]  
[270°, 360°]



# Computing the image gradient



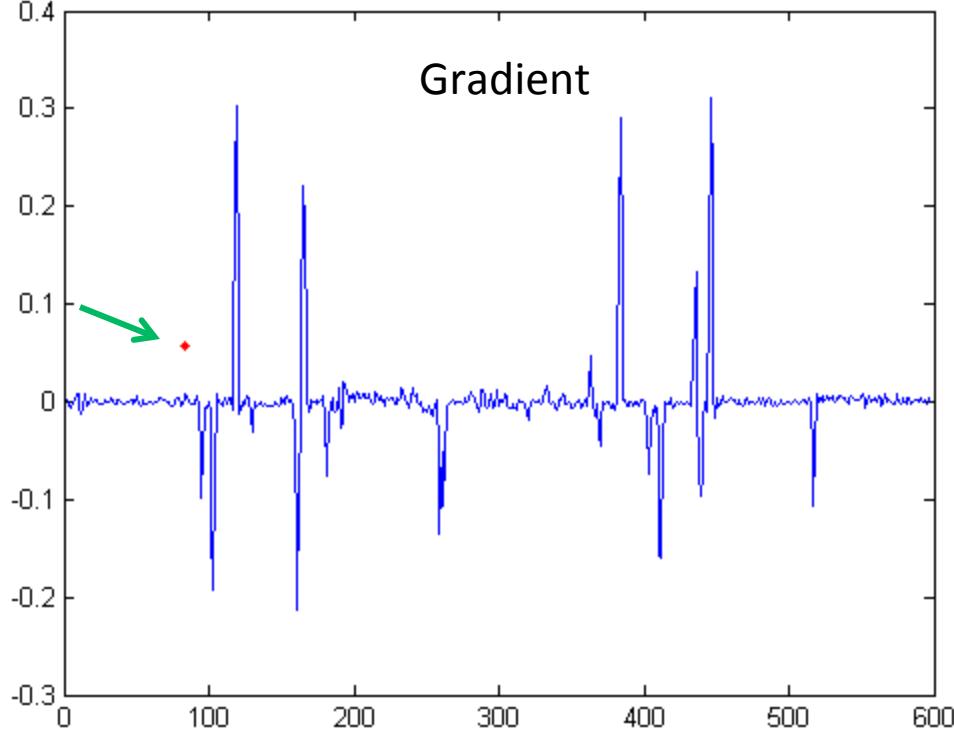
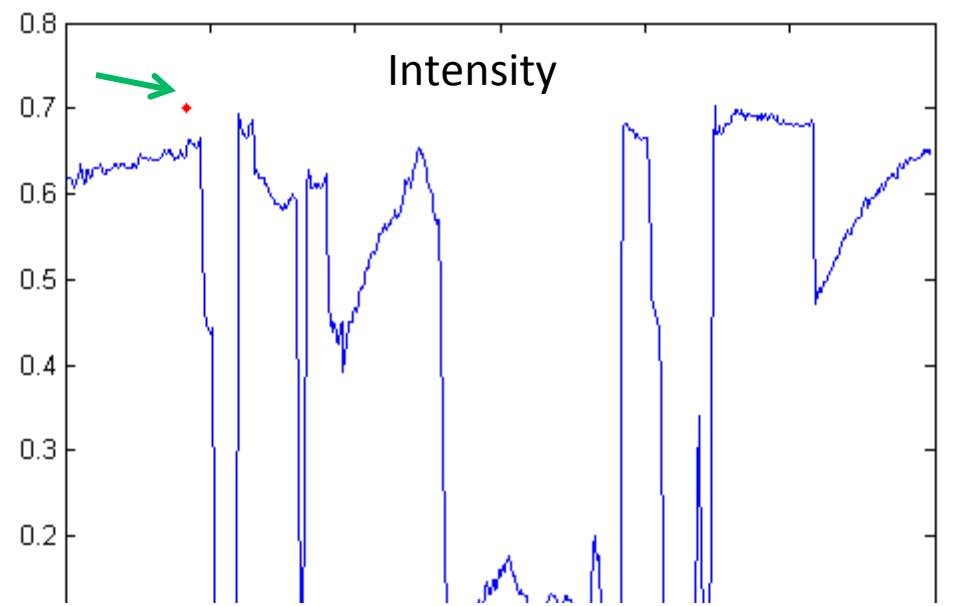
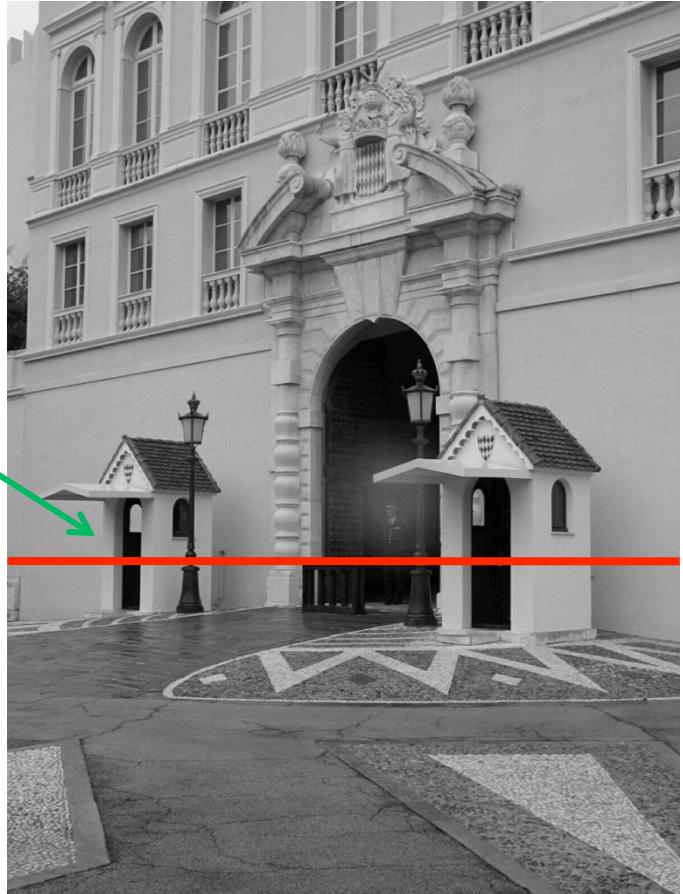
$$\frac{\partial f(x, y)}{\partial x}$$

$$\frac{\partial f(x, y)}{\partial y}$$

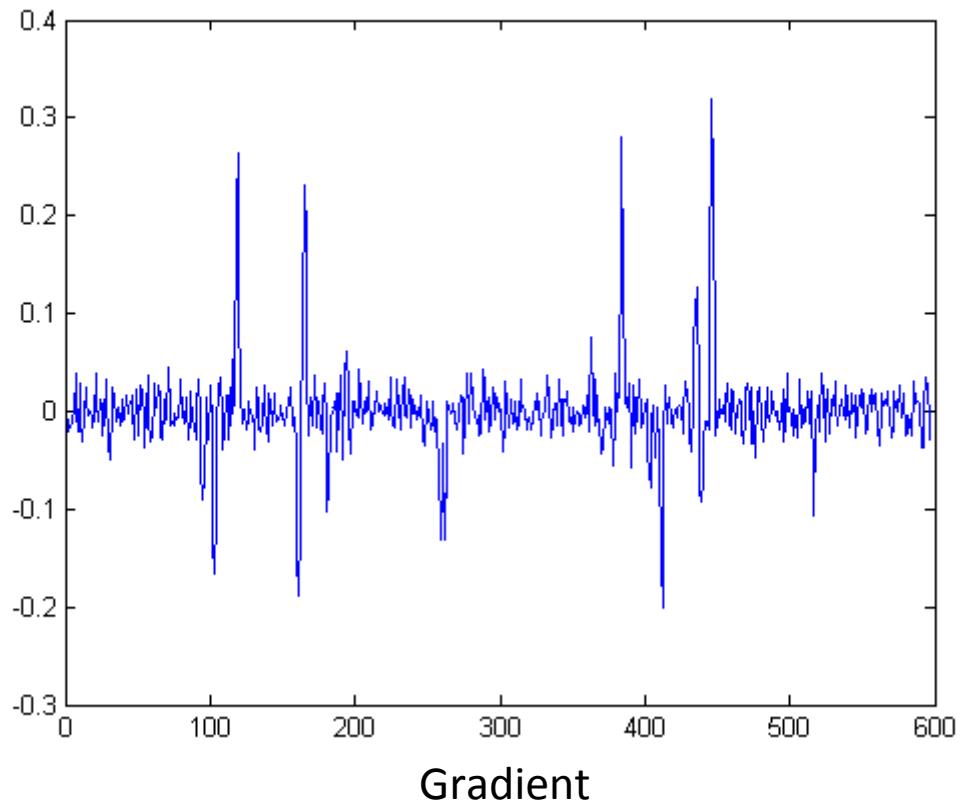
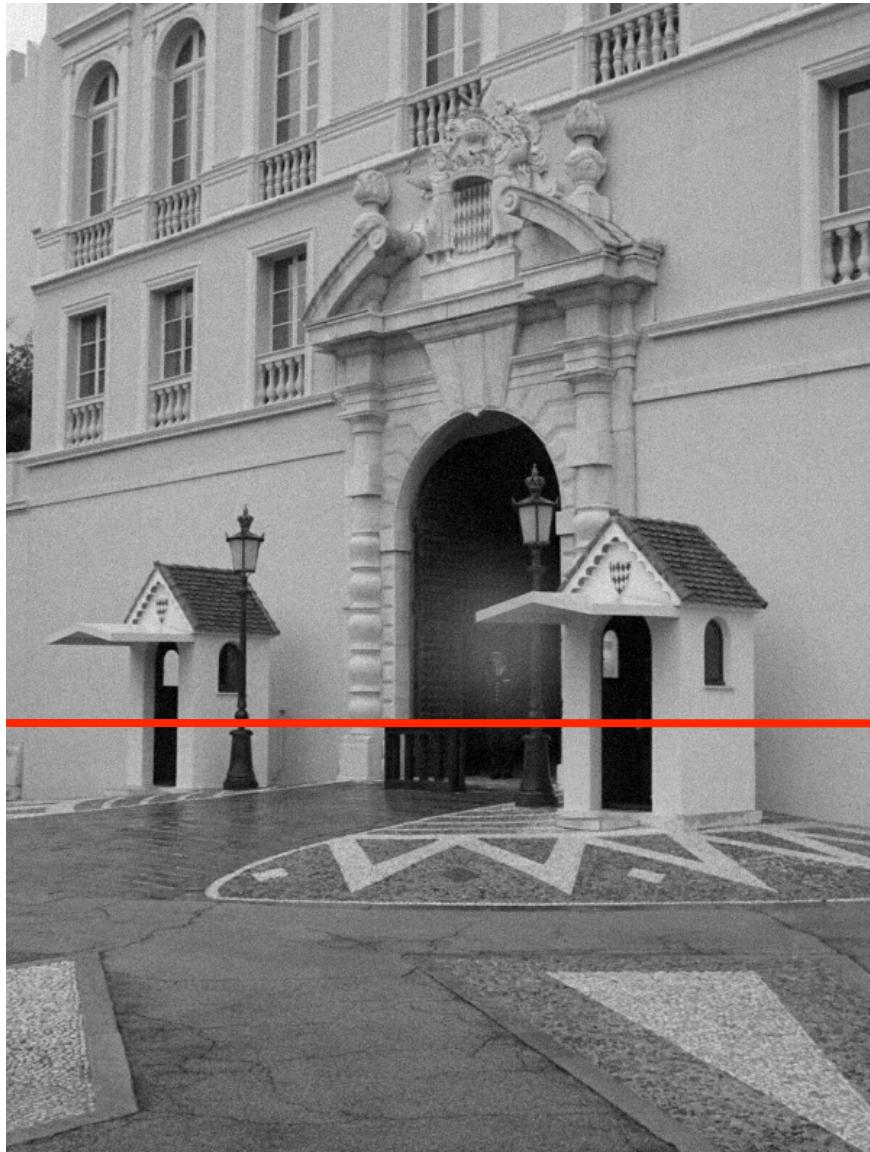
$$\sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

Gradient magnitude

# Intensity profile

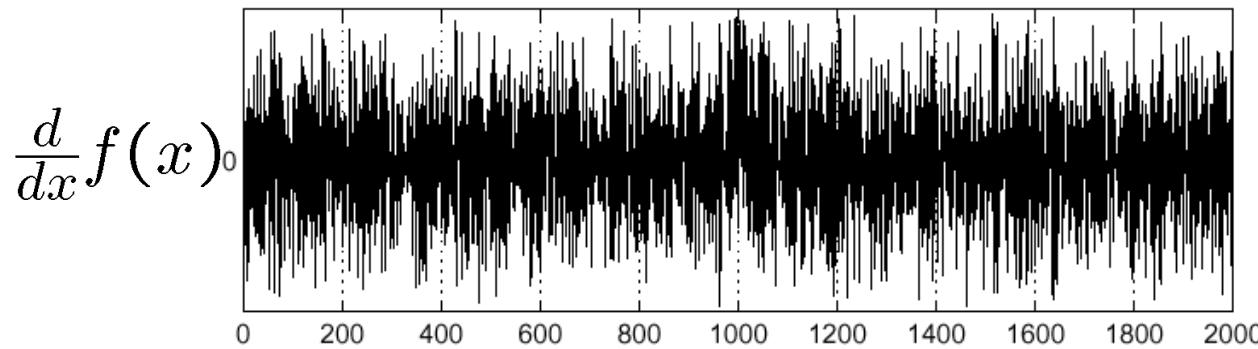
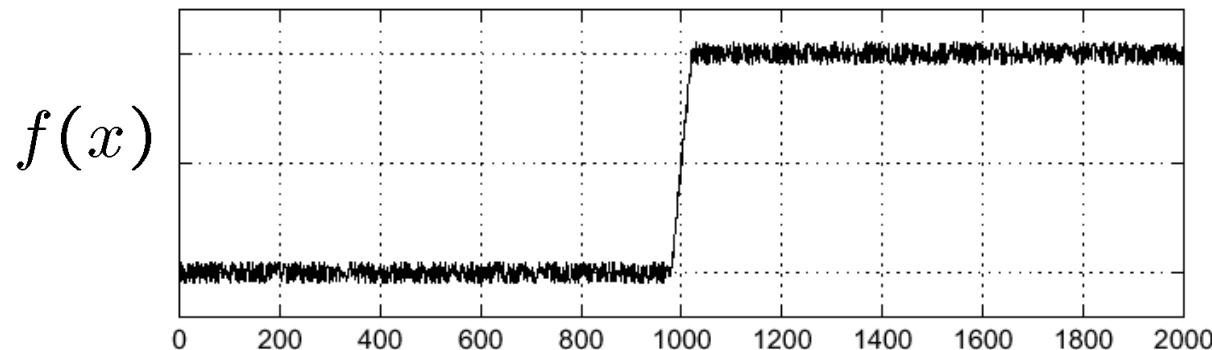


# With a little Gaussian noise



# Effects of noise

- Consider a single row or column of the image
  - Plotting intensity as a function of position gives a signal

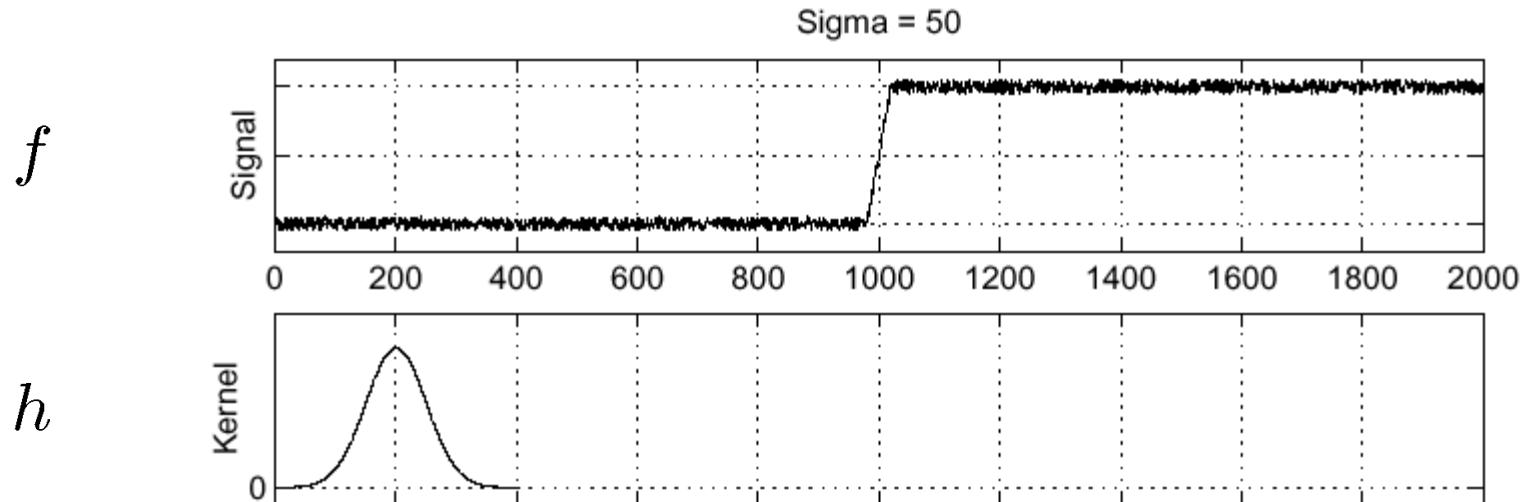


Where is the edge?

# Effects of noise

- Difference filters respond strongly to noise
  - Image noise results in pixels that look very different from their neighbors
  - Generally, the larger the noise the stronger the response
- What can we do about it?

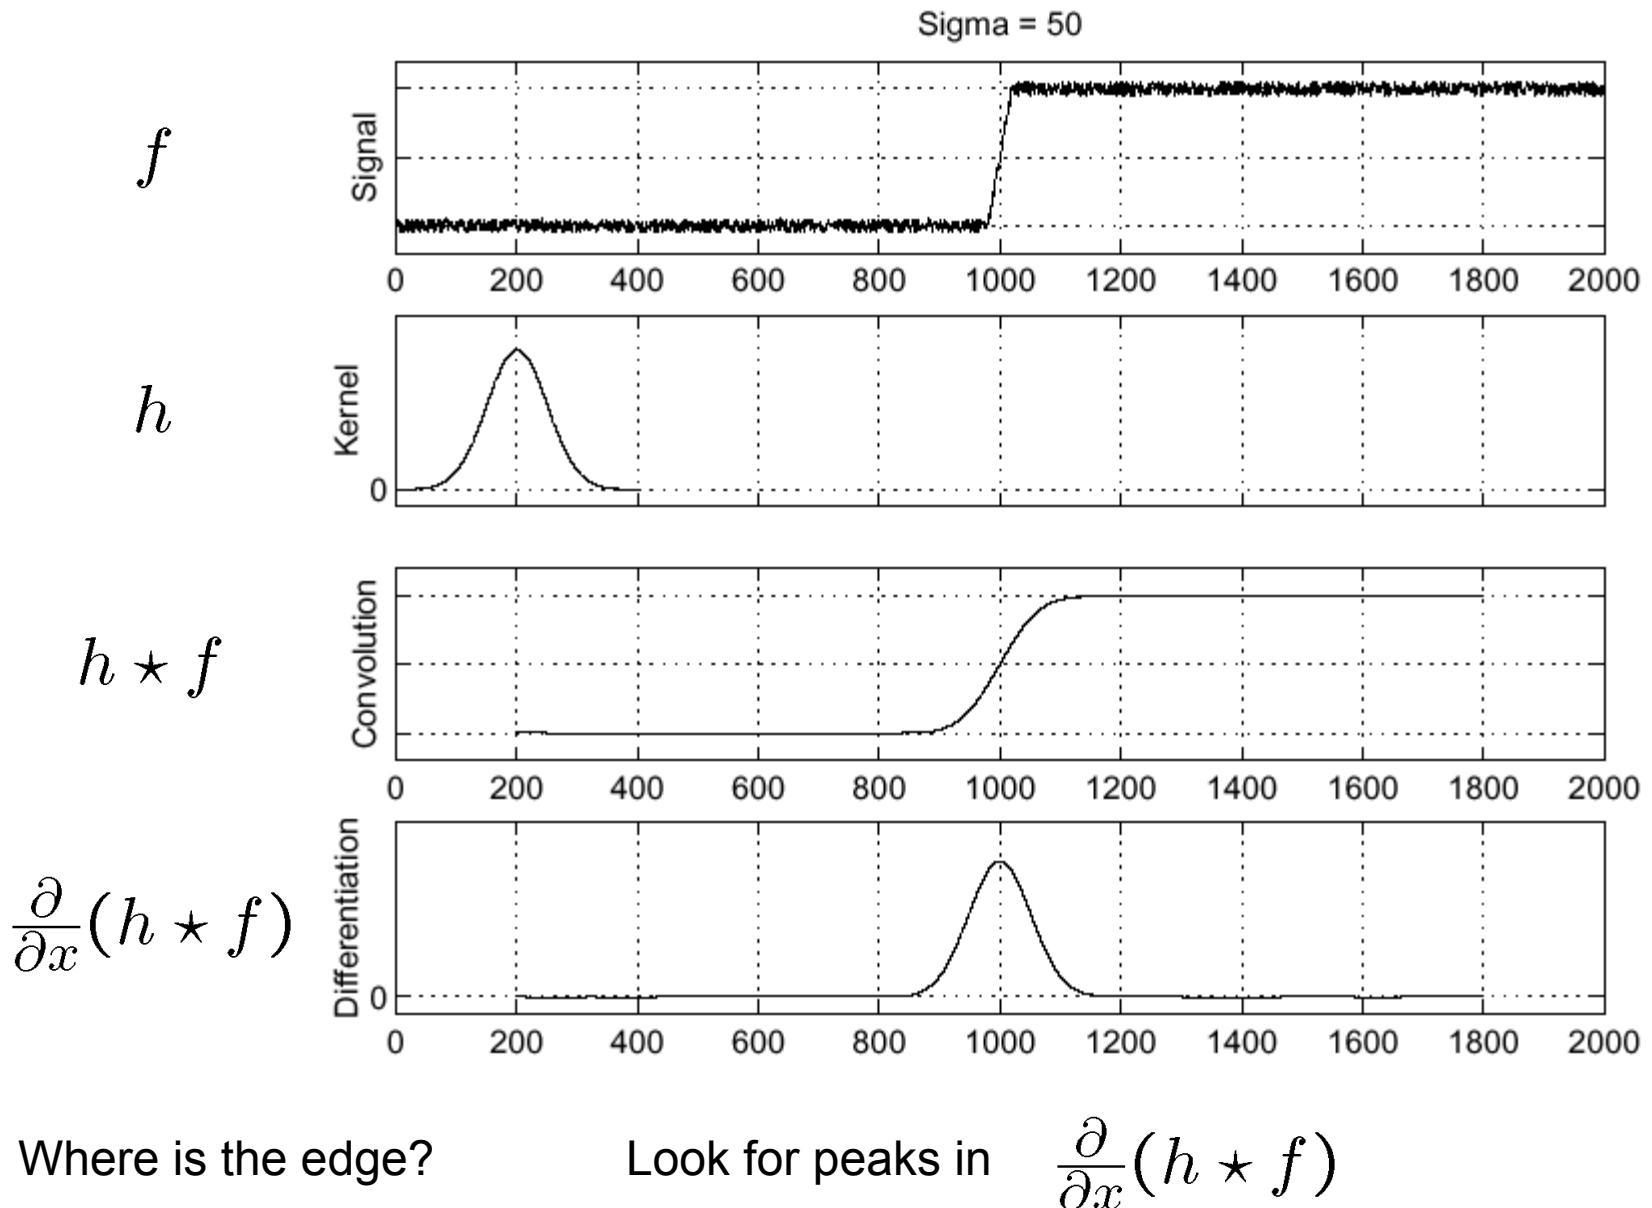
# Solution: smooth first



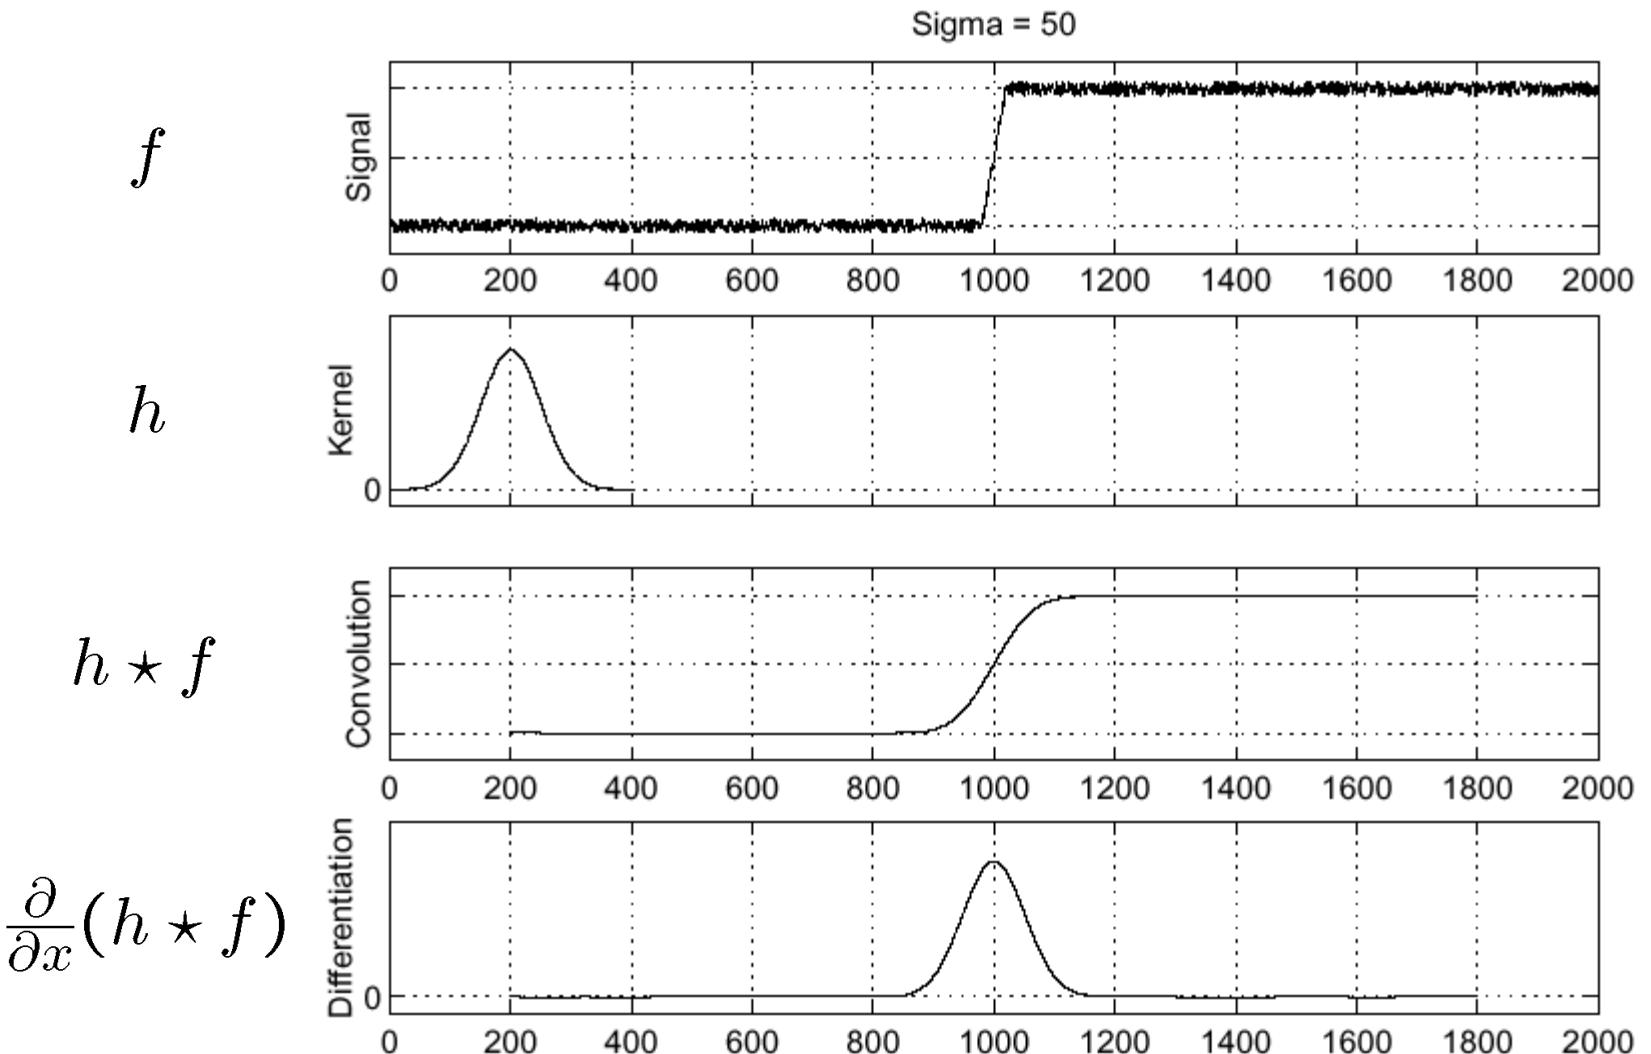
Where is the edge?

Look for peaks in  $\frac{\partial}{\partial x}(h \star f)$

# Solution: smooth first



# Solution: smooth first



Where is the edge?

Look for peaks in  $\frac{\partial}{\partial x}(h \star f)$  = zero's of the  $\frac{\partial^2}{\partial x^2}(h \star f)$  function

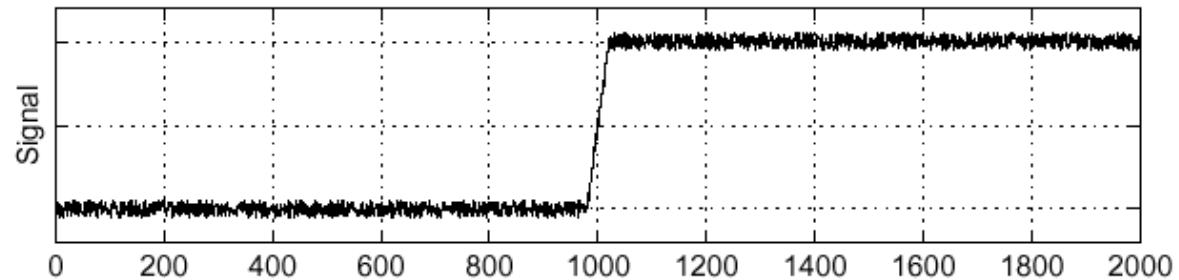
# Derivative theorem of convolution

$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

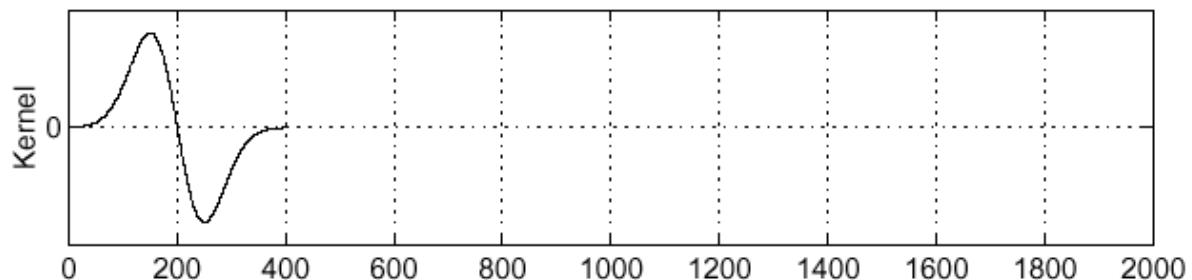
Differentiation property of convolution.

Sigma = 50

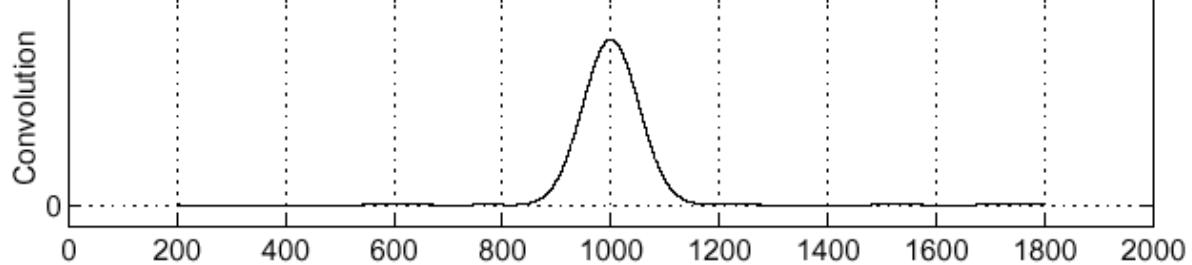
$f$



$\frac{\partial}{\partial x}h$

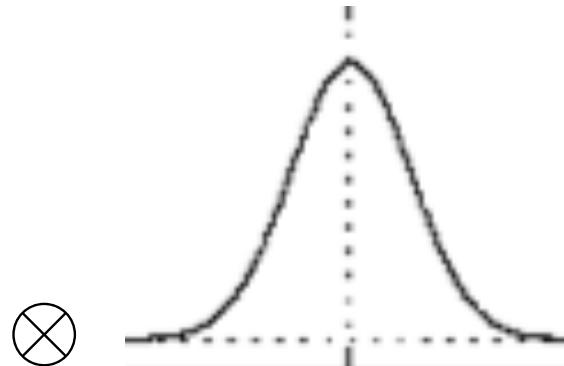


$(\frac{\partial}{\partial x}h) \star f$

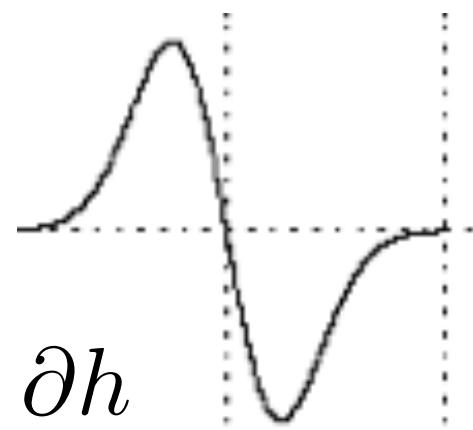


1D

-1	1
----	---



=



$d_x$

$h$

$\frac{\partial h}{\partial x}$

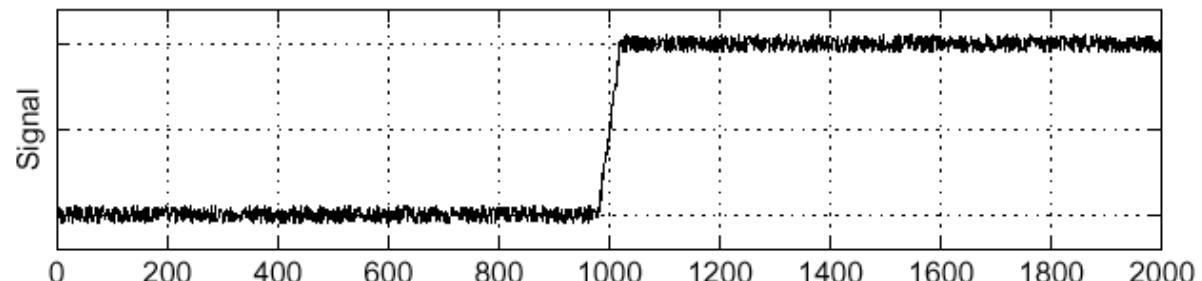
# Laplacian of a Gaussian – 1D

Consider

$$\frac{\partial^2}{\partial x^2}(h \star f)$$

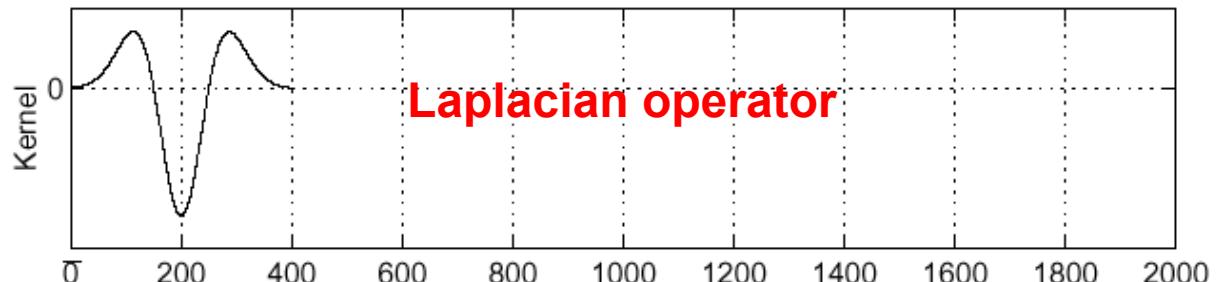
Sigma = 50

$f$

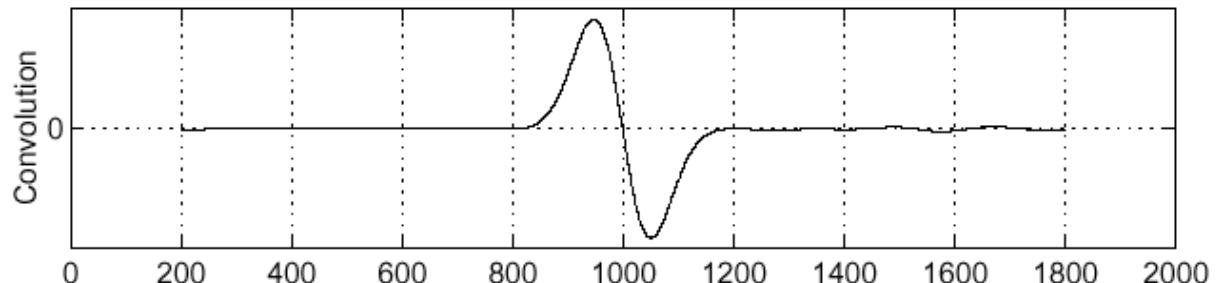


$$\frac{\partial^2}{\partial x^2} h$$

Laplacian operator



$$(\frac{\partial^2}{\partial x^2} h) \star f$$



Where is the edge?

Zero-crossing of the function  $(\frac{\partial^2}{\partial x^2} h) \star f$

# Derivative of Gaussian - 2D

$$d_x \otimes (h \otimes I) = (d_x \otimes h) \otimes I$$

derivation  
filter

Gaussian  
filter

image

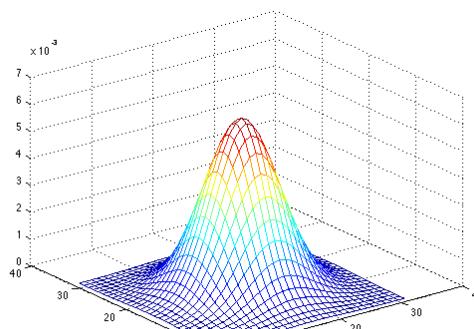
0	0	0
0	-1	1
0	0	0



$$\left[ \begin{array}{ccccc} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{array} \right]$$

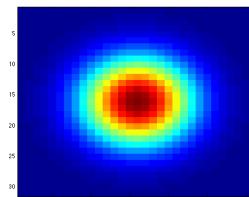


0	0	0
0	-1	1
0	0	0



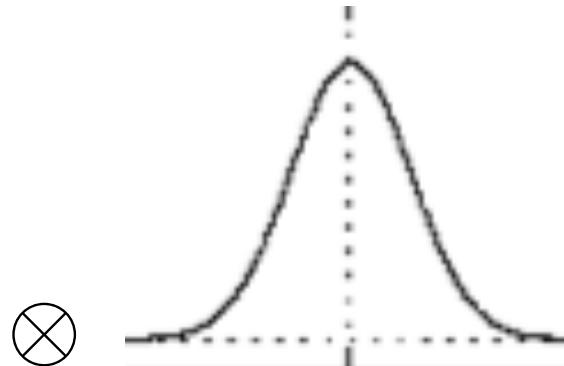
=

?

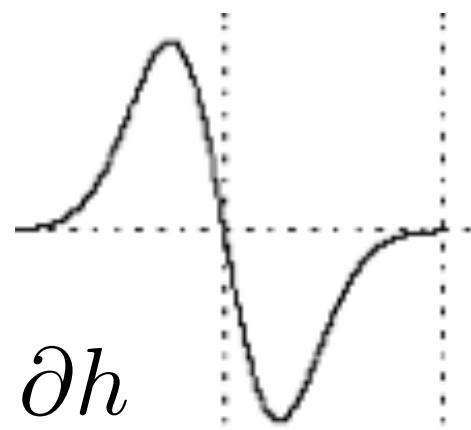


1D

-1	1
----	---



=



$d_x$

$h$

$\frac{\partial h}{\partial x}$

# Derivative of Gaussian filters - 2D

$$d_x \otimes (h \otimes I) = (d_x \otimes h) \otimes I$$

derivation  
filter

Gaussian  
filter

image

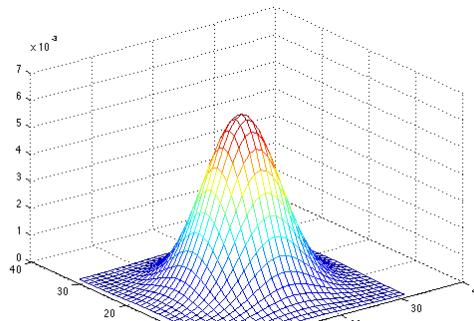
0	0	0
0	-1	1
0	0	0



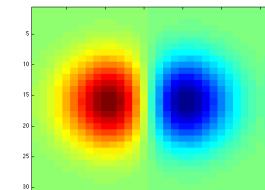
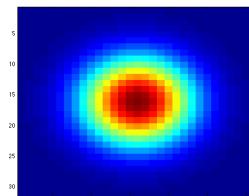
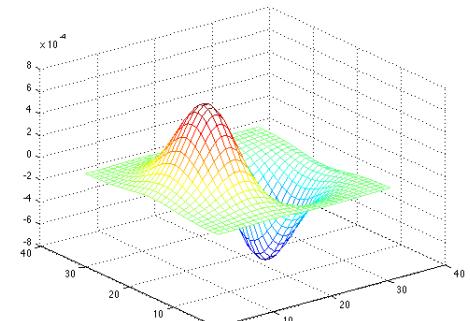
$$\left[ \begin{array}{ccccc} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{array} \right]$$



0	0	0
0	-1	1
0	0	0



=



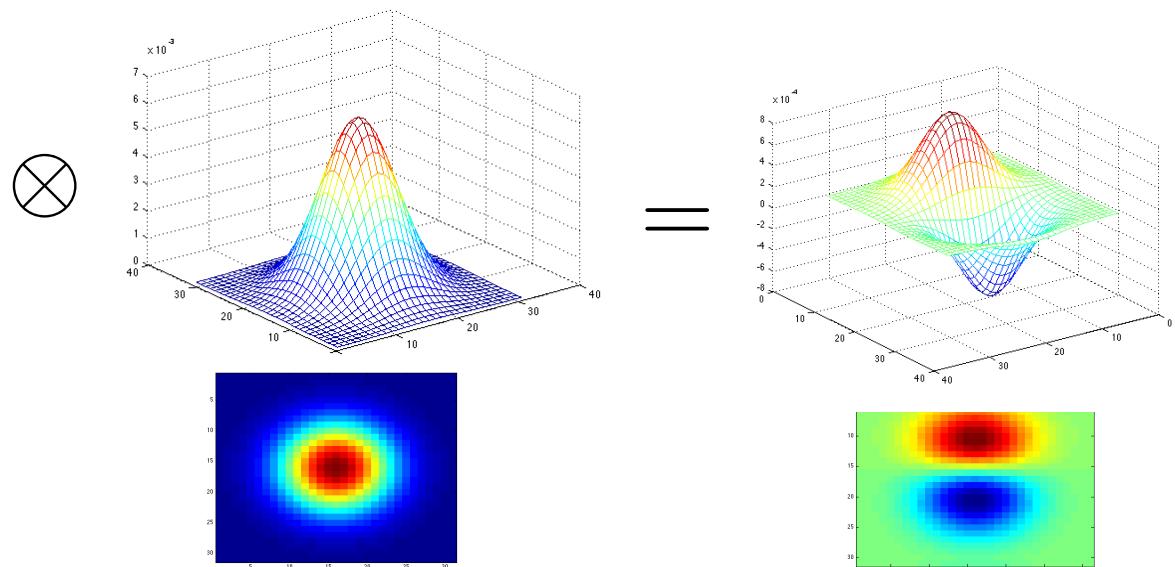
# Derivative of Gaussian filters - 2D

$$d_y \otimes (h \otimes I) = (d_y \otimes h) \otimes I$$

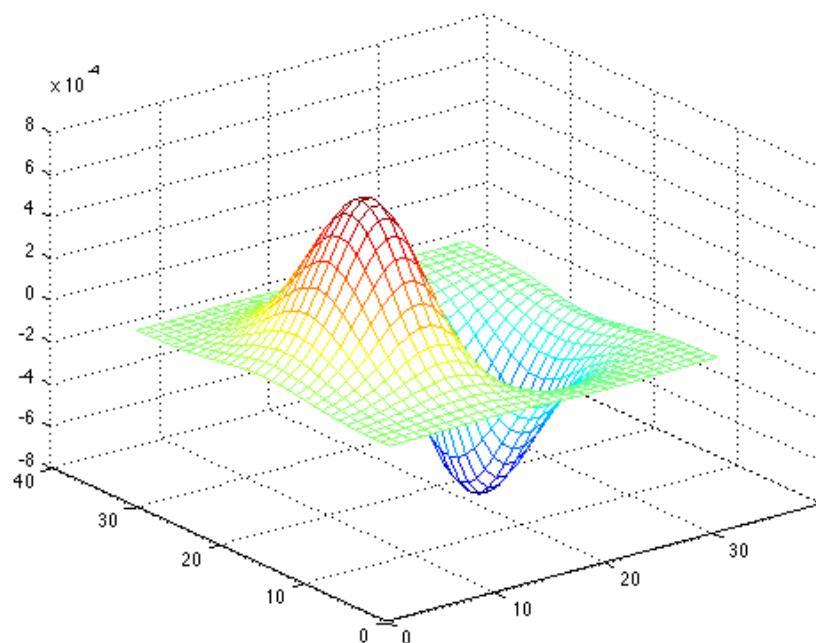
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

$$\otimes \begin{bmatrix} 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0219 & 0.0983 & 0.1621 & 0.0983 & 0.0219 \\ 0.0133 & 0.0596 & 0.0983 & 0.0596 & 0.0133 \\ 0.0030 & 0.0133 & 0.0219 & 0.0133 & 0.0030 \end{bmatrix}$$

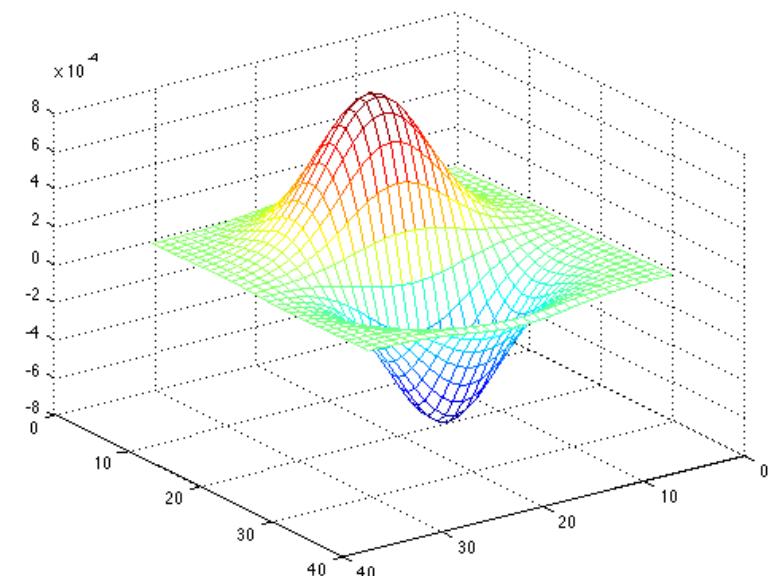
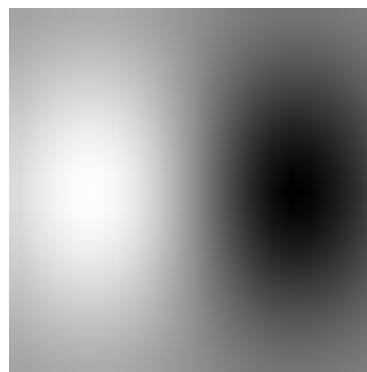
$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



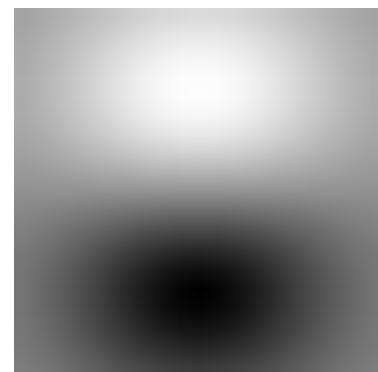
# Derivative of Gaussian filters



*x - direction*

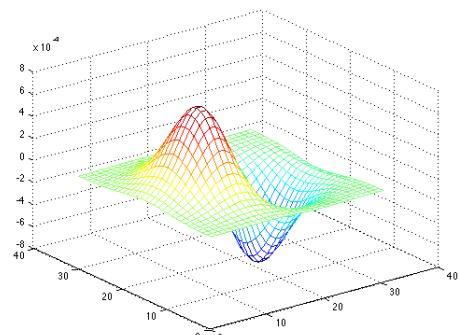


*y - direction*

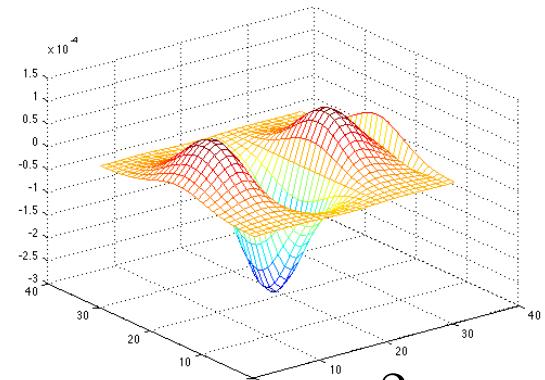


# Derivative of Gaussian filters - 2D

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$



=

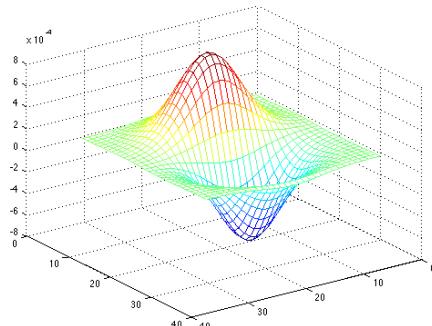


*direction x*

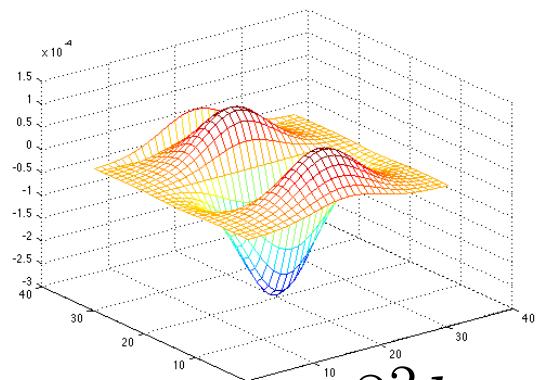
$$\frac{\partial h}{\partial x}$$

$$\frac{\partial^2 h}{\partial x^2}$$

$$\begin{bmatrix} 0 & 0 & 0 \\ 0 & -1 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$



=



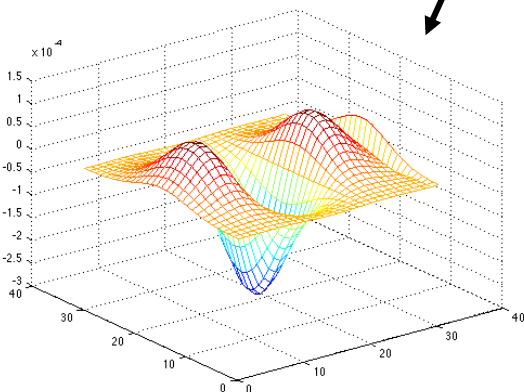
*direction y*

$$\frac{\partial h}{\partial y}$$

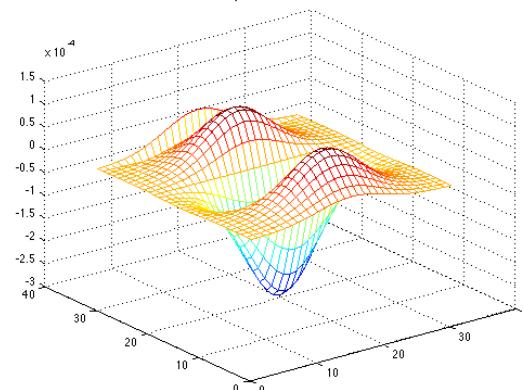
$$\frac{\partial^2 h}{\partial y^2}$$

# Laplacian of Gaussian

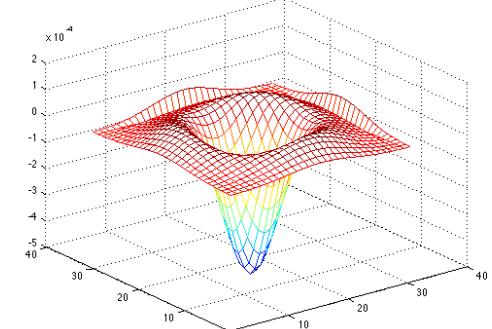
$$\nabla^2 h = \frac{\partial^2 h}{\partial x^2} + \frac{\partial^2 h}{\partial y^2}$$



+

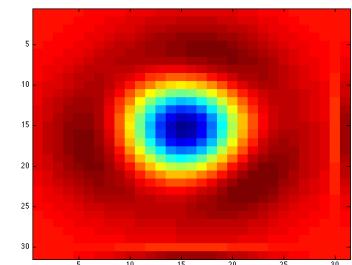


=

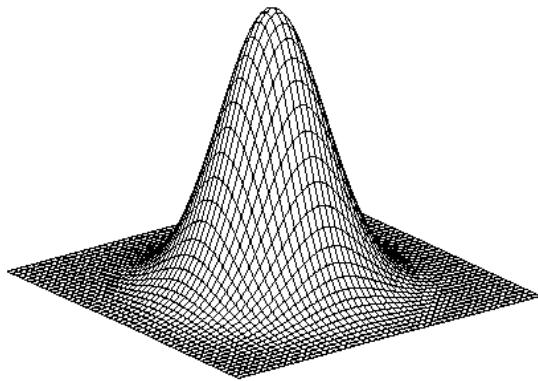


0	1	0
1	-4	1
0	1	0

Example of a  
Laplacian filter

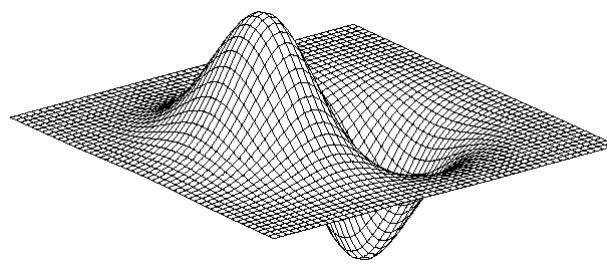


# 2D edge detection filters



**Gaussian**

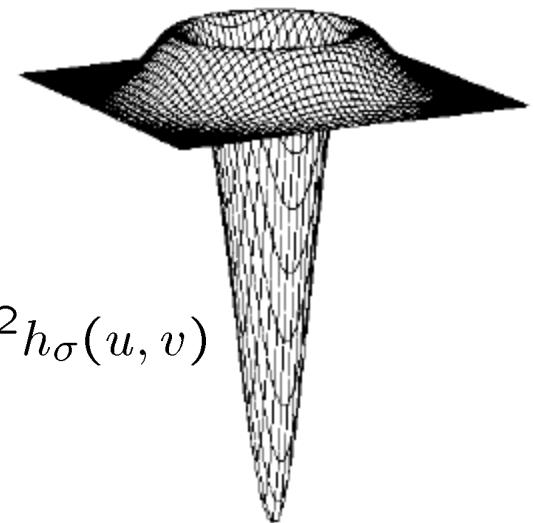
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



**derivative of Gaussian**

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

**Laplacian of Gaussian**



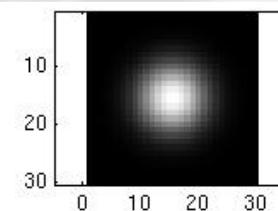
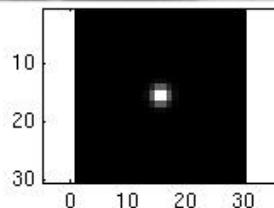
$$\nabla^2 h_\sigma(u, v)$$

- $\nabla^2$  is the Laplacian operator:

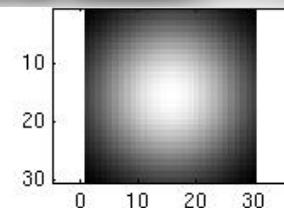
$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

# Smoothing with a Gaussian

Recall: parameter  $\sigma$  is the “scale” / “width” / “spread” of the Gaussian kernel, and controls the amount of smoothing.



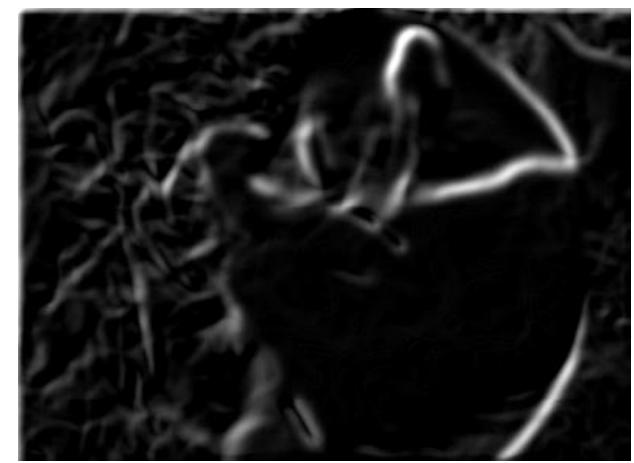
⋮ ⋮ ⋮



# Effect of $\sigma$ on derivatives



$\sigma = 1$  pixel



$\sigma = 3$  pixels

The apparent structures differ depending on Gaussian's scale parameter.

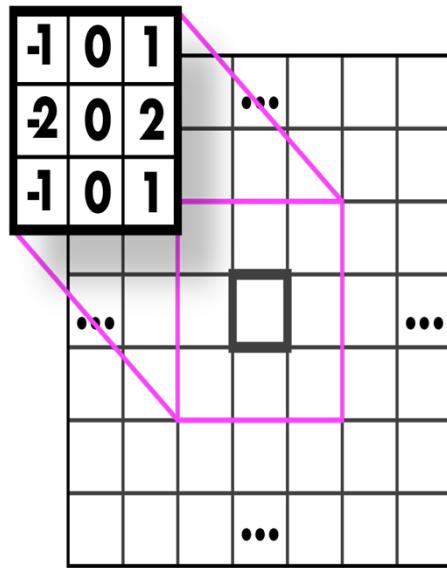
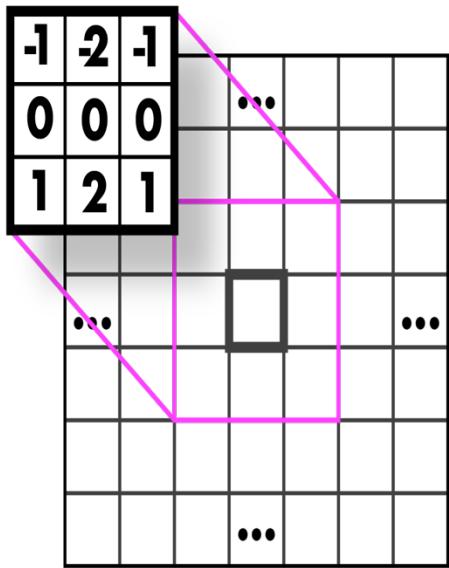
Larger values: larger scale edges detected  
Smaller values: finer features detected

# Mask properties

- Smoothing
  - Values positive
  - Sum to 1 → constant regions same as input
  - Amount of smoothing proportional to mask size
  - Remove “high-frequency” components; “low-pass” filter
- Derivatives
  - Opposite signs used to get high response in regions of high contrast
  - Sum to 0 → no response in constant regions
  - High absolute value at points of high contrast

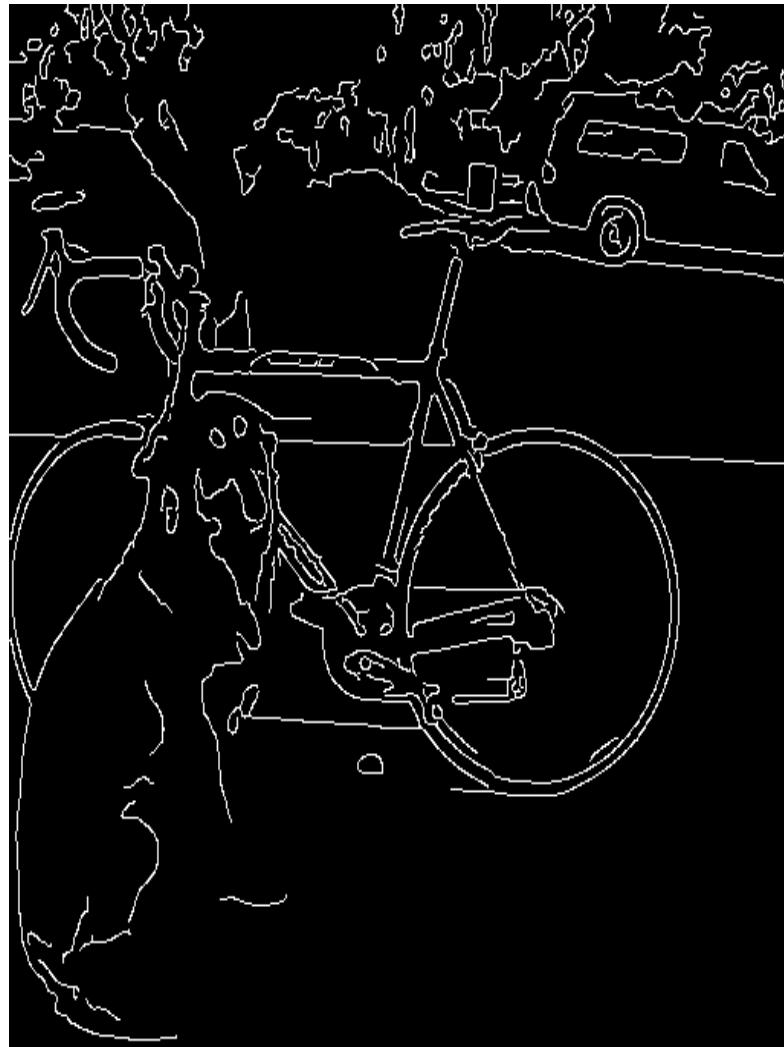
# Another approach: gradient magnitude

- Don't need 2nd derivatives
- Just use magnitude of gradient
- Are we done? No!





# What we really want: line drawing



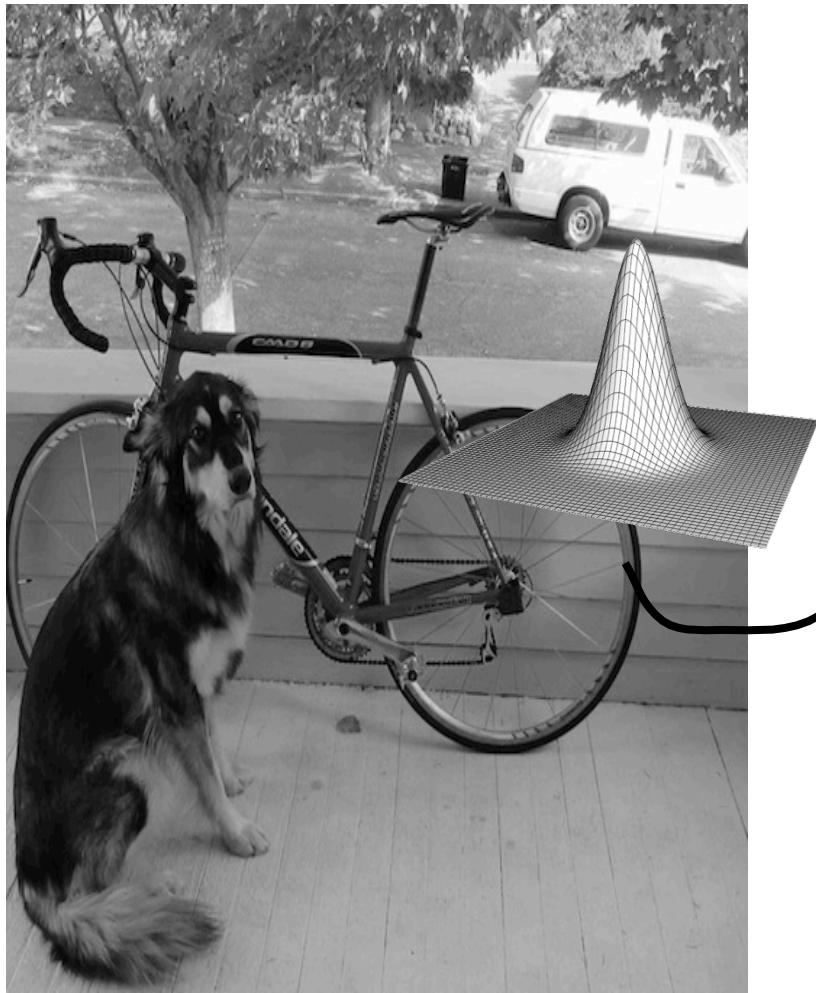
# Canny Edge Detection

## Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components

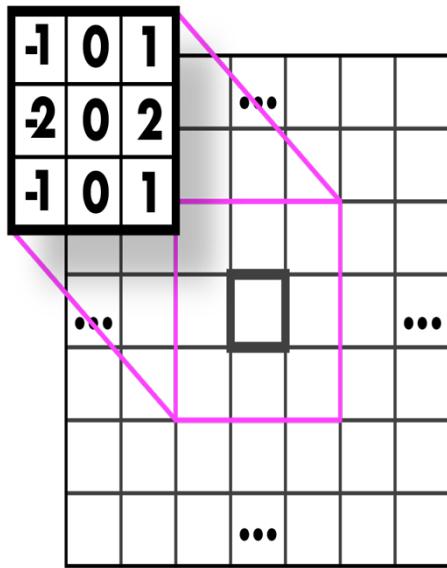
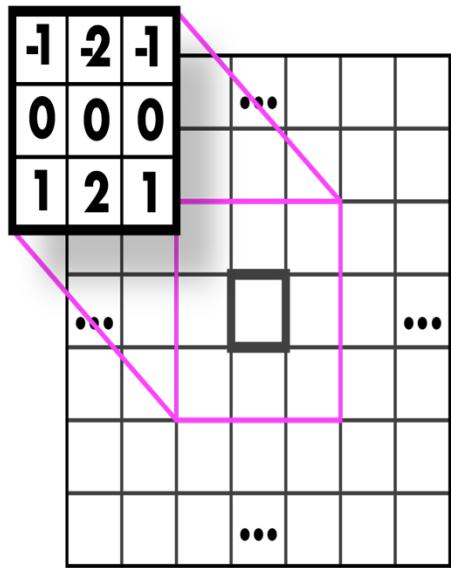
# Smooth image

- Apply a Gaussian filter



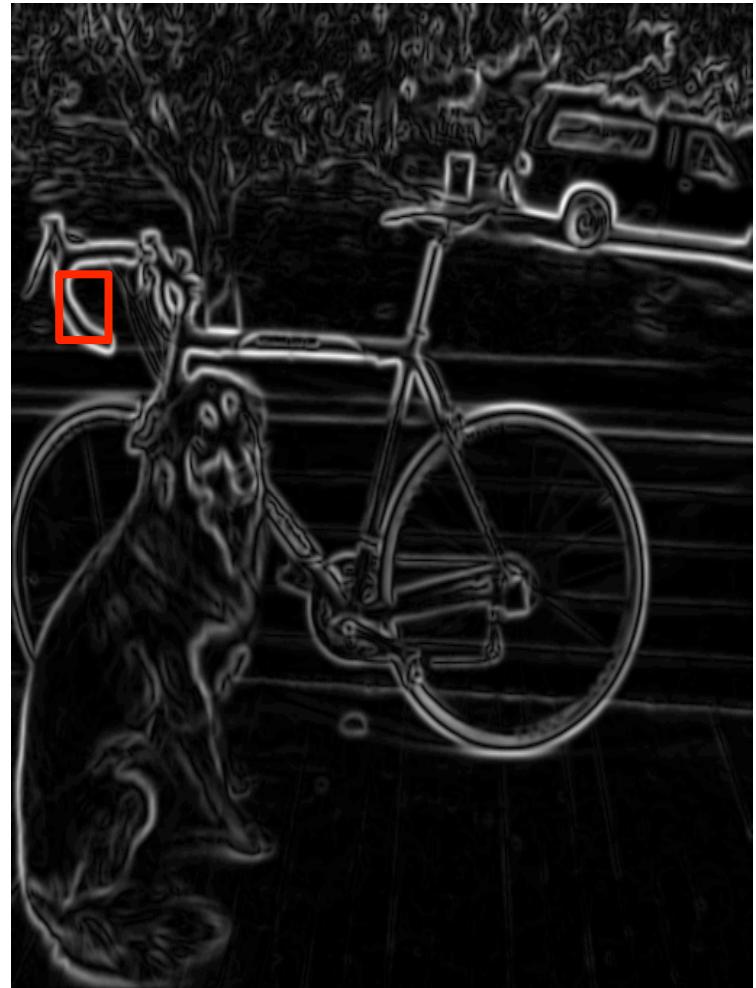
# Gradient magnitude and direction

- Sobel filter

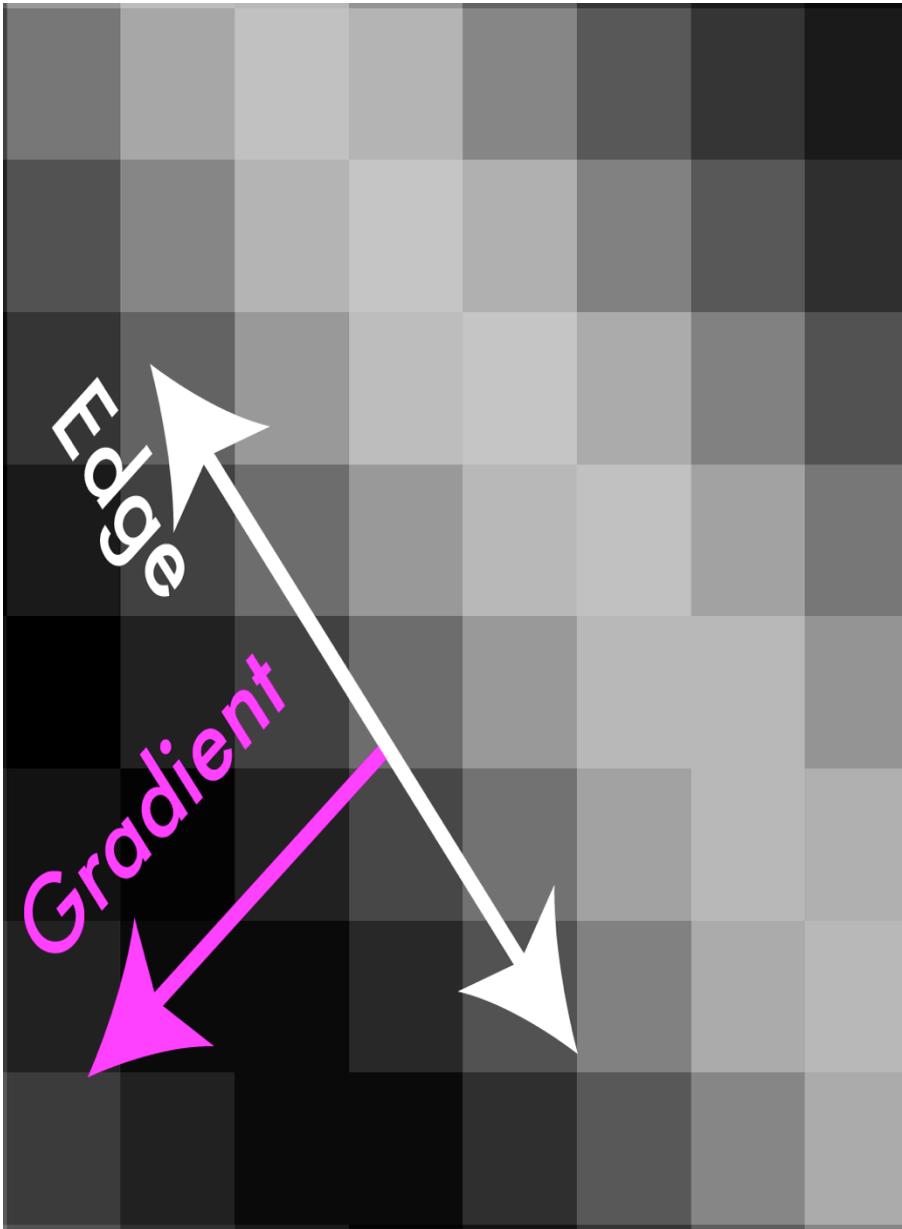


# Non-maximum suppression

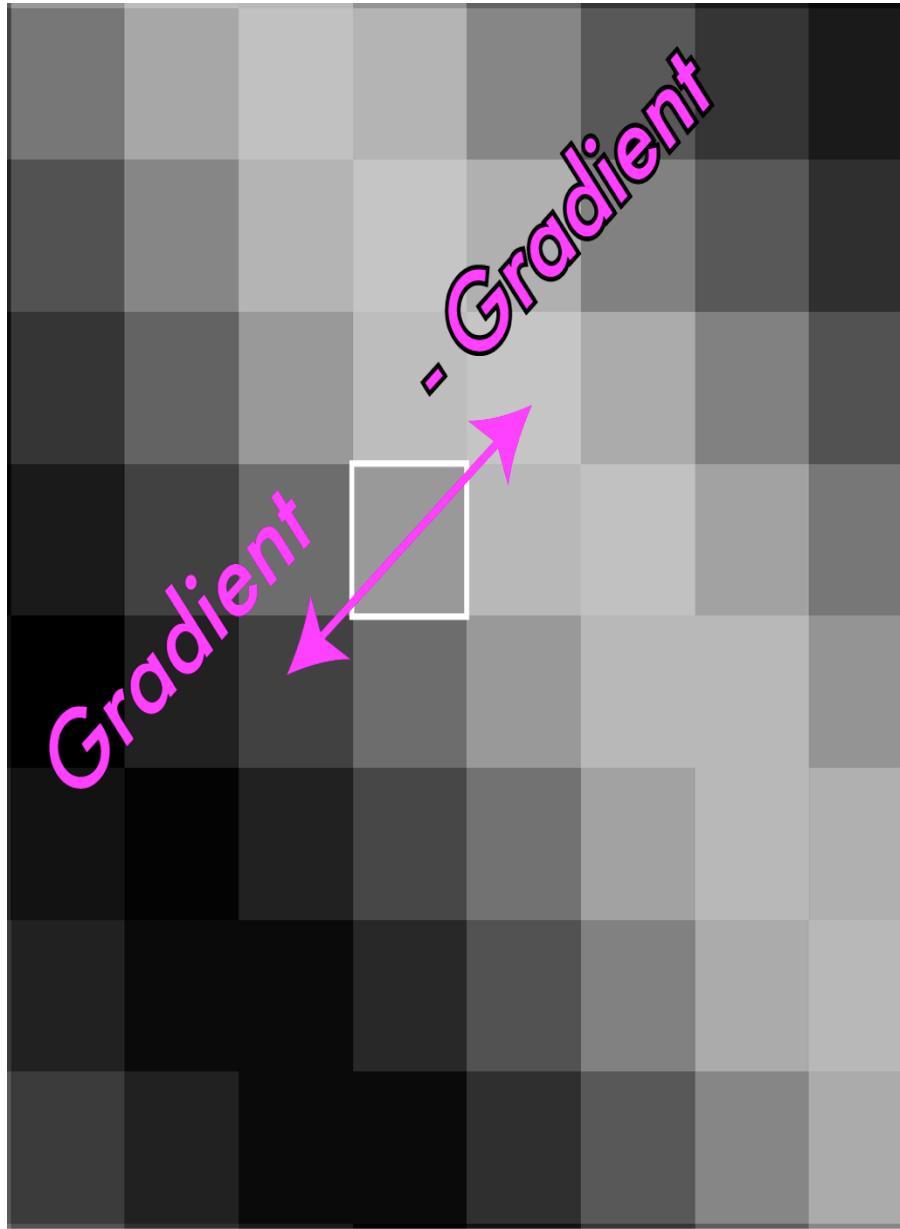
- Want single pixel edges, not thick blurry lines
- Need to check nearby pixels
- See if response is highest



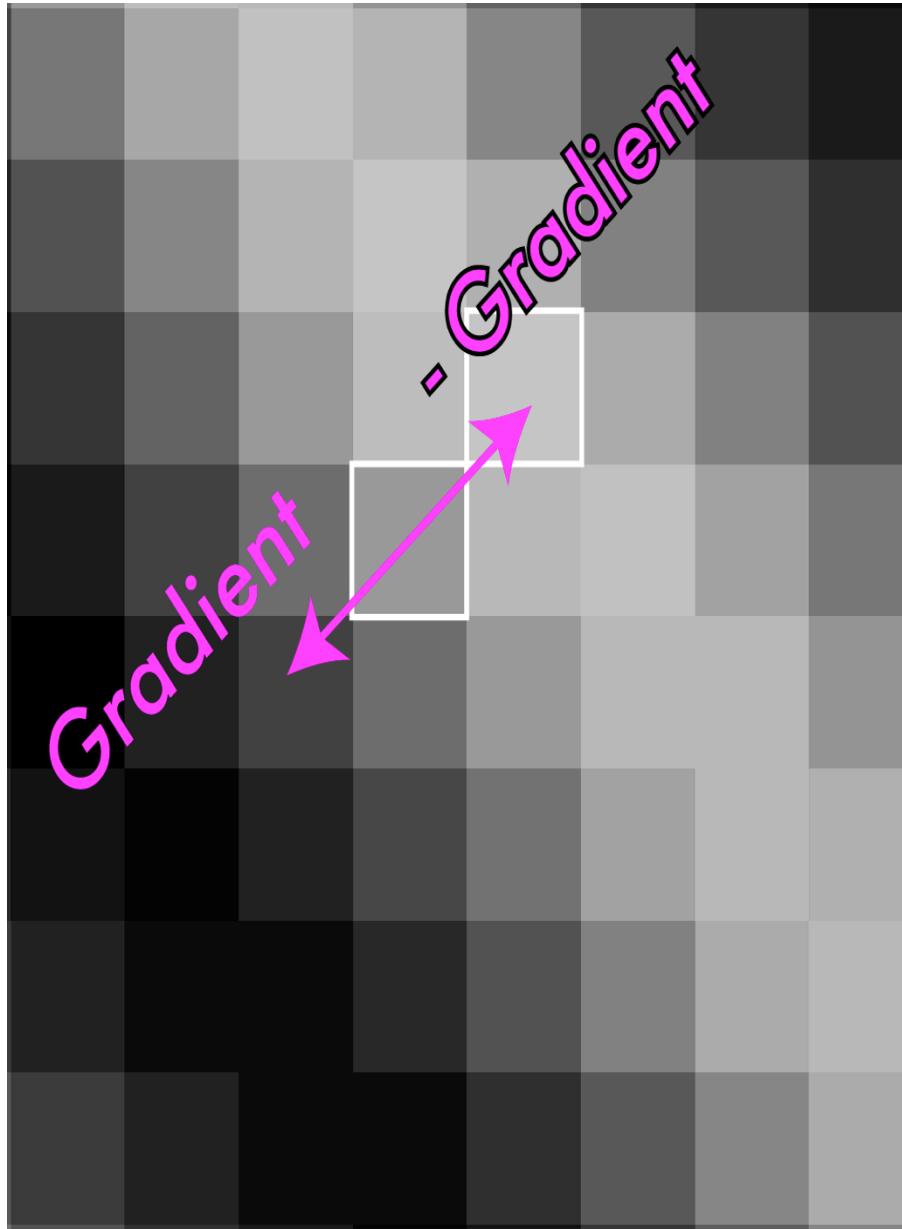
# Non-maximum suppression



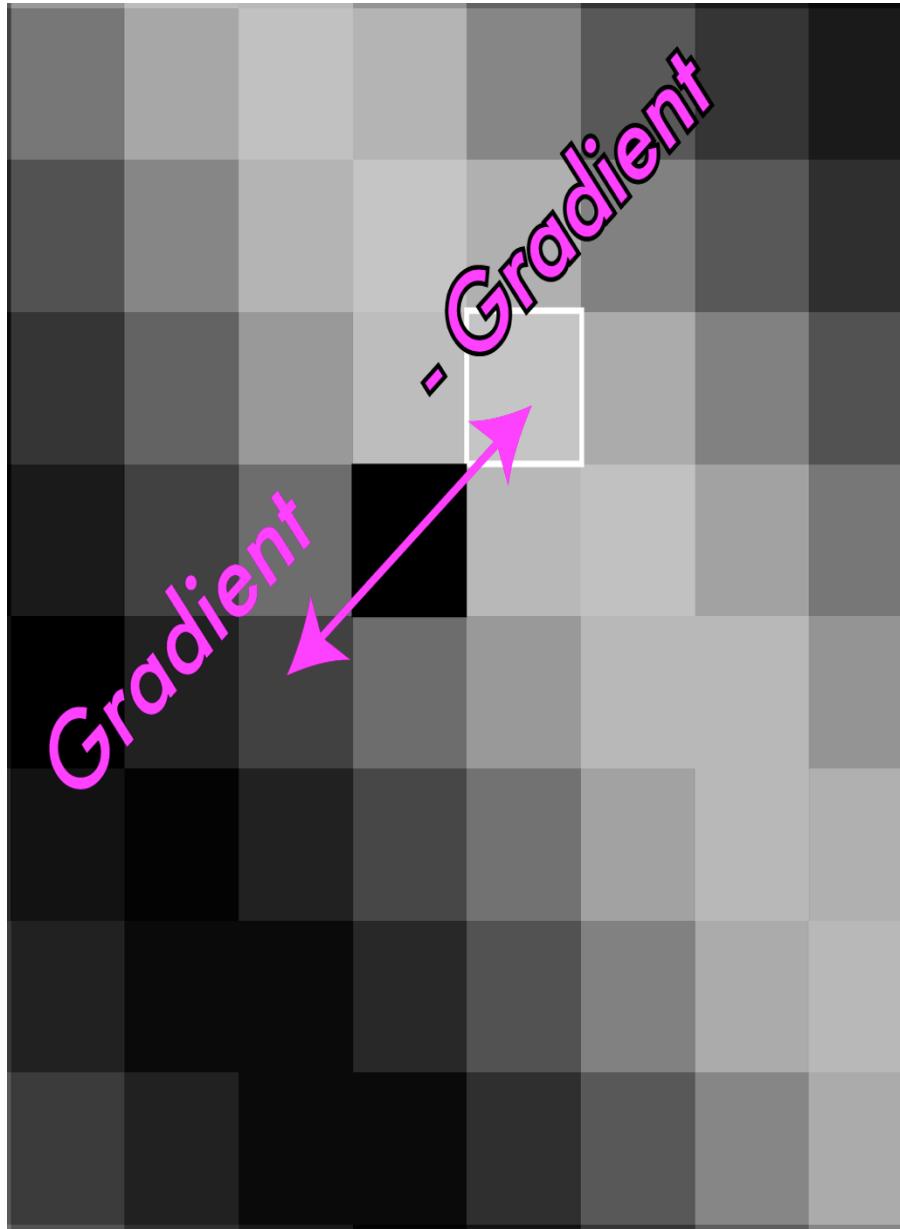
# Non-maximum suppression



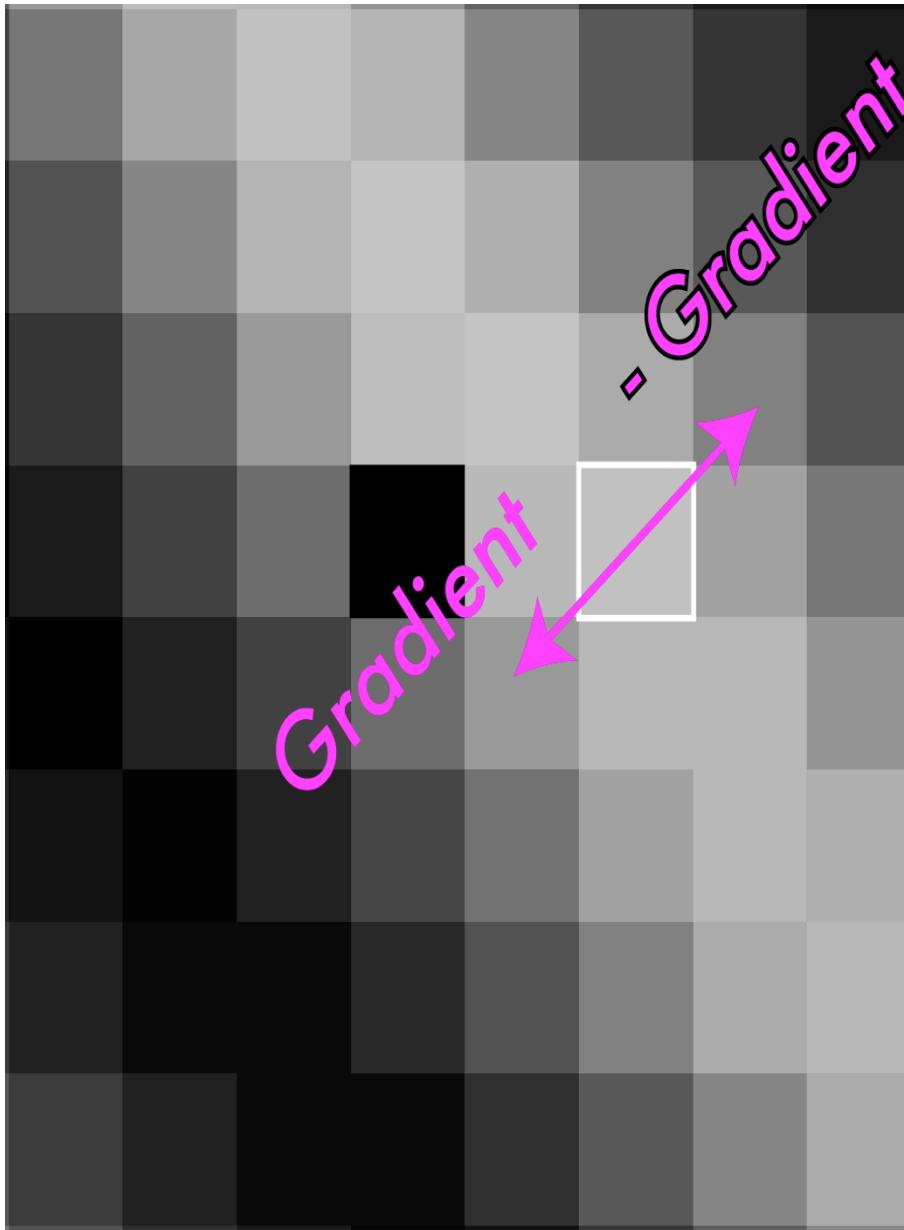
# Non-maximum suppression



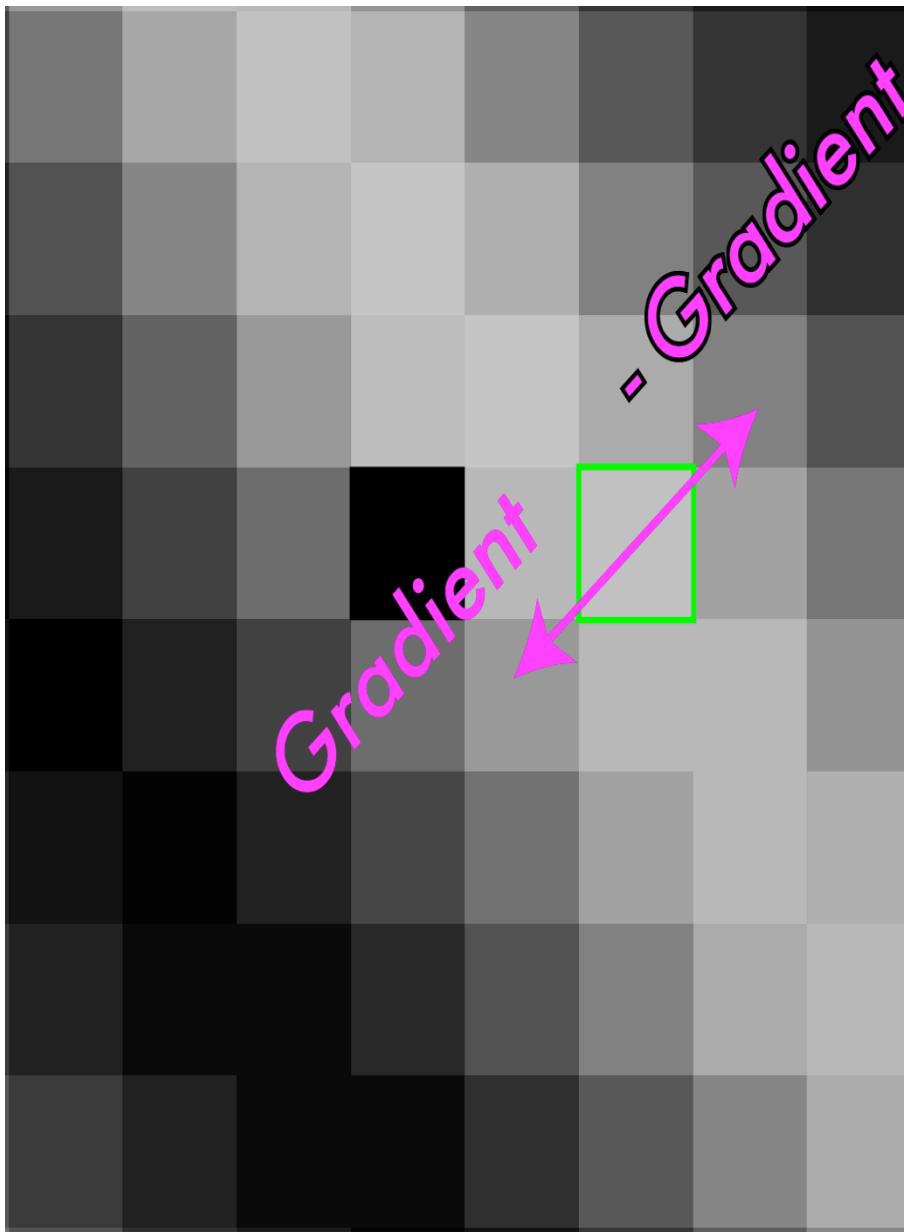
# Non-maximum suppression



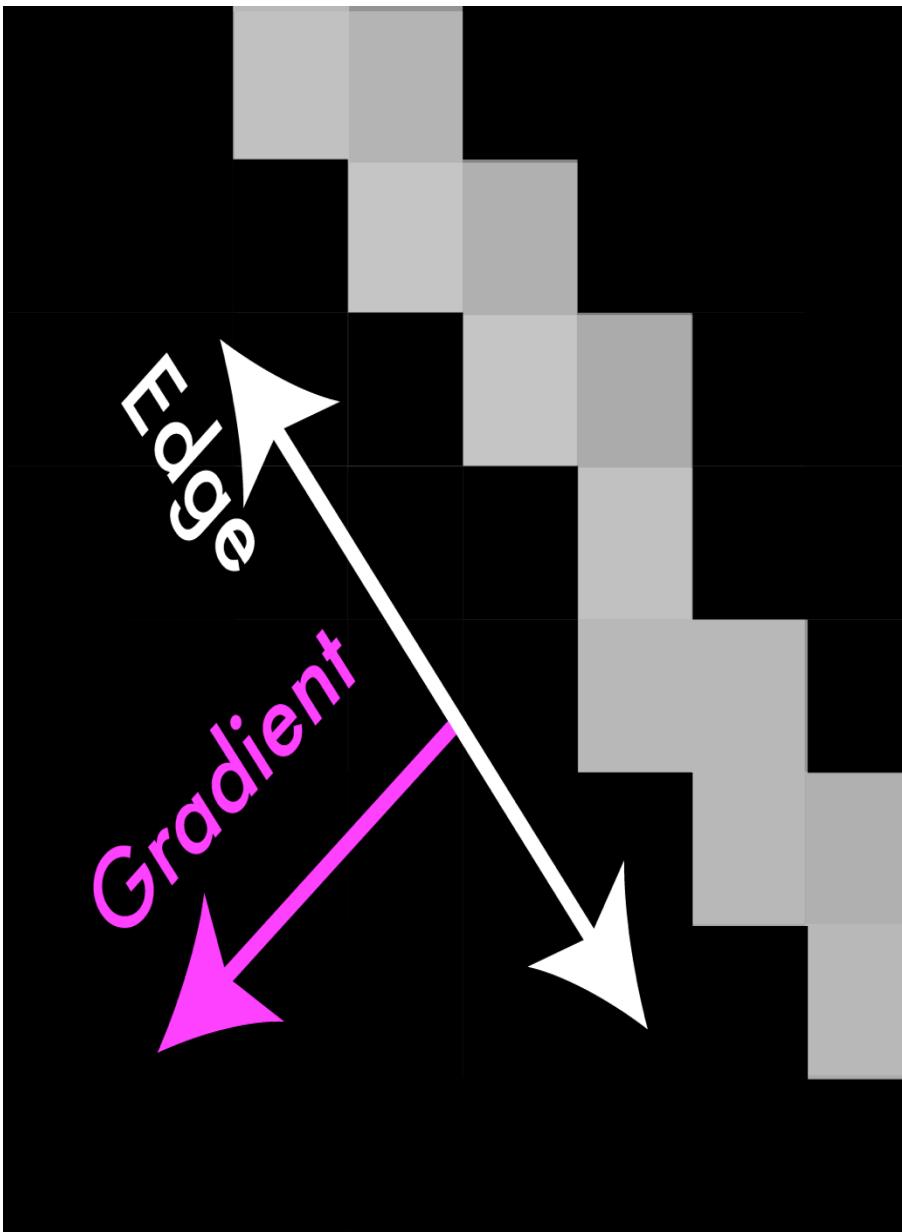
# Non-maximum suppression



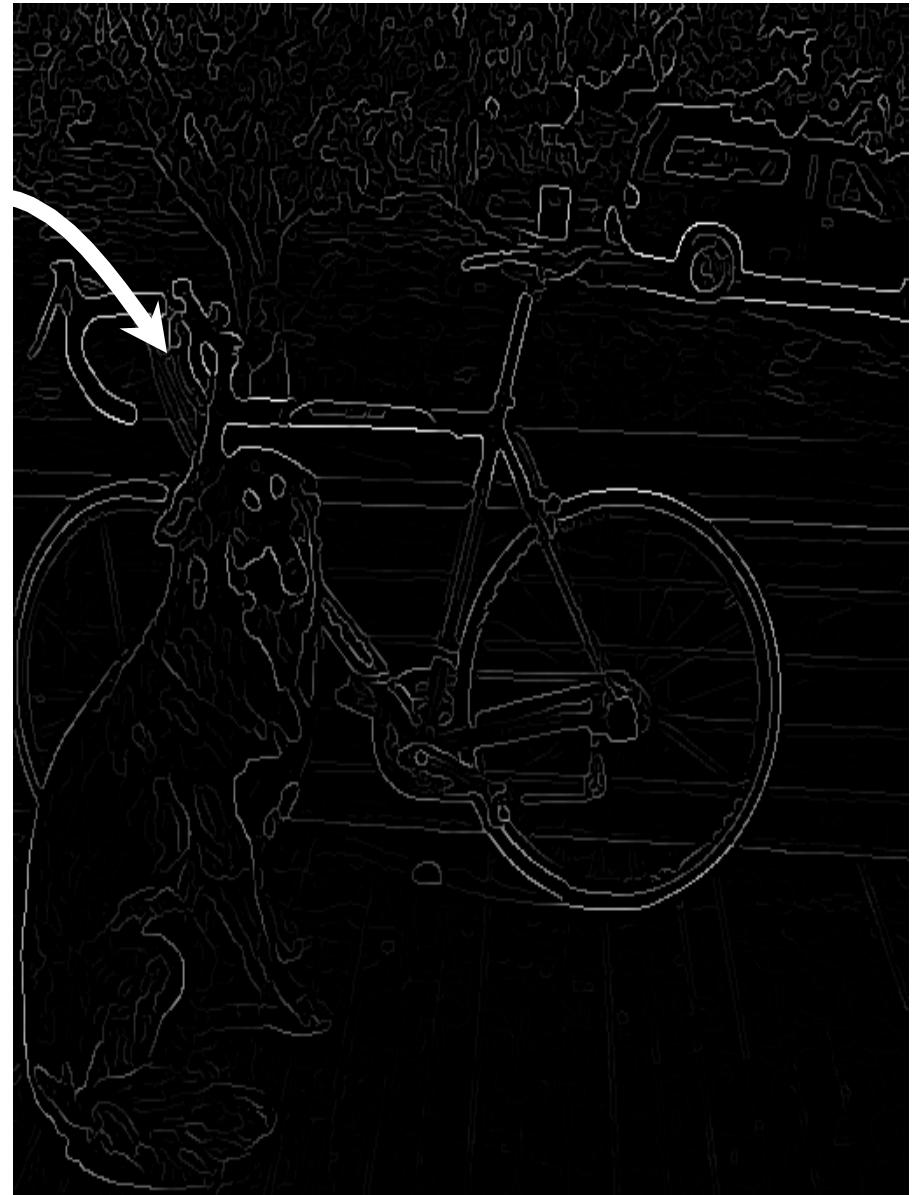
# Non-maximum suppression



# Non-maximum suppression

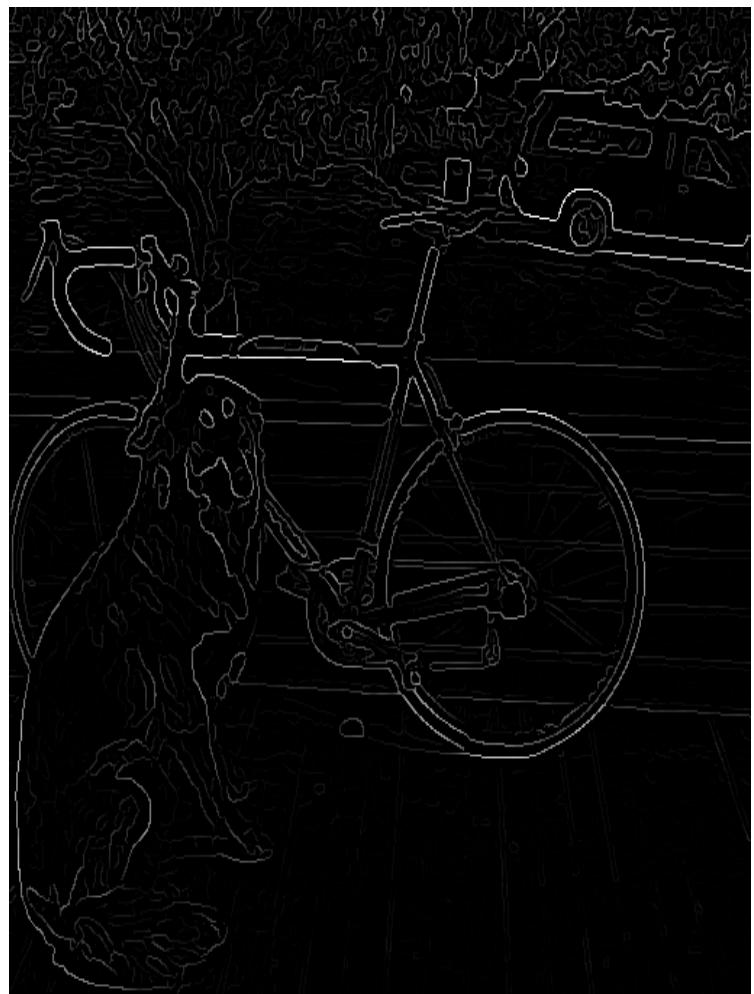


# Non-maximum suppression



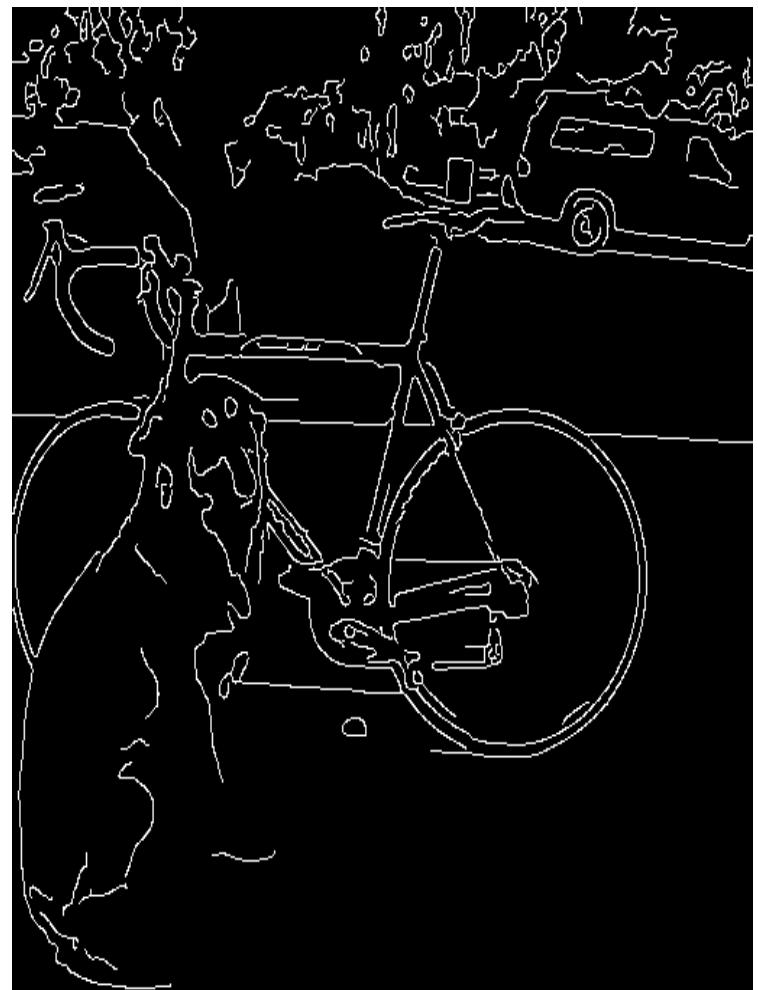
# Threshold edges

- Still some noise
- Only want strong edges
- 2 thresholds, 3 cases
  - $R > T$ : strong edge
  - $R < T$  but  $R > t$ : weak edge
  - $R < t$ : no edge
- Why two thresholds?



# Connect them up!

- Strong edges are edges!
- Weak edges are edges iff they connect to strong
- Look in some neighborhood (usually 8 closest)



# Canny Edge Detection

Algorithm:

- Smooth image (only want “real” edges, not noise)
- Calculate gradient direction and magnitude
- Non-maximum suppression perpendicular to edge
- Threshold into strong, weak, no edge
- Connect together components
- Tunable: Sigma, thresholds

# Canny Edge Detection

